

Data Integration and Large Scale Analysis

02 Data Warehousing and ETL

Dr. Lucas Iacono - 2025

Know Center Research GmbH & Graz University of Technology


Agenda

- Recap
- Data Warehousing (DWH)
- Extraction, Transformation, Loading (ETL)
- SQL/OLAP Extensions

Recap

Recap: Data Sources and Heterogeneity

Some important concepts and keywords




Register	Name	Class	Engine	Format
A	Corvette C7	GT3	Petrol	CSV
B	Porsche 911	GT3	Petrol	HDF5

- **Semantic** ---> Meaning (GT3 petrol race cars)
- **Ontology** ---> Conceptual Structure (Name/Class/Engine)
- **Format** ---> Physical Representation (Binary vs Text)

Recap: from data to models


AI Lifecycle



Recap: course goals

- Understand **major data integration architectures** and their role in modern data ecosystems.
- Apply **key techniques for data integration and cleaning** to ensure consistency, quality, and usability of data.
- Evaluate methods for **large-scale data storage and analysis**, with a focus on scalability and efficiency.

Today: DWH, ETL, SQL/OLAP Extensions



Data Warehousing

Complementary readings:

Wolfgang Lehner

Datenbanktechnologie für Data-Warehouse-Systeme

Konzepte und Methoden


Synthesis Lectures on Data Management

SYNTHESIS
COLLECTION OF TECHNOLOGY


Christian S. Jensen · Torben Bach Pedersen
Christian Thomsen

Multidimensional Databases and Data Warehousing


DWH: Integration Architecture




DWH: Integration Architecture




DWH: Integration Architecture



DWH: Integration Architecture



DWH: Integration Architecture




DWH: Motivations and Tradeoffs

- **Goal:** Queries over consolidated and cleaned data of several, potentially heterogeneous, data sources



DWH: Motivations and Tradeoffs

- **Goal:** Queries over consolidated and cleaned data of several, potentially heterogeneous, data sources






DWH: Motivations and Tradeoffs


- **Goal:** Queries over consolidated and cleaned data of several, potentially heterogeneous, data sources
- **Tradeoffs** (OLTP vs OLAP)
 - **Analytical query performance:** write vs read optimized data stores
 - **Virtualization:** overhead of remote access, source systems affected
 - **Consistency:** sync vs async changes, time regime -> up-to-date?
 - **Others:** flexibility, redundancy

DWH: Architecture


Operational source systems



DWH: Architecture



DWH: Architecture



DWH: Architecture

- **Data Warehouse (DWH):** “A data warehouse is a **subject-oriented, integrated, time-varying, non-volatile collection** of data in support of the management's **decision-making** process.” (Bill Inmon)
 - **Subject-oriented:** analysis-centric organization (e.g., sales) -> Data Mart
 - **Integrated:** consistent data from different data sources
 - **Time-varying:** History (snapshots of sources), and temporal modelling
 - **Non-volatile:** Read-only access, limited to periodic data loading by admin

DWH: Architecture


- **Data Warehouse (DWH):** “A data warehouse is a **subject-oriented, integrated, time-varying, non-volatile collection of data in support of the management's decision-making process.**” (Bill Inmon)
 - **Subject-oriented:** analysis-centric organization (e.g., sales) -> Data Mart
 - **Integrated:** consistent data from different data sources
 - **Time-varying:** History (snapshots of sources), and temporal modelling
 - **Non-volatile:** Read-only access, limited to periodic data loading by admin
- **DHW Instantiations:**
 - Data sources → staging areas → DWH → data marts (**top-down**)
 - Data marts → individual ETL → Data Sources → DWH (**bottom-up**)

DWH: Multi-dimensional Modeling: Data Cube




DWH: Multi-dimensional Modeling: Data Cube

- Central Metaphor: Data Cube
 - Qualifying data (categories, dimensions)
 - Quantifying data (cells)
 - Often sparse (0 for empty cells)





DWH: Multi-dimensional Modeling: Data Cube

- Central Metaphor: Data Cube
 - Qualifying data (categories, dimensions)
 - Quantifying data (cells)
 - Often sparse (0 for empty cells)
- Multi-dimensional Schema
 - Set of dimension hierarchies (D^1, \dots, D^n)
 - Set of measures (M^1, \dots, M^m)




DWH: Multi-dimensional Modeling: Data Cube

- Central Metaphor: Data Cube
 - Qualifying data (categories, dimensions)
 - Quantifying data (cells)
 - Often sparse (0 for empty cells)
- Multi-dimensional Schema
 - Set of dimension hierarchies (D^1, \dots, D^n)
 - Set of measures (M^1, \dots, M^m)
- Dimension Hierarchy
 - Partially-ordered set D of categorical attributes
 - Generic maximum element (Top_D)
 - Existing minimum element (Day)




DWH: Multi-dimensional Modeling: Data Cube

- Dimension Hierarchy, cont.
 - Classifying (categorical)
 - Define the hierarchy used for aggregations
 - Descriptive attributes
 - Describe data used for filtering
 - **Orthogonal dimensions:** there are no functional dependencies between attributes of different dimensions



DWH: Multi-dimensional Modeling: Data Cube


- Dimension Hierarchy, cont.
 - Classifying (categorical) vs descriptive attributes
 - **Orthogonal dimensions:** there are no functional dependencies between attributes of different dimensions
- Fact F
 - Base tuples w/ measures of summation type (**e.g units sold**)
 - Granularity
- Measure M
 - Metrics computed over non-empty subset of facts in schema
(e.g. units sold per year)
 - Scalar function (**temp offset**) vs aggregation function (**avg temp per day**)



DWH: Multi-dimensional Modeling: Operations


- **Slicing (cutting-out a single slice)**

- Select a “slice” of the cube by specifying a **filter condition on one of the dimensions** (categorical attributes)
- Same **data granularity** but subset of dimensions




- **Dicing (take a “chunk” of the cube)**

- Select a “sub-cube” by specifying a **filter condition on multiple dimensions**
- Complex Boolean expressions possible
- Sometimes slicing used synonym




DWH: Multi-dimensional Modeling: Operations

- **Roll-up (similar Merge) → Zooming-out**
 - Aggregation of facts or measures into **coarser-grained aggregates** (measures)
 - Same dimensions but different granularity
- **Drill-Down (similar Split) → Zooming-in**
 - Disaggregation of measures into **finer-grained** measures



DWH: Multi-dimensional Modeling: Operations

- **Pivot (How the data looks)**
 - Rotate cube by exchanging dimensions
 - Swap rows per columns



DWH: Aggregation Types

- **Recap: Classification of Aggregates**

- **Additive aggregation** functions (SUM, COUNT), can be summed across all cube dimensions [e.g. **Sales**]
- **Semi-additive aggregation** functions (MIN, MAX), can be summed across some cube dimensions but not in others like time [e.g. **Amount of Employees**]
- **Non-additive** functions, can't be summed across any dimension (AVG, STDDEV, VAR)

- **Summation Types of Measures**

- **FLOW**: arbitrary aggregation possible (e.g. **sales**)
- **STOCK**: aggregation possible, except over temporal dim (e.g. **bank account balance**)
- **VPU**: value-per-unit typically (e.g., **price per product**)

Prog	16/17	17/18	18/19	19/20	20/21	Total
CS	1153	1283	1321	1343	1368	?
SEM	928	970	939	944	985	?
ICE	804	868	846	842	849	?
Total	2855	3121	3106	3129	3202	?

DWH: Multi-dimensionality

- **MOLAP (Multi-Dim. OLAP)**
 - **OLAP server with native multi-dimensional data storage**
 - Dedicated query language: Multidimensional Expressions ([MDX](#))
 - [IBM Cognos](#) Powerplay, Essbase

DWH: Multi-dimensionality

- **MOLAP (Multi-Dim. OLAP)**
 - **OLAP server** with native multi-dimensional data storage
 - Dedicated query language: Multidimensional Expressions ([MDX](#))
 - [IBM Cognos](#) Powerplay, Essbase
- **ROLAP (Relation OLAP)**
 - **OLAP server** w/ storage in **RDBMS**
 - All commercial RDBMS vendors

DWH: Multi-dimensionality

- **MOLAP (Multi-Dim. OLAP)**
 - **OLAP server with native multi-dimensional data storage**
 - Dedicated query language: Multidimensional Expressions ([MDX](#))
 - [IBM Cognos](#) Powerplay, Essbase
- **ROLAP (Relation OLAP)**
 - **OLAP server w/ storage in RDBMS**
 - All commercial RDBMS vendors
- **HOLAP (Hybrid OLAP)**
 - **OLAP server w/ storage in RDBMS and multi-dimensional in-memory caches and data structures**

Requires mapping
to relational
model

DWH: Relational Data Model (RECAP)

- **Domain D**: set of possible values (INT, CHAR[20], BOOL).
- **Relation R (table)**:
 - **Relation Schema (RS)**: set of attributes $\{A_1, \dots, A_n\}$ with **domains**.
- **Additional Terminology**
 - **Tuple**: row of k elements in a relation. [1 with 3]
 - **Cardinality** (relational): number of tuples in the relation. [4]
 - **Rank (degree)**: number of attributes. [3]
 - **Semantics**:
 - Set semantics → no duplicates
 - Bag semantics → duplicates allowed (practical use)
 - **Order irrelevant**: neither tuples nor attributes have a defined order.

Attribute

A1 INT	A2 INT	A3 BOOL
3	7	T
1	2	T
3	4	F
1	7	T

Tuple

cardinality: 4
rank: 3

DWH: Relational Data Model (RECAP)

- **Domain (D)**: set of possible values (INT, CHAR[20], BOOL).
- **Relation (R)**: subset of the Cartesian product of domains.
 - Relation Schema (RS): set of attributes $\{A_1, \dots, A_n\}$ with domains.
- **Additional Terminology**
 - **Tuple**: row of k elements in a relation.
 - **Cardinality (relational)**: number of tuples in the relation.
 - **Rank (degree)**: number of attributes.
 - **Semantics**:
 - Set semantics → no duplicates
 - Bag semantics → duplicates allowed (practical use)
 - **Order irrelevant**: neither tuples nor attributes have a defined order.

Attribute

A1 INT	A2 INT	A3 BOOL
3	7	T
1	2	T
3	4	F
1	7	T

Tuple

cardinality: 4
rank: 3

DWH: Relational Data Model (RECAP)

- **Domain (D)**: set of possible values (INT, CHAR[20], BOOL).
- **Relation (R)**: subset of the Cartesian product of domains.
 - Relation Schema (RS): set of attributes $\{A_1, \dots, A_n\}$ with domains.
- **Additional Terminology**
 - **Tuple**: row of k elements in a relation.
 - **Cardinality (relational)**: number of tuples in the relation.
 - **Rank (degree)**: number of attributes.
 - **Semantics**:
 - Set semantics → no duplicates
 - Bag semantics → duplicates allowed (practical use)
 - **Order irrelevant**: neither tuples nor attributes have a defined order.

Attribute

A1 INT	A2 INT	A3 BOOL
3	7	T
1	2	T
3	4	F
1	7	T

Tuple

cardinality: 4
rank: 3


In ER Model its a relationship between entities
e.g. 1:N (one department -> many employees)

DWH: ROLAP - Star Schema






DWH: ROLAP - Snowflake Schema



DWH: Evolution of DWH/OLAP Workloads




Batch

DWH: Evolution of DWH/OLAP Workloads



DWH: Evolution of DWH/OLAP Workloads




DWH: Evolution of DWH/OLAP Workloads





DWH: Star schema example

Task: Given below ER diagram, create a **ROLAP star schema**. Data types can be ignored, but indicate **PK** and **FK** constraints.



DWH: Star schema example

Task: Given below ER diagram, create a ROLAP star schema. Data types can be ignored, but indicate PK and FK constraints.



Extraction, Transformation, Loading (ETL)

ETL: Overview

- **Concept**

- ETL process refers to the overall process of obtaining data from the source systems, cleaning and transforming it, and loading it into the DWH.
- Includes many integration and cleaning techniques

ETL: Overview

- **Concept**
 - ETL process refers to the overall process of obtaining data from the source systems, cleaning and transforming it, and loading it into the DWH.
 - Includes many integration and cleaning techniques
- **ETL**
 - Extract data from **heterogeneous sources**
 - **Transform** data via dedicated data flows or in **staging area**
 - **Load** cleaned and transformed data into **DWH**


ETL: Overview

- **Concept**
 - ETL process refers to the overall process of obtaining data from the source systems, cleaning and transforming it, and loading it into the DWH.
 - Includes many integration and cleaning techniques
- **ETL**
 - Extract data from heterogeneous sources
 - **Transform** data via dedicated data flows or in staging area
 - **Load** cleaned and transformed data into DWH
- **ELT**
 - Extract data from **heterogeneous sources**
 - **Load** raw data directly into DWH
 - **Transform** inside the DWH via SQL


ETL: TYPES OF HETEROGENEITY

Heterogeneity


ETL: TYPES OF HETEROGENEITY



ETL: TYPES OF HETEROGENEITY




ETL: TYPES OF HETEROGENEITY




- Synonyms / Homonyms: (Instructor vs lecturer)
- Simple Mapping: (12hr vs 24 hs)
- Language Expressions: (Database vs Datenbank)

ETL: TYPES OF HETEROGENEITY



ETL: TYPES OF HETEROGENEITY



ETL: Corrupted data

- **Heterogeneity of Data Sources**
 - Data comes from many different systems \longrightarrow Disorder.
- **Human Errors**
 - Data is **incomplete, missing, or wrongly labeled** (sometimes due to laziness or bias!).
- **Shifting Formats:**
 - The way data is saved **changes over time** \longrightarrow Inconsistency. (e.g. YYYY MM DD, DD MM YYYY)
- **Equipment Failure:**
 - Faulty hardware (like batteries or sensors)

ETL: Corrupted data

- **Heterogeneity of Data Sources**
 - Data comes from many different systems ----> Disorder.
- **Human Errors**
 - Data is incomplete, missing, or wrongly labeled (sometimes due to laziness or bias!).
- **Shifting Formats:**
 - The way data is saved changes over time ----> Inconsistency.
- **Equipment Failure:**
 - Faulty hardware (like batteries or sensors)

[Credit: Felix Naumann]

Uniqueness & duplicates		Contradictions & wrong values			Missing Values		Ref. Integrity
ID	Name	BDay	Age	Sex	Phone	Zip	Zip
3	Smith, Jane	05/06/1975	44	F	999-9999	98120	98120
3	John Smith	38/12/1963	55	M	867-4511	11111	San Jose
7	Jane Smith	05/06/1975	24	F	567-3211	98120	90001

Typos

ETL: Planning and Design Phase

- **Architecture, Flows, and Schemas**

- Plan requirements, architecture, tools
- Design high-level integration flows (systems, integration jobs)
- Data understanding (copy/code books, metadata)
- Design dimension & loading (static, dynamic incl keys)
- Design fact table & loading

ETL: Planning and Design Phase

- **Architecture, Flows, and Schemas**

- Plan requirements, architecture, tools
- Design high-level integration flows (systems, integration jobs)
- Data understanding (copy/code books, metadata)
- Design dimension & loading (static, dynamic incl keys)
- Design fact table & loading

- **Data Integration and Cleaning**

- Types of data sources (snapshot, APIs, query language, logs)
- Prepare schema mappings → see **04 Schema Matching and Mapping**
- Change data capture and incremental loading (diff, aggregates)
- Transformations, enrichments, and deduplication → **05 Entity Linking**
- Data validation and cleansing → see **06 Data Cleaning and Data Fusion**

ETL: Events and Change Data Capture

- **Architecture, Flows, and Schemas**

- Plan requirements, architecture, tools
- Design high-level integration flows (systems, integration jobs)
- Data understanding (copy/code books, metadata)
- Design dimension loading (static, dynamic incl keys)
- Design fact table loading

- **Data Integration and Cleaning**

- Types of data sources (snapshot, APIs, query language, logs)
- Prepare schema mappings → see **04 Schema Matching and Mapping**
- Change data capture and incremental loading (diff, aggregates)
- Transformations, enrichments, and deduplication → **05 Entity Linking**
- Data validation and cleansing → see **06 Data Cleaning and Data Fusion**

- **Optimization**

- Partitioning schemes for loaded data (e.g., per month)
- Maintenance

ETL: Events and Change Data Capture

Goal: Monitoring operations of data sources for detecting changes

- **Explicit Messages/Triggers**

- Setup update propagation from the source systems to middleware (e.g. every night at 23:00)
- **Asynchronously** propagate the updates into the DWH

ETL: Events and Change Data Capture

Goal: Monitoring operations of data sources for detecting changes

- **Explicit Messages/Triggers**
 - Setup update propagation from the source systems to middleware (e.g. every night at 23:00)
 - **Asynchronously** propagate the updates into the DWH
- **Log-based Capture**
 - Parse system logs / provenance to **retrieve changes since last loading**
 - Leverage explicit audit columns or internal timestamps

ETL: Events and Change Data Capture


Goal: Monitoring operations of data sources for detecting changes

- **Explicit Messages/Triggers**
 - Setup update propagation from the source systems to middleware
 - Asynchronously propagate the updates into the DWH
- **Log-based Capture**
 - Parse system logs / provenance to retrieve changes since last loading
 - Sometimes combined w/ replication → **03 MoM, EAI, and Replication**
 - Leverage explicit audit columns or internal timestamps
- **Snapshot Differences**
 - Compute **difference between old and new snapshot** (e.g., files) before loading
 - Broadly applicable but **more expensive**

ETL: Example Flow



Simitsis [et.al.](#) XPAD: a platform for analytic data flows. In Proceedings of the 2013 ACM SIGMOD International Conference on Management of Data (SIGMOD '13). Association for Computing Machinery,



SQL/OLAP Extensions

SQL/OLAP Extensions: Multi-Groupings

- **Recap: GROUP BY**

- Group tuples by categorical variables (e.g. year)
- Aggregate per group (SUM of Revenue)

`SELECT Year, SUM(Revenue)
FROM Sales
GROUP BY Year`

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

→

Year	SUM
2004	60
2005	30

SQL/OLAP Extensions: Multi-Groupings

- **Recap: GROUP BY**

- Group tuples by categorical variables
- Aggregate per group


`SELECT Year, SUM(Revenue)
FROM Sales
GROUP BY Year`

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

→

Year	SUM
2004	60
2005	30

- **Grouping Extensions**




SQL/OLAP Extensions: Multi-Groupings - Grouping Sets

- **Semantics**

- **What It Does:** It allows you to generate aggregated results (sums, averages, etc.) for multiple, specific combinations of columns within a single query.
- **Why Use It?** It's a great shortcut that often offers better performance than writing many separate GROUP BY clauses and manually joining them with UNION ALL.
- **In Short:** You get multiple reports (views) with just one command.


- **Example**



SQL/OLAP Extensions: Multi-Groupings - Grouping Sets

- Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY GROUPING SETS
    ((), (Year), (Year,Quarter))
```




Granularity

Year	Quarter	Revenue																																																															
2004	1	10	<table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th style="background-color: #e6f2ff;">Year</th> <th style="background-color: #e6f2ff;">Quarter</th> <th style="background-color: #e6f2ff;">SUM</th> <th colspan="2"></th> </tr> </thead> <tbody> <tr> <td>-</td> <td>-</td> <td>90</td> <td colspan="2" style="text-align: right;">()</td> </tr> <tr> <td>2004</td> <td>-</td> <td>60</td> <td colspan="2" rowspan="2" style="text-align: right;">(Year)</td> </tr> <tr> <td>2005</td> <td>-</td> <td>30</td> </tr> <tr> <td>2004</td> <td>1</td> <td>10</td> <td colspan="2" rowspan="5" style="text-align: right; vertical-align: middle;"> <table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th colspan="2"></th> <th colspan="2"></th> <th style="background-color: #e6f2ff;">(Y, Q)</th> </tr> </thead> <tbody> <tr> <td colspan="2">2004</td> <td>1</td> <td>20</td> <td>(Y, Q)</td> </tr> <tr> <td colspan="2">2004</td> <td>2</td> <td>10</td> <td></td> </tr> <tr> <td colspan="2">2004</td> <td>3</td> <td>20</td> <td></td> </tr> <tr> <td colspan="2">2004</td> <td>4</td> <td>30</td> <td></td> </tr> </tbody> </table> </td> </tr> <tr> <td>2004</td> <td>2</td> <td>20</td> </tr> <tr> <td>2004</td> <td>3</td> <td>10</td> </tr> <tr> <td>2004</td> <td>4</td> <td>20</td> </tr> <tr> <td>2005</td> <td>1</td> <td>30</td> </tr> </tbody> </table>			Year	Quarter	SUM			-	-	90	()		2004	-	60	(Year)		2005	-	30	2004	1	10	<table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th colspan="2"></th> <th colspan="2"></th> <th style="background-color: #e6f2ff;">(Y, Q)</th> </tr> </thead> <tbody> <tr> <td colspan="2">2004</td> <td>1</td> <td>20</td> <td>(Y, Q)</td> </tr> <tr> <td colspan="2">2004</td> <td>2</td> <td>10</td> <td></td> </tr> <tr> <td colspan="2">2004</td> <td>3</td> <td>20</td> <td></td> </tr> <tr> <td colspan="2">2004</td> <td>4</td> <td>30</td> <td></td> </tr> </tbody> </table>						(Y, Q)	2004		1	20	(Y, Q)	2004		2	10		2004		3	20		2004		4	30		2004	2	20	2004	3	10	2004	4	20	2005	1	30
Year	Quarter	SUM																																																															
-	-	90				()																																																											
2004	-	60				(Year)																																																											
2005	-	30																																																															
2004	1	10	<table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th colspan="2"></th> <th colspan="2"></th> <th style="background-color: #e6f2ff;">(Y, Q)</th> </tr> </thead> <tbody> <tr> <td colspan="2">2004</td> <td>1</td> <td>20</td> <td>(Y, Q)</td> </tr> <tr> <td colspan="2">2004</td> <td>2</td> <td>10</td> <td></td> </tr> <tr> <td colspan="2">2004</td> <td>3</td> <td>20</td> <td></td> </tr> <tr> <td colspan="2">2004</td> <td>4</td> <td>30</td> <td></td> </tr> </tbody> </table>						(Y, Q)	2004		1	20	(Y, Q)	2004		2	10		2004		3	20		2004		4	30																																					
					(Y, Q)																																																												
2004		1			20	(Y, Q)																																																											
2004		2			10																																																												
2004		3			20																																																												
2004		4	30																																																														
2004	2	20																																																															
2004	3	10																																																															
2004	4	20																																																															
2005	1	30																																																															

SQL/OLAP Extensions: Multi-Groupings - Rollup

- **Semantics**


- **What it does:** It's a command that generates sequential subtotals across a hierarchy you define, culminating in a **Grand Total** of all the data.
- **Why use It?** Perfect for reports where you need to see sales by country, then by country and city, and finally the worldwide total.
- **In short:** The ROLLUP syntax is just a much cleaner way to write out a long and repetitive GROUPING SETS command.



SQL/OLAP Extensions: Multi-Groupings - Rollup

- Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY ROLLUP(Year,Quarter)
```



Granularity

Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

→


Year	Quarter	SUM
-	-	90
2004	-	60
2005	-	30
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

(Y, Q) → (Year)

SQL/OLAP Extensions: Multi-Groupings - Cube

- **Semantics**

- **What it does:** It creates all possible subtotals from the columns you define, including the **grand total**. It gives you **every single combination of the data**.
- **Why use It?** It's essential for Multi-Dimensional Analysis (OLAP) where users might need to **slice the data by any combination of fields**.
- **In short:** Like ROLLUP, the CUBE syntax is a concise, high-performance way to write out an enormous, exhaustive GROUPING SETS command.




SQL/OLAP Extensions: Multi-Groupings - Cube

- Example

```
SELECT Year, Quarter, SUM(Revenue)
FROM R
GROUP BY CUBE(Year,Quarter)
```

= GROUP BY (), (Y), (Q), (Y,Q))



Year	Quarter	Revenue
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30



Year	Quarter	SUM
-	-	90
2004	-	60
2005	-	30
-	1	40
-	2	20
-	3	10
-	4	20
2004	1	10
2004	2	20
2004	3	10
2004	4	20
2005	1	30

Summary and Q&A

Summary and Q&A

- **Data Warehousing (DWH)**
 - DWH architecture
 - Multidimensional modeling
- **Extraction, Transformation, Loading (ETL)**
 - ETL process, errors, (and data flows → )
- **SQL/OLAP Extensions**
 - Multi-grouping operations
- **Next lecture: Data Integration Architectures**
 - October 17. Message-oriented Middleware, EAI, and Replication + **Project Presentation**

Many thanks!