

Data Integration and Large Scale Analysis

09- Distributed Data Storage

Dr. Lucas Iacono - 2025

Agenda

- Announcements
- Motivation and Terminology
- Data Lakes
- Object Stores
- Distributed File Systems
- Key-Value Stores and Cloud DBMS (+ eWarehouses)

Announcements

Announcements

- **Project**
 - January 10. **Project Submission Deadline**
- **Final Exams**
 - January 30th (First Exam)
 - March 20th (Second Exam)

Motivation and Terminology

Motivation and Terminology

Overview Distributed Data Storage: **Distributed DBS (#03)**

- **What?**
 - A DBS is a virtual (logical) database that appears as a **single database** to the user but is composed by **multiple physical databases** located in **different physical locations**.
- **Why?**
 - Store and process data efficiently (e.g. data spread across different geographical locations).
 - Lets users access and work with data as if it were all in one place, even though it's distributed.

Motivation and Terminology

Overview Distributed Data Storage: Components for Global Query Processing

What if you run a query in a DBS?

1. **Identify** Figure out which physical databases contain the requested information.
2. **Unify** Combine the data from multiple databases as if it came from a single source.
3. **Optimize** Ensure the distributed system is fast and efficient when handling queries.

Motivation and Terminology

Overview Distributed Data Storage: DBS Types

- **Virtual DBS (homogeneous):**
 - All databases use the same technology and schema.
 - **Example:** Several PostgreSQL databases distributed across locations, all with same tables, attributes and relationships.
- **Federated DBS (heterogeneous):**
 - The databases can use different technologies or schemas.
 - **Example:** Two PostgreSQL databases with different entities, relationships and attributes.

Motivation and Terminology

Overview Distributed Data Storage: Cloud & Distributed Storage

Why?

1. **Large-scale:** handle **very large amounts** of data.
2. **Semi-structured/nested.** Data doesn't align with traditional rows and columns (**JSON, XML**).
3. **Fault tolerance:** Data **available** and **reliable** despite system components' failures

Motivation and Terminology

Overview Distributed Data Storage: Cloud & Distributed Storage

Types: Cloud Storage

1. Block Storage (e.g. AWS EBS):

- a. Data splitted into **blocks**, which can be individually read or written.
- b. Used for systems that need fast, low-level access to data.

Motivation and Terminology

Overview Distributed Data Storage: Cloud & Distributed Storage

Types: Cloud Storage

1. Object Storage (e.g. AWS S3):
 - a. Data stored as objects (data, metadata, and UID).
 - b. Ideal for storing unstructured data like media files, backups, or large datasets.
 - c. Objects of a limited size (e.g., 5TB in AWS S3).

Motivation and Terminology

Overview Distributed Data Storage: Cloud & Distributed Storage

Types: **Distributed file systems**

1. **Distributed File Systems (e.g. HDFS):**
 - a. File systems built on top of **block** or **object** storage to manage files across multiple servers.
 - b. Allow for large-scale file sharing and processing.

Motivation and Terminology

Overview Distributed Data Storage: Cloud & Distributed Storage

Types: Database as a Service (DBaaS) - 1

NoSQL Databases (e.g. Redis, MongoDB):

- a. Target: Designed for flexibility and scalability.
- b. Types:
 - i. Key-Value Stores: Store **data** as **key-value** pairs
 - ii. Document Stores: Store **data** as **documents**

Motivation and Terminology

Overview Distributed Data Storage: Cloud & Distributed Storage

Types: Database as a Service (DBaaS) - 2

Cloud DBMSs (e.g. Amazon RDS, Google Cloud SQL):

- a. **Target:** handle OLTP and OLAP workloads.
- b. Combine SQL databases with cloud scalability and flexibility.

Motivation and Terminology

Central Data Abstractions: **Files and Objects**

- **File:** large and continuous block of data saved in a specific format (CSV, JSON, Binary, etc.).
- **Object:** like a file, but binary and it includes metadata (E.g. Images on S3)

Motivation and Terminology

Central Data Abstractions: **Distributed Collections**

- **Logical multiset (bag)** of key-value pairs (unsorted collection)
- **Different physical representations** key-value pairs can be stored in various ways (e.g., database, across files, or in memory).
- **Easy Distribution via Horizontal Partitioning.** Data divided into "chunks" (shards or partitions) based on the keys. Each chunk stored on a different machine (easier to handle large-scale data).
- **How collections are created:** from single file with data or a folder of files (even if they're messy and unsorted).

Key	Value
173010	100
173025	110
173045	150
173058	160
175500	70
173110	190
175515	0

Data Lakes

Data Lakes


Data Lake concept: a massive collection of datasets that may...

- be **hosted** in different storage systems
- **vary** in their formats
- **not** be accompanied by any useful **metadata** or may use different formats to describe their metadata
- **change** autonomously over time




Nargesian, F., Zhu, E., Miller, R. J., Pu, K. Q., & Arocena, P. C. (2019). Data lake management: challenges and opportunities. *Proceedings of the VLDB Endowment*, 12(12), 1986-1989.

Data Lakes



Data Lake Management System [*]



[*] Nargesian, F., Zhu, E., Miller, R. J., Pu, K. Q., & Arocena, P. C. (2019). Data lake management: challenges and opportunities. *Proceedings of the VLDB Endowment*, 12(12), 1986-1989.

Data Lakes

PROS

- **Store Everything:**
 - Store all kinds of data, no matter its structure.
 - Data added as-is (append-only) → not modified in place.
- **No Pre-Planning Required:**
 - No need for defining a schema before adding data.
 - Useful for situations where analysis to perform is not yet clear.

Data Lakes

PROS

- **File-Based Storage:**
 - Data stored as raw files, open formats (CSV, JSON, etc).
 - Files can serve as **inputs** or **intermediate outputs** for further processing.
- **Scalable and Agile:**
 - Data lakes rely on **distributed storage** to handle **large datasets**.
 - They support **distributed analytics** for processing data quickly and efficiently. **(Not a Warehouse!)**

Data Lakes

CONS “Data Swamp”

- **Low Data Quality** Without a schema data might be incomplete, incorrect, duplicated or inconsistent.
- **Missing Metadata** hard to search and understand what the data is for.
- **No Data Catalog** Without a clear catalog, it's challenging to locate specific datasets in the lake.
- **Solution:** data curation, metadata management, data catalog, governance, provenance.



Halevy et.al.
(2016, June).
Goods: Organizing
google's datasets.
In Proceedings of
the 2016
International
Conference on
Management of Data
(pp. 795-806).

Object Stores

Object Storage

Key-Value stores

- **Key-value mapping**
 - **Values** can be of a variety of data types
 - E.g. “**250**” {"sensor": “Speed_FW_Left”, “raw”: 150}
- **Scalability using Sharding:**
 - Datasets splitted into smaller chunks (shards) across multiple machines.
 - Each shard handles a subset of the **key-value** pairs, enabling the system to scale efficiently.
- **APIs for CRUD**
 - Enable entities to interact with the **key-value** store to perform these operations
(Copy-Read-Update-Delete)


Object Storage

Object Store. Similar to key-value stores, but **optimized for large objects** (videos, backups, etc.).

- **Structure:**
 - **Object Identifier (Key)**: UID to retrieve the object.
 - **Metadata** (e.g., size, type, creation date).
 - **Object (BLOB)**: The actual data, stored as a Binary Large Object. (E.g. Image at Azure Blob)
- **APIs:**
 - REST APIs: HTTP-based interfaces for CRUD
 - SDKs: Programming libraries for easier integration with applications.

Object Storage

Key Techniques



Object Storage

Examples: Amazon Simple Storage Service (S3)

- Reliable object store for photos, videos, documents or any binary data
- **Bucket:** Uniquely named, static data container
- **arn:**aws:s3:::distributed-storage
- **https://distributed-storage.s3.us-west-2.amazonaws.com/Temperature-processed.csv**
- Single (5GB)/ multi-part (5TB)
- AWS CLI Upload / Download
- Storage classes: STANDARD, STANDARD_IA, GLACIER

Distributed File Systems

Hadoop Distributed File System (HDFS)



Brief Hadoop History

- Google's GFS + MapReduce [ODSI'04] → Apache Hadoop (2006).
- Spark (HDFS + Others)


HDFS Overview

- Hadoop's distributed file system → **clusters and large datasets**
- Implemented in **Java**, w/ native libraries for **compression, I/O, CRC32**
- Files split into **128MB blocks**, replicated (**3x**), and distributed

Hadoop Distributed File System (HDFS)

How HDFS works:

- Split files into **Blocks**
- Store **blocks** across **DataNodes**
- Replicate **blocks** across **DataNodes** (reliability)
- **NameNode** coordinates this process



Hadoop Distributed File System (HDFS)

NameNode

- Coordinator daemon that manage file system **namespace (metadata)**
- **Metadata** (replication, permissions, sizes, block ids, directory structure, etc)

DataNode

- Worker daemon per **cluster node** that manages block storage (list of disks)
- Block **creation, deletion, replication as individual files** in local FS
- **On startup:** scan local blocks and send block report to **NameNode**
- Serving data block **read and write** requests
- **Send heartbeats** to **NameNode** (capacity, current transfers) and **receives replies** (replication, removal of block replicas) from to **NameNode**

HDFS InputFormats and RecordReaders

Overview HDFS API

- **InputFormats** Implements access to distributed files (data collections)
- **Split** Logical division (w/offsets) of the input data aligned with the block size of HDFS (**128 MB**).
- **RecordReader:** API for reading key-value pairs from file splits (**block creation**)
 - Keys: offsets, Values: data between offsets
- Examples: FileInputFormat, TextInputFormat, SequenceFileInputFormat


HDFS InputFormats and RecordReaders

Overview InputFormats Example

```

FileInputFormat.addInputPath(job,
path); # path: dir/file
TextInputFormat informat = new
TextInputFormat();
InputSplit[] splits =
informat.getSplits(job, numSplits);
LongWritable key = new LongWritable();
Text value = new Text();
for(InputSplit split : splits) {
RecordReader<LongWritable,Text> reader
= informat
.getRecordReader(split, job,
Reporter.NULL);
while( reader.next(key, value) )
... //process individual text lines
}
    
```


Key-Value	
Byte offset	Line content




HDFS InputFormats and RecordReaders

Sequence Files (store **key-value** pairs efficiently)

- **Binary File format** (reducing storage overhead compared to plain text files)
- **Optimal Compression** Record-level and block-level compression
- **Input & Output** in MapReduce/Spark jobs
- **Intermediate storage** during MapReduce workflows
- **Splittable**: they can be splittable allowing parallel processing in Hadoop.



HDFS Read




1. Open file
2. Metadata (Block Locations, Replicas)
3. Read Block 1 (DataNode 1)
4. Read Block 2 (Data Node 3)

HDFS Write

HDFS Client → NameNode

1. Create a new file in HDFS.
2. The **NameNode** checks whether the **file already exists** and if the client has the **necessary permissions** to create it.
3. If everything is valid, the NameNode:
 - a. Allocates a **lease** to the client, granting it the right to write to the file.
 - b. Identifies a set of **DataNodes** (+replica nodes) where the blocks of the file will be stored.




HDFS Write

Writing Blocks to DataNodes

1. The client begins writing data in **blocks** to the first DataNode in the pipeline.
2. Once the first DataNode receives a block, it forwards a copy of the block to the second DataNode in the pipeline.
3. The second DataNode then forwards the block to the third DataNode, completing the **replication pipeline**.

DataNodes Report to the NameNode

1. After receiving and storing the data blocks, each DataNode sends a **heartbeat**. It includes a report confirming that the DataNode has **successfully stored** the block and is **ready** for future tasks.



HDFS Data Locality

HDFS is generally rack-aware (node-local, rack-local, other)


Scheduler reads from closest data node

Replica placement (x3): local DN, other-rack DN, same-rack DN

Key-Value Stores, Cloud DBMS, EWarehouses)

Key-Value Stores Motivation

- **Complex data models** (e.g., JSON) can be transformed into simple **key-value** pairs.
- Designed to operate reliably at **very-large scale** using commodity hardware and distributing the workload across multiple servers.
- **Key-value maps**, where values can be of a variety of data types
- **APIs for CRUD operations** (create, read, update, delete)
- **Scalability via sharding** (horizontal partitioning)



Key-Value Stores: Example

Amazon DynamoDB Simple, highly-available data storage for small objects in ~1MB range (data, shopping carts)

- Aims to achieve Service Level Agreements (**SLAs**) that guarantee **99.9%** of requests are served with low latency, even under high loads.
- **System interface:** get and put operations
- **Partitioning** using consistent hashing where Nodes are organized in a ring structure and each node is responsible for a specific range of **keys**.
- **Replication**, Each data is **replicated N times** to provide fault tolerance.
- **Eventual consistency**

Cloud Databases (DBaaS): Motivation and Key-aspects

DBaaS: A cloud service model that allows users to access, manage, and scale databases without dealing with the underlying infrastructure.

Providers handle database configuration, maintenance, security, updates, and availability.



Key aspects

- SQL and NoSQL (MongoDB, DynamoDB).
- Auto-scaling, high availability, and support for multiple data types.
- Mainly used for OLTP (web apps, IoT, etc.)

Examples: Amazon RDS, Google Cloud SQL, Azure SQL Database

Elastic Data Warehouses

- Data warehousing + Cloud Computing + Distributed Storage
- Used for OLAP



Summary and Q&A

Summary and Q&A

- Motivation and Terminology
- Data Lakes
- Object Stores (S3)
- Distributed File Systems (HDFS)
- Key-Value Stores, Cloud DBMS, Elastic Data Warehouses
- Next Lecture: Distributed, Data-Parallel Computation **[Dec 19]**

Vielen Dank!