

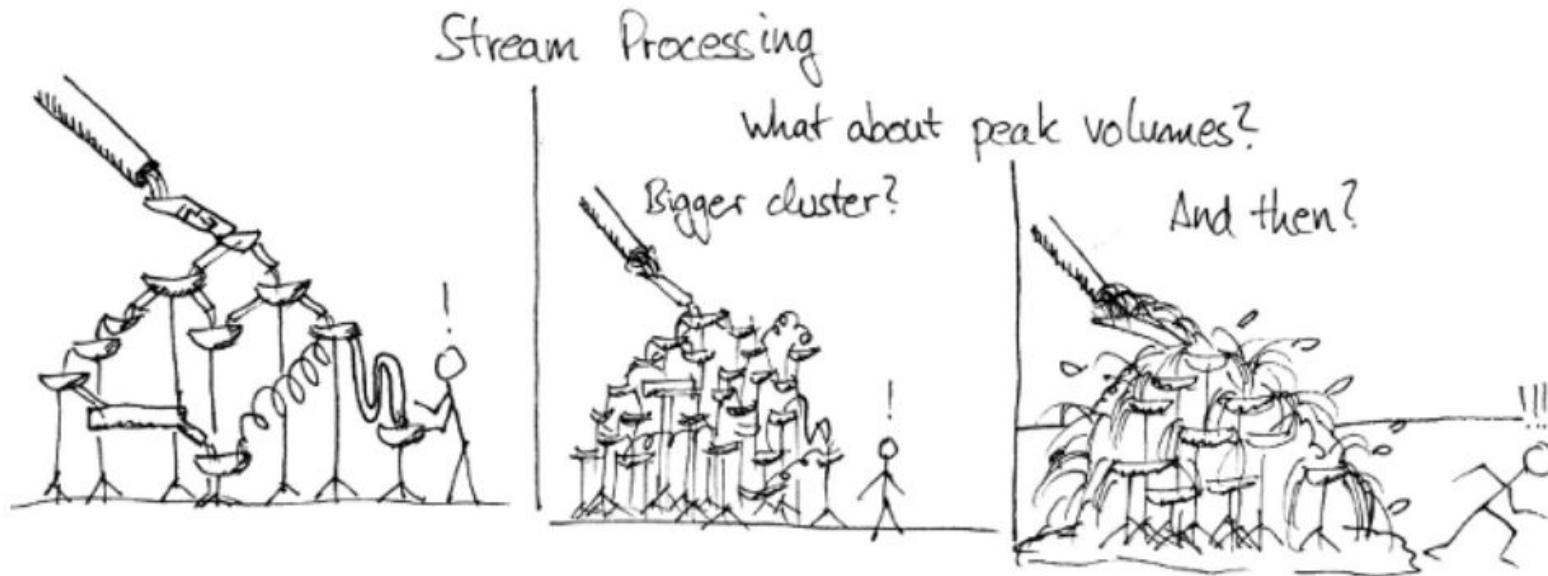
# Data Integration and Large Scale Analysis

## 11- Stream Processing

Lucas Iacono. PhD. - 2026

# Data Stream Processing

# Data Stream Processing



# Data Stream Processing: Terminology

## Ubiquitous Data Streams

**Event and message streams** (e.g., clicks, X, IoT, Machines)

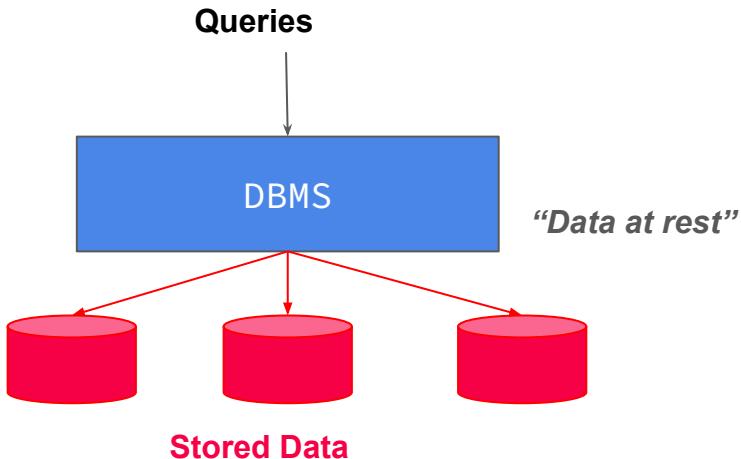


Characteristics	
Event Streams	Message Streams
<ul style="list-style-type: none"><li>• Time-oriented.</li><li>• Immutable</li><li>• Real-time processing</li></ul>	<ul style="list-style-type: none"><li>• System communication</li><li>• Asynchronous</li><li>• Structured</li></ul>

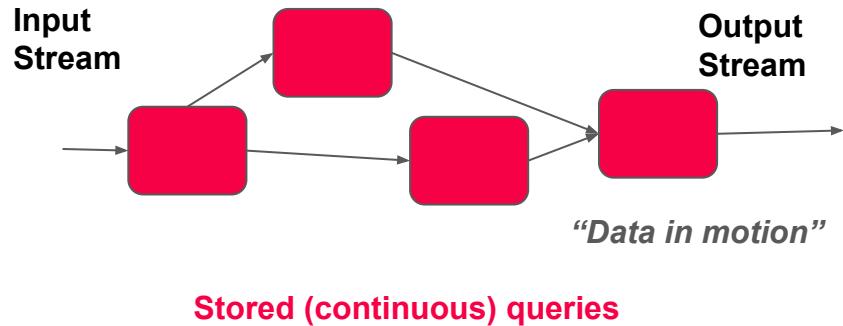
# Data Stream Processing: Terminology

## Stream Processing Architecture

- **Infinite input streams**, often with windows semantics
- Continuous queries



## Stream Processing Engines



# Data Stream Processing: Terminology

## Use Cases

- **Monitoring and alerting** (notifications on events / patterns)
- **Real-time reporting** (aggregate statistics for dashboards)
- **Real-time ETL** and event-driven data updates
- Real-time decision making (fraud detection)
- **Data stream mining** (summary **statistics** w/ limited memory)

## Data Stream

- Unbounded stream of data tuples  $S = (s_1, s_2, \dots, s_n)$ . Each data tuple  $(s_i) \rightarrow s_i = (t_i, d_i)$
- $t_i$  = timestamp,  $d_i$  = data
- $S = \{(10:30:00 09082025, 22.5), (10:30:00 09082025, 22.7), (10:30:00 09082025, 22.8), \dots\}$

# Data Stream Processing: Terminology

## Real-time Latency Requirements

- **Real-time:** guaranteed task completion by a given deadline (30 fps)
- **Near Real-time:** few milliseconds to seconds
- In practice → used with much weaker meaning

## Challenges in Real-Time Systems:

- Resource Constraints (memory, computing, storage)
- Latency (e.g. data in a queue waiting for been processed)

# Data Stream Processing: History of Stream Processing Systems

2000s

- **Data stream management systems:** **STREAM** (Stanford'01), **Aurora** (Brown/MIT/Brandeis'02), **TelegraphCQ** (Berkeley'03) → mostly unsuccessful in industry/practice
- **Message-oriented middleware** and **Enterprise Application Integration (EAI)**: IBM Message Broker, SAP eXchange Infra

# Data Stream Processing: History of Stream Processing Systems

2010s

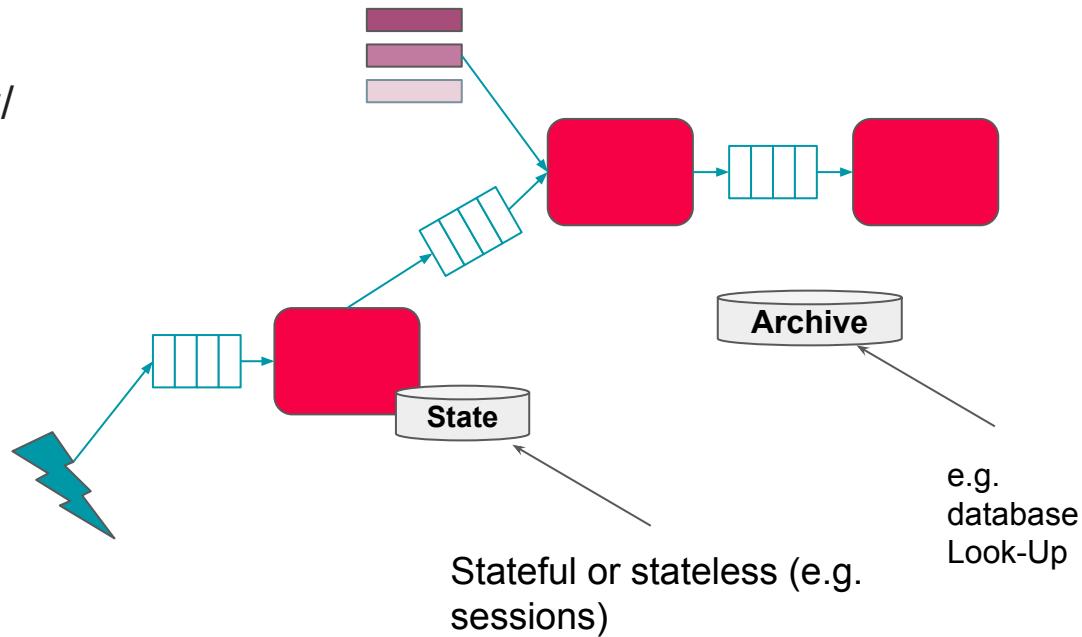
- **Distributed stream processing engines**, and “unified” batch/stream processing
- **Proprietary systems**: Google Cloud Dataflow, MS StreamInsight / Azure Stream Analytics, IBM InfoSphere Streams / Streaming Analytics, AWS Kinesis
- **Open-source systems**: Apache Spark Streaming (Databricks), **Apache Flink** (Data Artisans), **Apache Kafka** (Confluent), **Apache Storm**



# System Architecture - Native Streaming

## Basic System Architecture

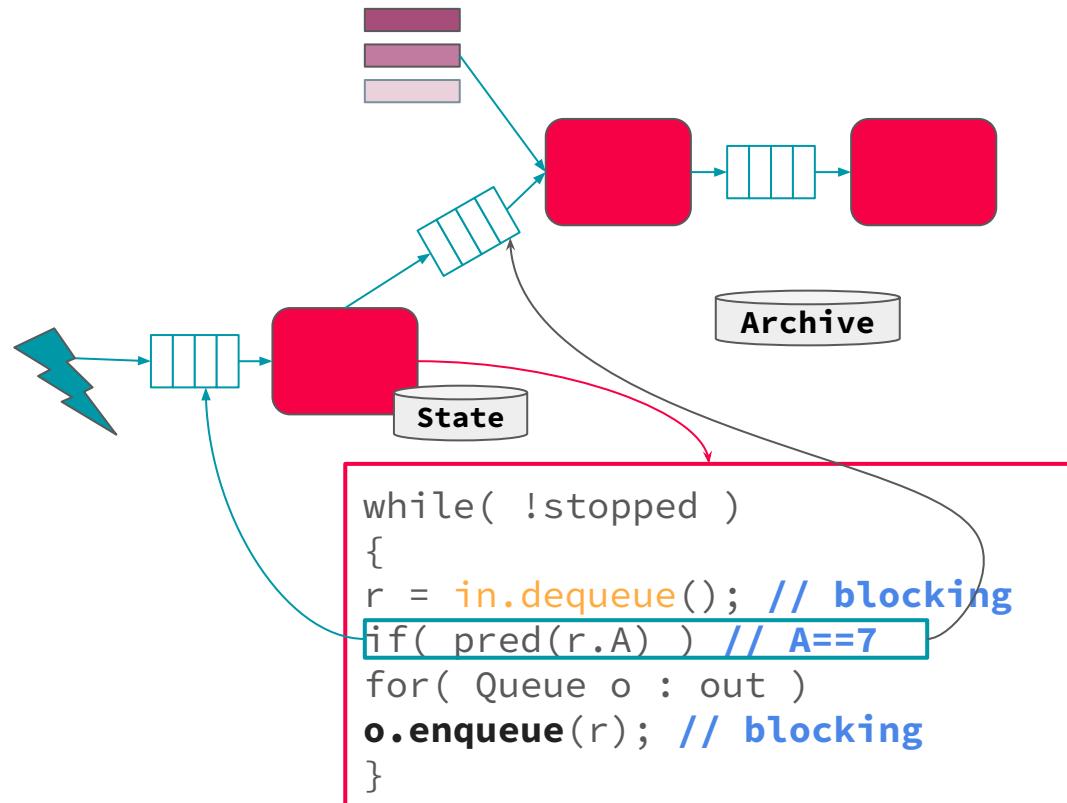
- Data flow graphs (potentially w/ multiple consumers)
- **Nodes** asynchronous ops (w/ state) (e.g., separate threads)
- **Edges** data dependencies (tuple/message streams)
- **Push model** data production controlled by source



# System Architecture - Native Streaming

## Operator Model

- Read from input queue
- Write to potentially many output queues



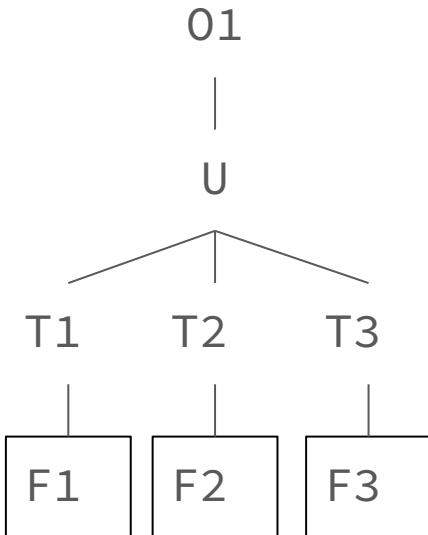
# System Architecture - Sharing

## Multi-Query Optimization

- Given **set of continuous queries** compile minimal DAG w/o redundancy -> **subexpression elimination** -> **avoid redundant operations and share intermediate results between queries to improve system efficiency.**

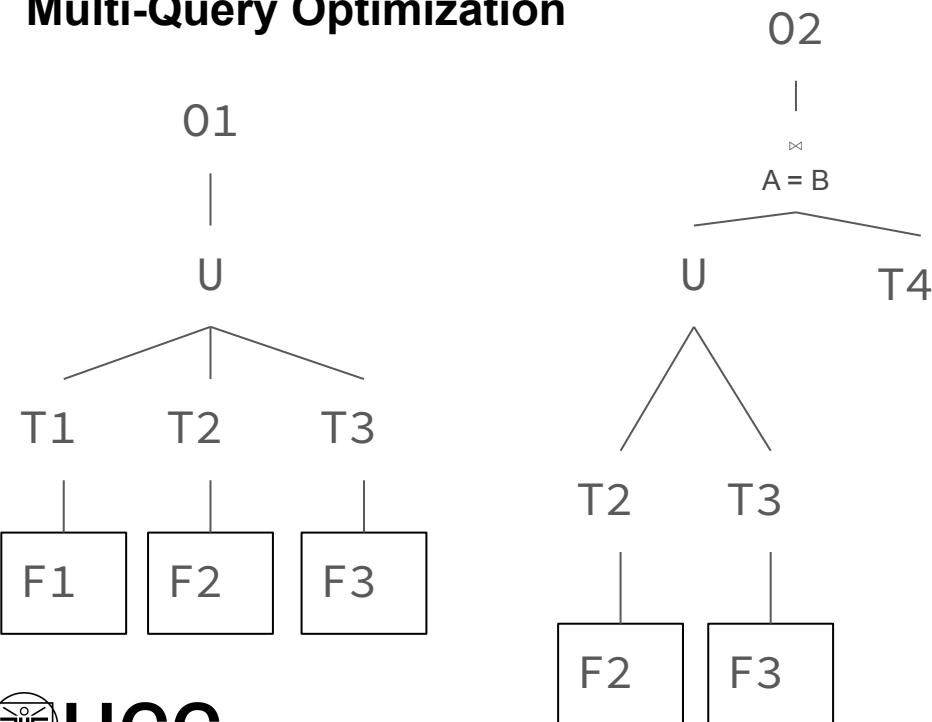
# System Architecture - Sharing

## Multi-Query Optimization



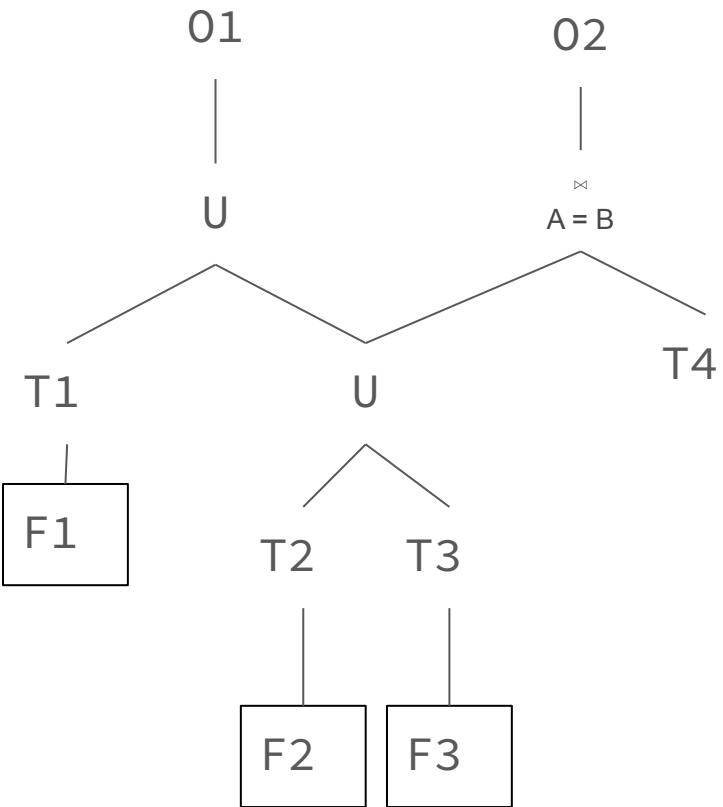
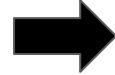
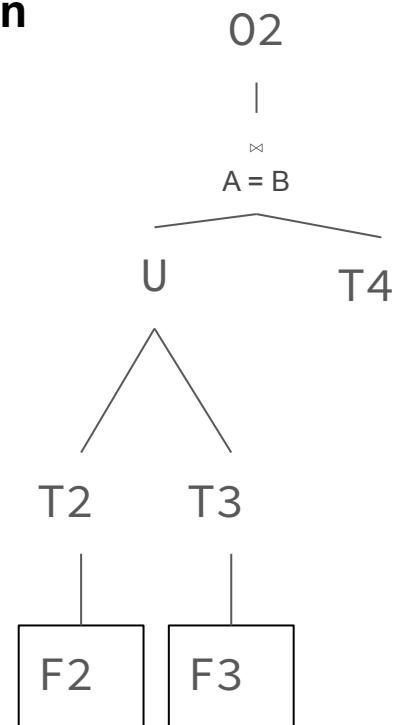
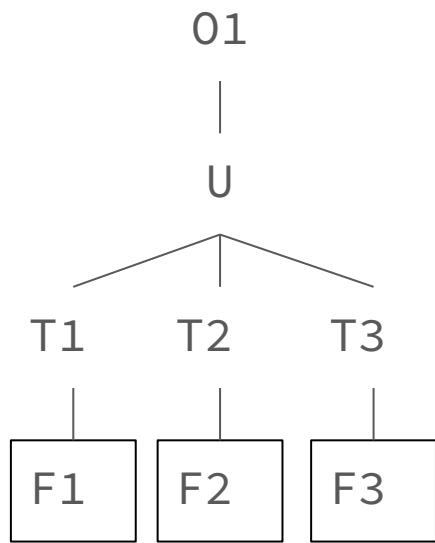
# System Architecture - Sharing

## Multi-Query Optimization



# System Architecture - Sharing

## Multi-Query Optimization

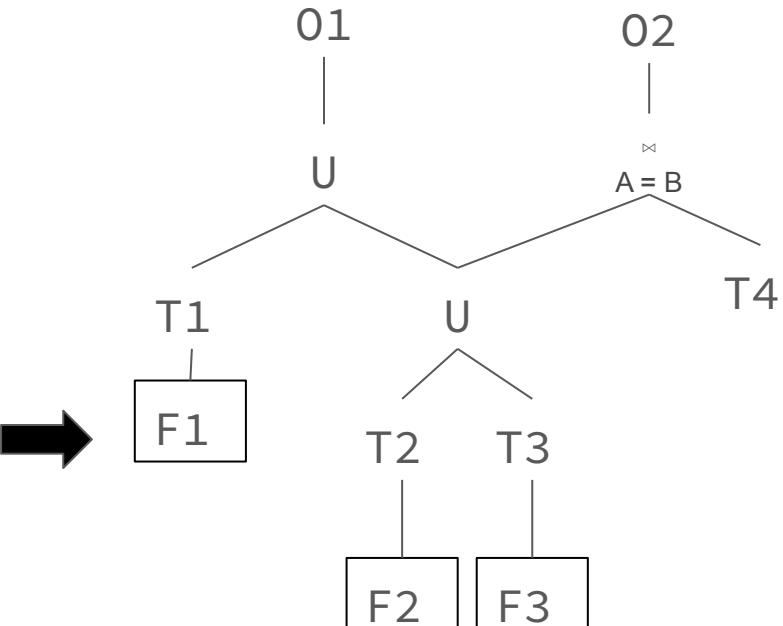
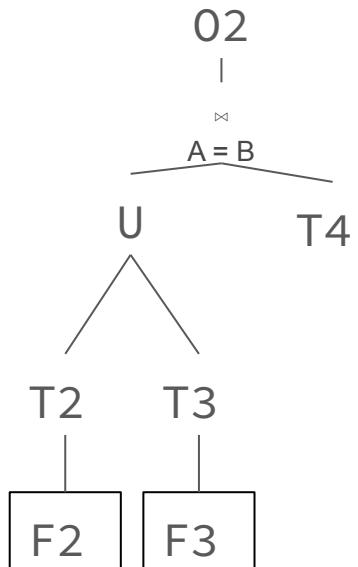
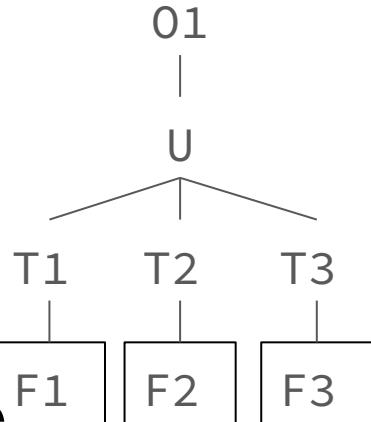


# System Architecture - Sharing

Operator Sharing: complex ops w/ multiple predicates for adaptive reordering

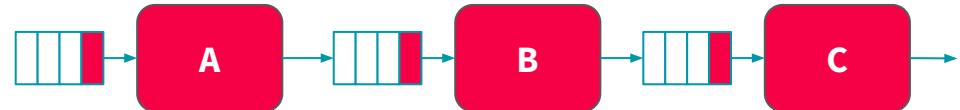
## Multi-Query Optimization

Queue Sharing: share results with multiple queries



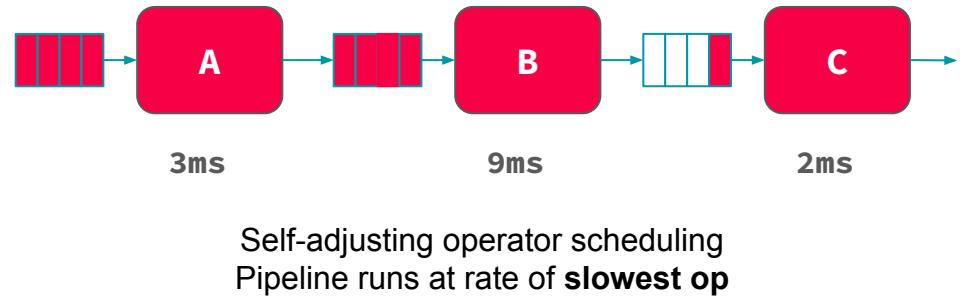
# System Architecture - Handling Overload

- **Back Pressure**
  - Graceful handling of overload w/o data loss
  - **Slow down sources**
  - E.g. blocking queues



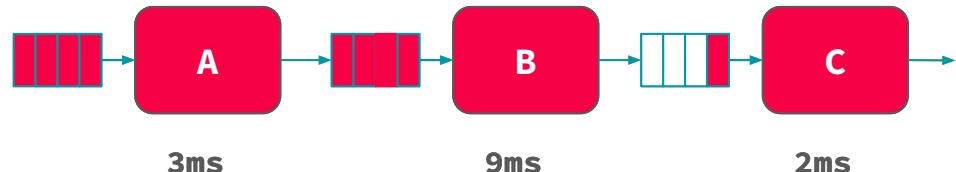
# System Architecture - Handling Overload

- **Back Pressure**
  - Graceful handling of overload w/o data loss
- **Slow down sources**
- E.g. blocking queues



# System Architecture - Handling Overload

- **Back Pressure**
  - Graceful handling of overload w/o data loss
  - **Slow down sources**
  - E.g. blocking queues
- **Load Shedding**
  - **Random-sampling-based** load shedding
  - **Relevance-based** load shedding



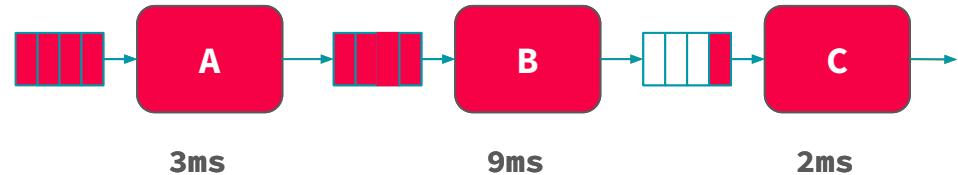
Self-adjusting operator scheduling  
Pipeline runs at rate of **slowest op**



[Nesime Tatbul et al: Load Shedding in a Data Stream Manager. VLDB 2003]

# System Architecture - Handling Overload

- **Back Pressure**
  - Graceful handling of overload w/o data loss
  - **Slow down sources**
  - E.g. blocking queues
- **Load Shedding**
  - **Random-sampling-based** load shedding
  - **Relevance-based** load shedding
  - **Summary-based** load shedding (synopses)
  - Given SLA, select queries and shedding placement that minimize error and satisfy constraints
- **Distributed Stream Processing**
  - Data flow partitioning (distribute the **query**)
  - Key range partitioning (distribute the **data stream**)



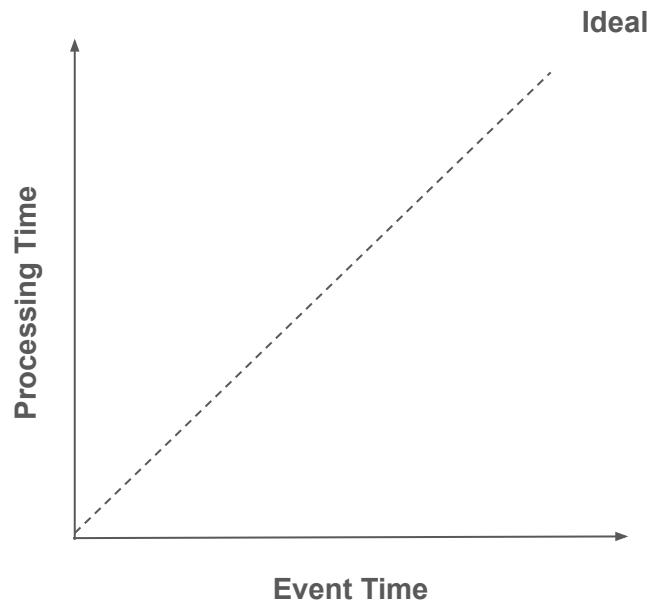
Self-adjusting operator scheduling  
Pipeline runs at rate of **slowest op**



[Nesime Tatbul et al: Load Shedding in a Data Stream Manager. VLDB 2003]

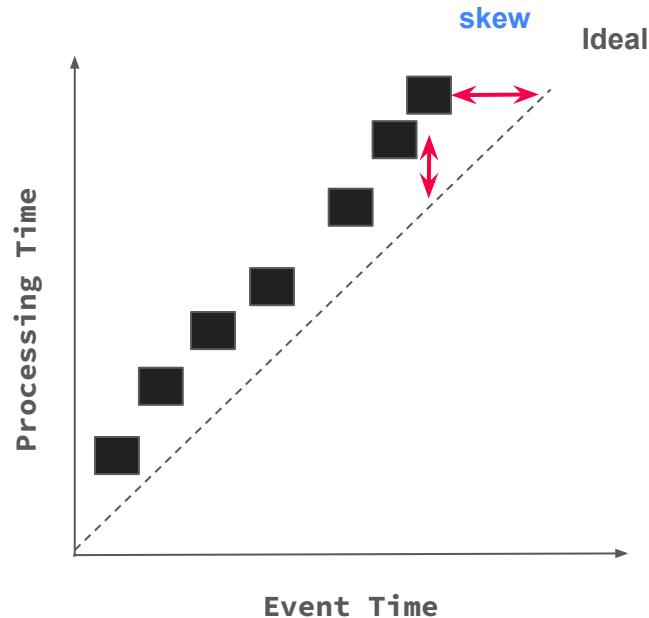
# Time (Event, System, Processing)

- **Event Time**
  - Real time when the **event/data item was created**
- **Ingestion**
  - System time when the **data item was received**
- **Processing Time**
  - System time when the **data item is processed**



# Time (Event, System, Processing)

- **Event Time**
  - Real time when the **event/data item was created**
- **Ingestion**
  - System time when the **data item was received**
- **Processing Time**
  - System time when the **data item is processed**
- **In practice**
  - **Delayed and unordered** data items
  - Use of heuristics (e.g **watermarks = delays threshold**)
  - Use of more complex triggers (late results)



# Durability and Consistency Guarantees

## 03 Message-oriented Middleware, EAI, and Replication

- At Most Once
  - “Send and forget”, ensure data is never counted twice
  - Might cause data loss on failures

# Durability and Consistency Guarantees

- **At Most Once**
  - “**Send and forget**”, ensure data is never counted twice
  - Might cause data loss on failures
- **At Least Once**
  - “**Store and forward**” deliver the message until reception of the acknowledgements from receiver
  - Might create incorrect state (processed multiple times)

# Durability and Consistency Guarantees

- **At Most Once**
  - “**Send and forget**”, ensure data is never counted twice
  - Might cause data loss on failures
- **At Least Once**
  - “**Store and forward**” deliver the message until reception of the acknowledgements from receiver
  - Might create incorrect state (processed multiple times)
- **Exactly Once**
  - “**Store and forward**” w/ **guarantees** regarding state updates and sent msgs
  - Often via dedicated transaction mechanisms (**hand-shaking protocols**)

# Window Semantics

- **Windowing Approach**
  - Many operations like joins/aggregation **undefined over unbounded streams**
  - Compute operations over windows of **(a) time** or **(b) elements counts**

# Window Semantics

- **Windowing Approach**
  - Many operations like joins/aggregation **undefined over unbounded streams**
  - Compute operations over windows of **(a) time** or **(b) elements counts**
  
- **Tumbling Window**
  - Every data item is only part of a single window
  - **Aka Jumping window**

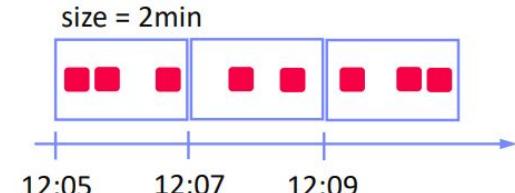


# Window Semantics

- **Windowing Approach**
  - Many operations like joins/aggregation **undefined over unbounded streams**
  - Compute operations over windows of (a) time or (b) elements counts

- **Tumbling Window**
  - Every data item is only part of a single window
  - **Aka Jumping window**

- **Sliding Window**
  - Time- or tuple-based sliding windows
  - Insert new and expire old data items



# Stream Joins I

## Basic Stream Join

- **Tumbling window:** use classic join methods
- **Sliding window** (symmetric for both R and S) → Applies to arbitrary join pred

# Stream Joins I

## Basic Stream Join

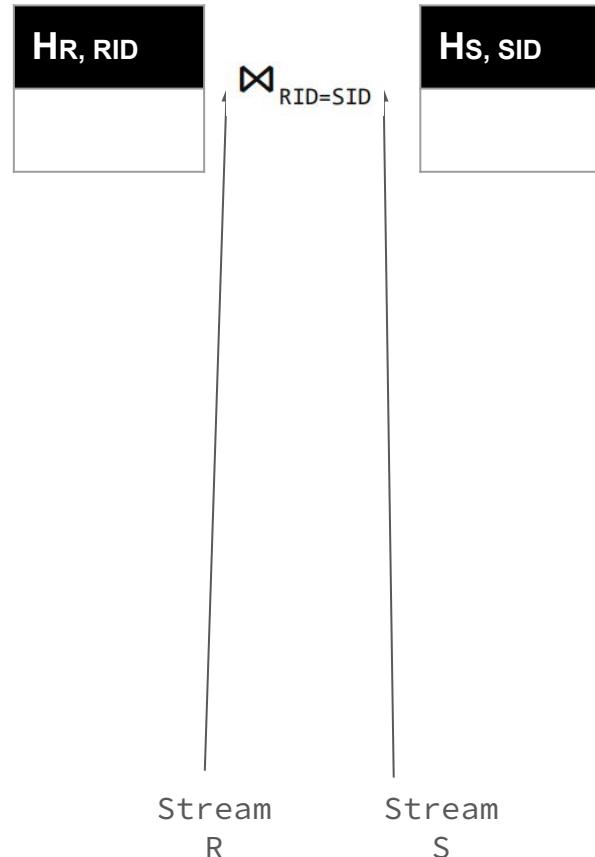
- **Tumbling window:** use classic join methods
- **Sliding window** (symmetric for both R and S) → Applies to arbitrary join pred

For each new  $r$  in R:

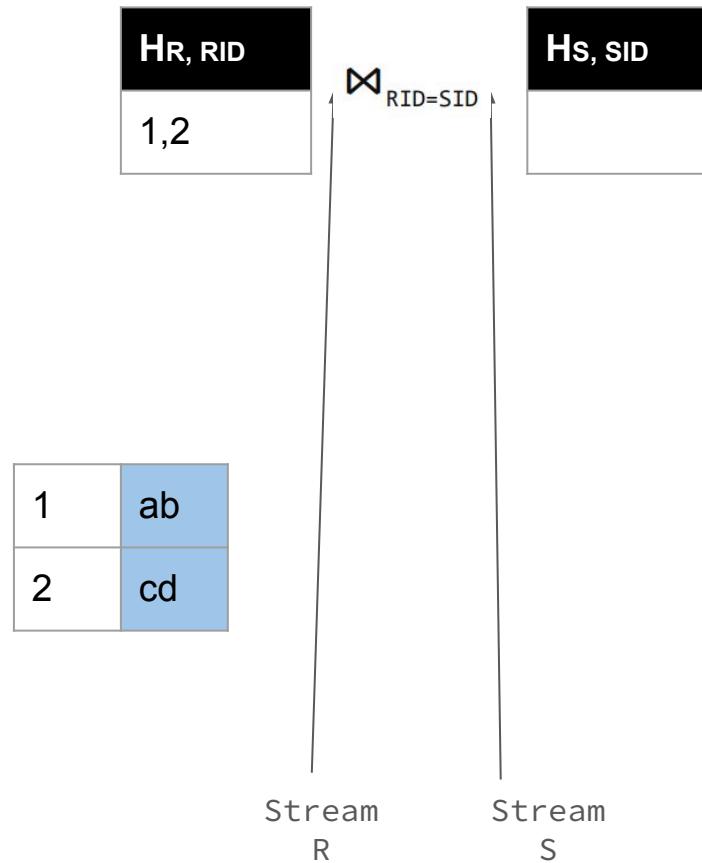
- a. **Scan** window of stream S to find match tuples
- b. **Insert** new  $r$  into window of stream R
- c. **Invalidate** expired tuples in window of stream R

# Stream Joins II

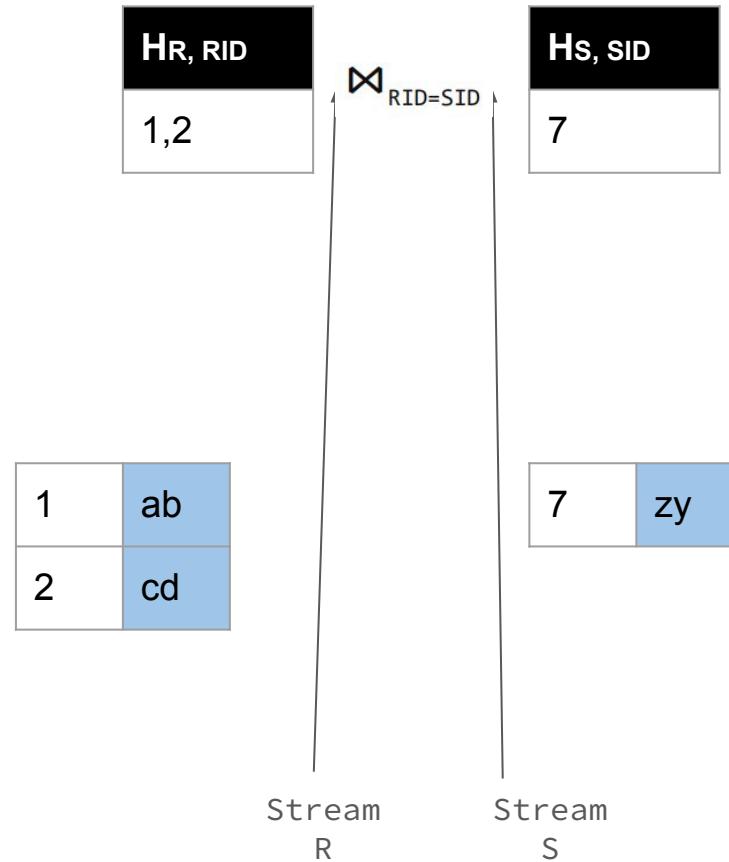
- **Double-Pipelined Hash Join**
  - Join of bounded streams (or unbounded w/ invalidation)
  - Equijoin predicate, symmetric and non-blocking
  - For every incoming tuple (e.g. left): probe (right)+emit, and build (left)



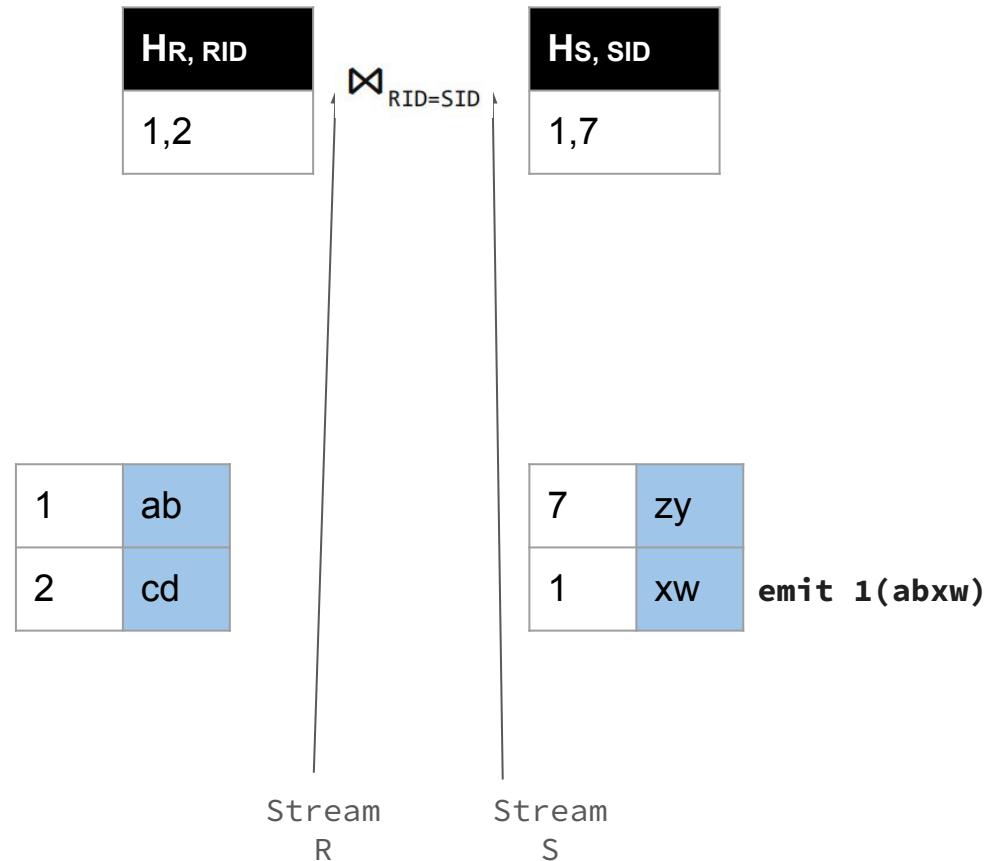
# Stream Joins II



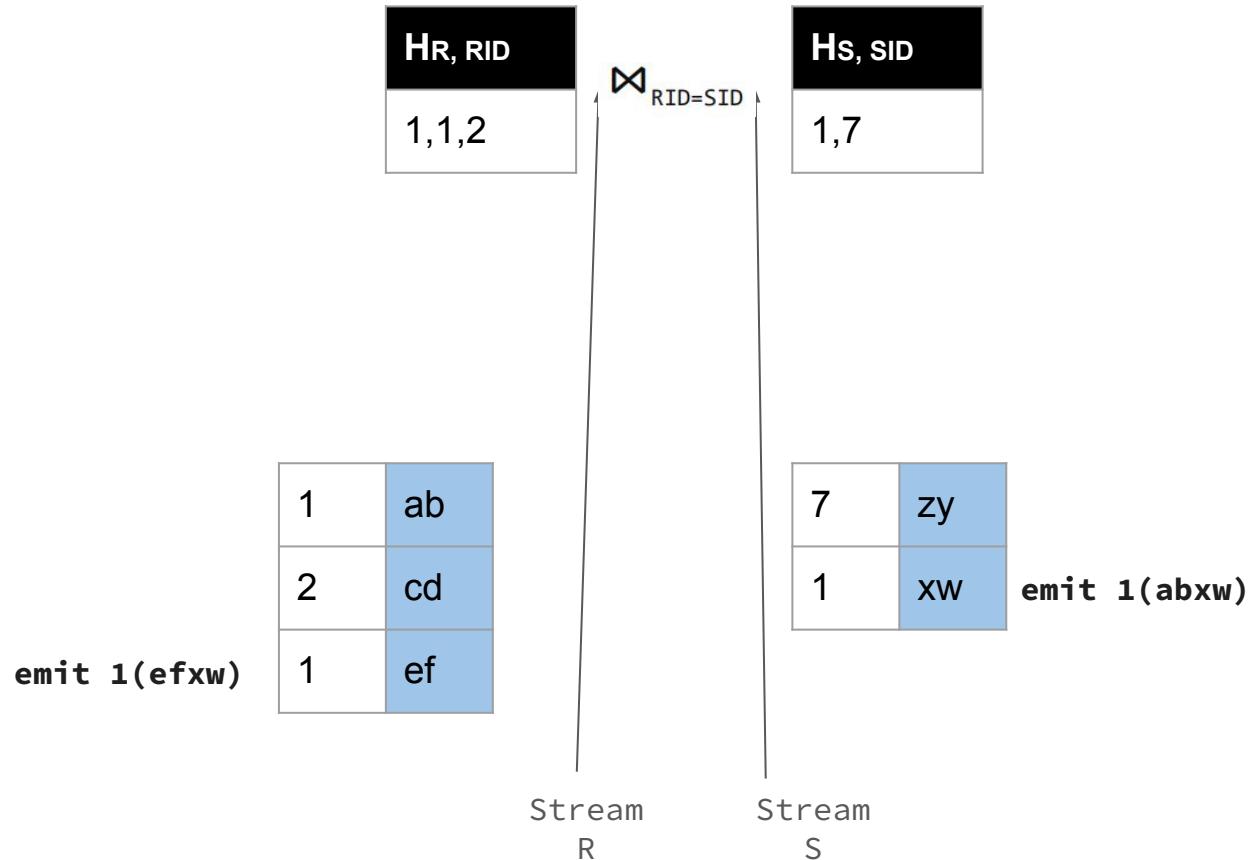
# Stream Joins II



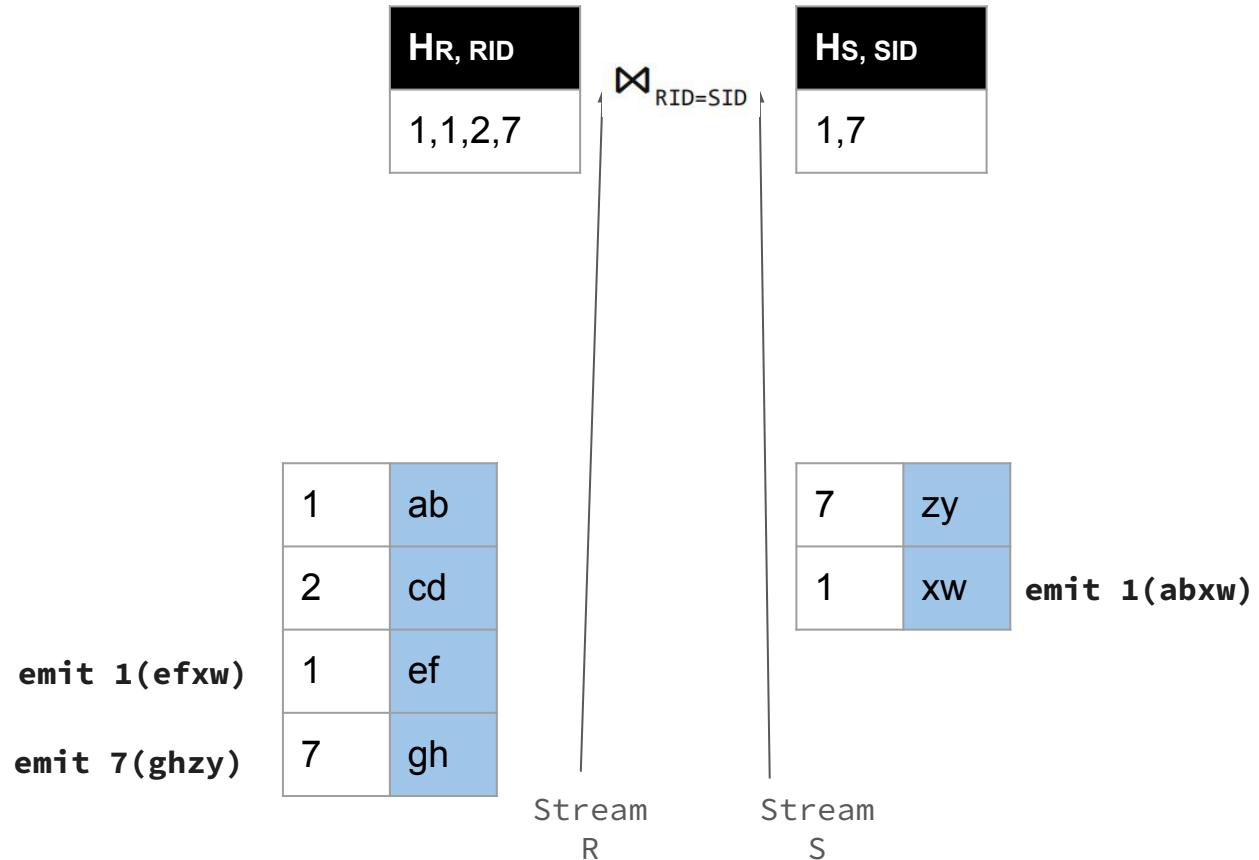
# Stream Joins II



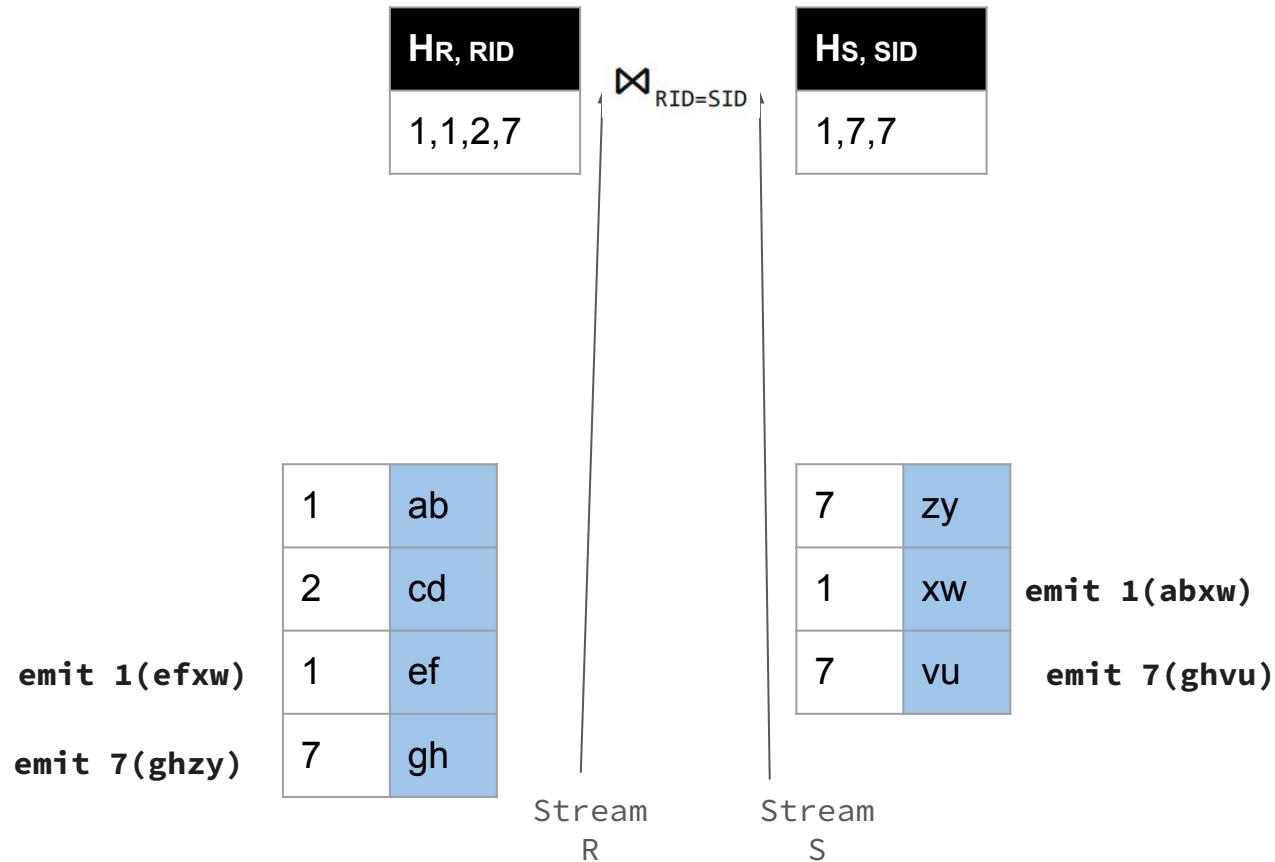
# Stream Joins II



# Stream Joins II



# Stream Joins II



# Distributed Stream Processing

# Query-Aware Stream Partitioning

## Example Use Case

- AT&T network monitoring
- **112M packets/s → 26 cycles/tuple** on **3Ghz CPU**
- Complex query sets (apps w/ **~50 queries**) and massive data rates



T. Johnson et.al,  
Query-aware  
partitioning for  
monitoring  
massive network data  
streams. **SIGMOD 2008**

# Query-Aware Stream Partitioning

T. Johnson et.al, Query-aware partitioning for monitoring massive network data streams. SIGMOD 2008



## Baseline Query Execution Plan

Self join	$\bowtie_{tb=tb+1}$	<b>Query FLOW PAIRS:</b> users generating heavy-loads during longer periods (no peaks)
High-level aggregation	$\gamma_2$	<b>Query HEAVY FLOWS:</b> heaviest flows for each IP
Low-level aggregation	$\gamma_1$	<b>Query FLOWS:</b> how many requests are generated per each IP in the network (per minute)
Low-level filtering	$\sigma$	
	TCP	

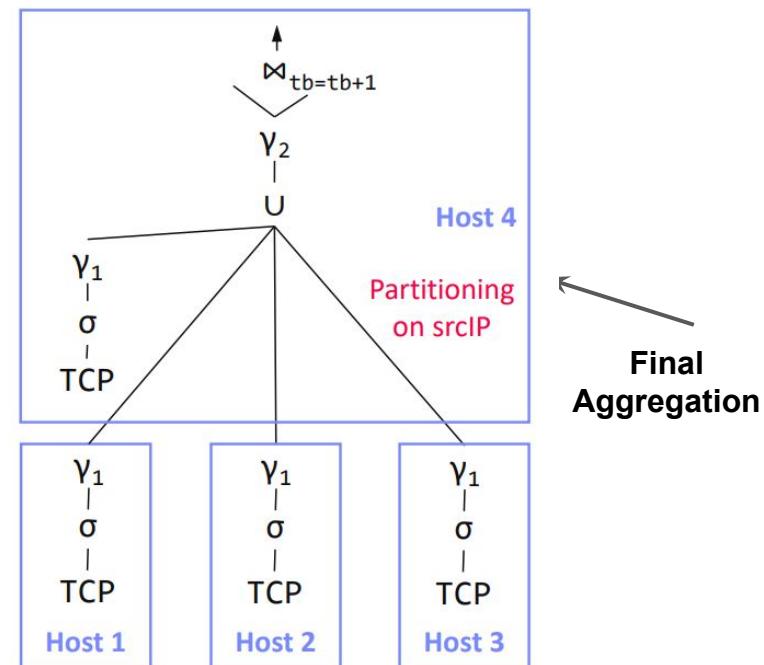
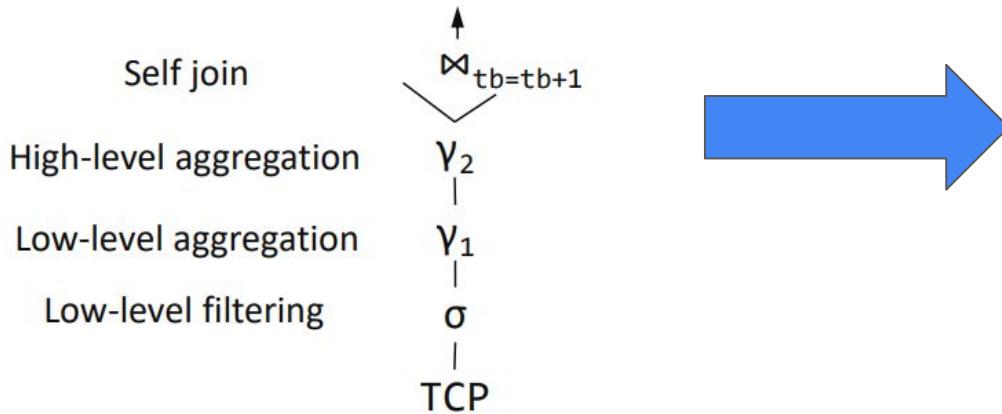
# Query-Aware Stream Partitioning

T. Johnson et.al, Query-aware partitioning for monitoring massive network data streams. SIGMOD 2008

**Solution** divide in sub-queries and distribute

**Optimized Plan:**

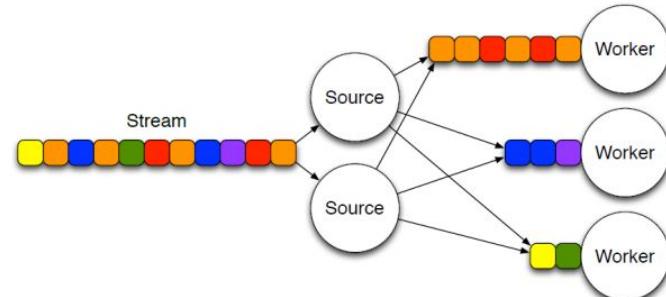
- Distributed Plan operators
- Pipeline and task parallelism
- **Not always enough**



# Stream Group Partitioning

## Large-Scale Stream Processing

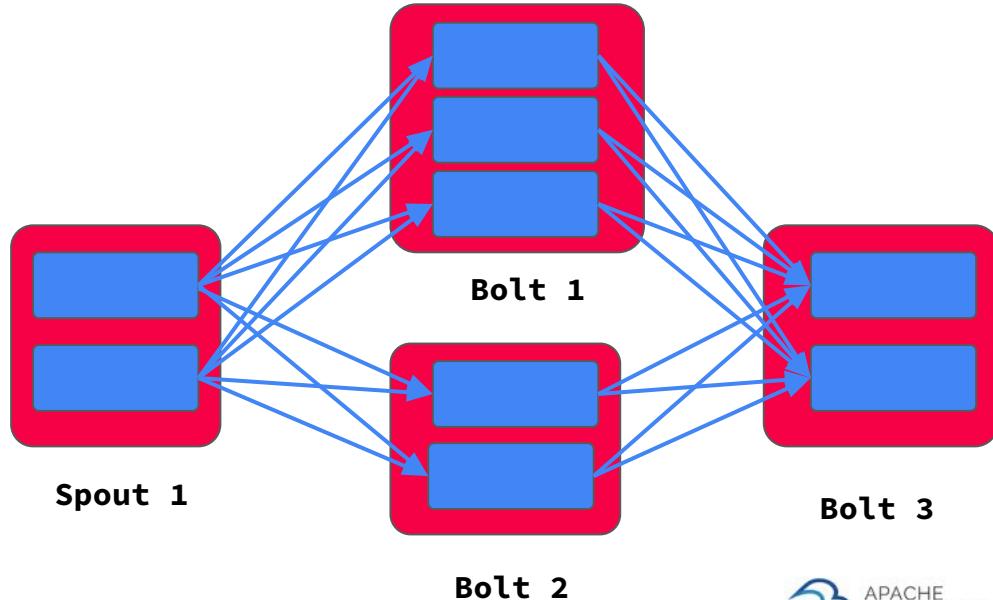
- Limited pipeline parallelism and task parallelism (independent subqueries)
- **Combine with data-parallelism over stream groups**
  - **Shuffle Grouping**
    - Tuples are randomly distributed across consumer tasks
    - Good load balance
  - **Fields Grouping**
    - Tuples partitioned by grouping attributes (keys)
    - Guarantees order within keys, but load imbalance if skew



# Example Apache Storm

- **Example Topology DAG**

- **Spouts:** sources of streams
- **Bolts:** UDF compute ops
- Tasks mapped to worker processes and executors (threads)



```

Config conf = new Config();
conf.setNumWorkers(3);
topBuilder.setSpout("Spout1", new FooS1(), 2);
topBuilder.setBolt("Bolt1", new FooB1(), 3).shuffleGrouping("Spout1");
topBuilder.setBolt("Bolt2", new FooB2(), 2).shuffleGrouping("Spout1");
topBuilder.setBolt("Bolt3", new FooB3(), 2)
    .shuffleGrouping("Bolt1").shuffleGrouping("Bolt2");
StormSubmitter.submitTopology(..., topBuilder.createTopology());
  
```

# Example Twitter Heron

- Motivation
  - Heavy use of Apache Storm at Twitter
  - Issues: debugging, performance, shared cluster resources, back pressure mechanism

- Twitter Heron

- API-compatible distributed streaming engine
- De-facto streaming engine at Twitter since 2014

- Dhalion (Heron Extension)

- Automatically reconfigure Heron topologies to meet throughput SLO

- Now back pressure implemented in Apache Storm 2.0 (May 2019)

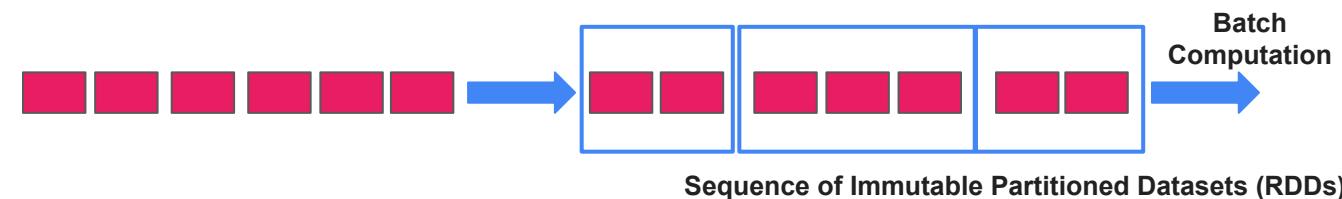


Sanjeev Kulkarni et al:  
Twitter Heron:  
Stream  
Processing at  
Scale.  
SIGMOD 2015

# Discretized Stream (Batch) Computation

- **Motivation**
  - **Fault tolerance** (low overhead, fast recovery)
  - Combination w/ **distributed batch analytics**
- **Discretized Streams (DStream)**
  - Batching of input tuples (100ms – 1s) based on ingest time.
  - Periodically run distributed jobs of stateless, deterministic tasks → **DStreams**
  - State of all tasks materialized as RDDs, recovery via lineage
- **Criticism: High latency, required for batching**

Matei Zaharia et al:  
 Discretized streams:  
 fault-tolerant  
 streaming  
 computation at  
 scale. SOSP 2013



# Unified Batch/Streaming Engines

- **Apache Spark Streaming (Databricks)**
  - Micro-batch computation with exactly-once guarantee
  - Back-pressure and water mark mechanisms
  - Structured streaming via SQL (2.0), continuous streaming (2.3)
- **Apache Flink (Data Artisans, now Alibaba)**
  - Tuple-at-a-time with exactly-once guarantee
  - Back-pressure and water mark mechanisms
  - Batch processing viewed as special case of streaming



# Summary and Q&A

- **Summary and Q&A**

- Data Stream Processing
- Distributed Stream Processing

- **Next Lectures**

- Distributed Machine Learning **[Jan 16]**
- **Written Exam [Jan 30]**

# Vielen Dank!