

Math Preliminaries & MLE

DEEP LEARNING KU (DAT.C302UF)

Simon Hitzinger

Oct 22, 2025

Institute of Machine Learning and Neural Computing
Graz University of Technology, Austria

- Strengthen knowledge about mathematical foundations of deep learning
- Become familiar with PyTorch
- Acquire a solid understanding of modern deep learning models
 - Architectures, Training Paradigms, Model Selection, Evaluation
- Being able to apply deep learning to practical problems

Math Preliminaries

Linear Algebra

- **Extremely important** from a computational perspective
- We'll use PyTorch to write "Linear Algebra Programs"
 - Core object: **Tensors** (same in Tensorflow and JAX)

Linear Algebra

- **Extremely important** from a computational perspective
- We'll use PyTorch to write "Linear Algebra Programs"
 - Core object: **Tensors** (same in Tensorflow and JAX)

Calculus

- Minimizing continuous, differentiable **loss functions**
- Will often be **high-dimensional** → Multivariate Calculus (Matrix Calculus)

Linear Algebra

- **Extremely important** from a computational perspective
- We'll use PyTorch to write "Linear Algebra Programs"
 - Core object: **Tensors** (same in Tensorflow and JAX)

Calculus

- Minimizing continuous, differentiable **loss functions**
- Will often be **high-dimensional** → Multivariate Calculus (Matrix Calculus)

Probability Theory

- **"Optimal Reasoning under Uncertainty"**
 - Noisy data, uncertainty in model parameters etc.
- Generative Modeling, Statistical Learning Theory, ...

Linear Algebra

Dot Product

- Given vectors

$$\mathbf{x} = (x_1, x_2, \dots, x_N)^T \in \mathbb{R}^N, \quad \mathbf{y} = (y_1, y_2, \dots, y_N)^T \in \mathbb{R}^N$$

- We define the **dot product** as

$$\mathbf{x}^T \mathbf{y} = \mathbf{y}^T \mathbf{x} = \sum_{i=1}^N x_i y_i$$

- PyTorch: `x.T @ y`

Important Norms

For $\mathbf{x} = (x_1, \dots, x_N)^T \in \mathbb{R}^N$, we define

- The ℓ_2 norm

$$\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^N x_i^2} = \sqrt{\mathbf{x}^T \mathbf{x}}.$$

- Thus, $\|\mathbf{x}\|_2^2 = \sum_{i=1}^N x_i^2 = \mathbf{x}^T \mathbf{x}$

- The ℓ_1 norm $\|\mathbf{x}\|_1 = \sum_{i=1}^N |x_i|$

Norms

Important Norms

For $\mathbf{x} = (x_1, \dots, x_N)^T \in \mathbb{R}^N$, we define

- The ℓ_2 norm
$$\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^N x_i^2} = \sqrt{\mathbf{x}^T \mathbf{x}}.$$
 - Thus, $\|\mathbf{x}\|_2^2 = \sum_{i=1}^N x_i^2 = \mathbf{x}^T \mathbf{x}$
- The ℓ_1 norm $\|\mathbf{x}\|_1 = \sum_{i=1}^N |x_i|$

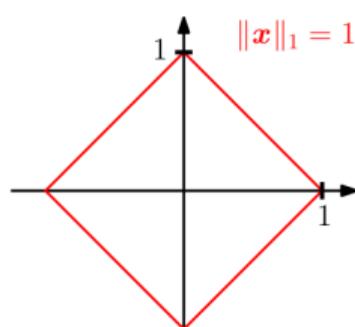
PyTorch

- `torch.linalg.norm(x, ord=p)`
- $p = 1$ or $p = 2$ in this case

Important Norms

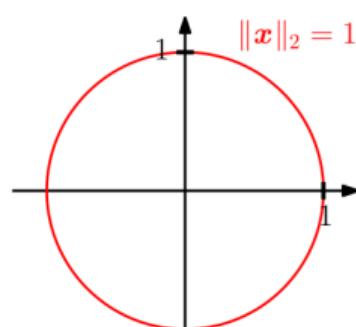
For $\mathbf{x} = (x_1, \dots, x_N)^T \in \mathbb{R}^N$, we define

- The ℓ_2 norm
$$\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^N x_i^2} = \sqrt{\mathbf{x}^T \mathbf{x}}$$
 - Thus, $\|\mathbf{x}\|_2^2 = \sum_{i=1}^N x_i^2 = \mathbf{x}^T \mathbf{x}$
- The ℓ_1 norm $\|\mathbf{x}\|_1 = \sum_{i=1}^N |x_i|$



PyTorch

- `torch.linalg.norm(x, ord=p)`
- $p = 1$ or $p = 2$ in this case



- Matrix $A \in \mathbb{R}^{M \times N}$
 - Can be seen as **collection** of N column vectors $\in \mathbb{R}^M$

Matrix-vector product

$$\mathbf{y} = A\mathbf{x}, \quad \mathbf{x} \in \mathbb{R}^N, \mathbf{y} \in \mathbb{R}^M$$

- PyTorch: `y = A @ x`
- A is a **linear map** that transforms $\mathbf{x} \in \mathbb{R}^N$ into $\mathbf{y} \in \mathbb{R}^M$

Differential Calculus

- Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a **univariate function**
- If it exists, **the derivative of f w.r.t. x** is defined as

$$\frac{df}{dx}(x) = f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

- Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a **univariate function**
- If it exists, **the derivative of f w.r.t. x** is defined as

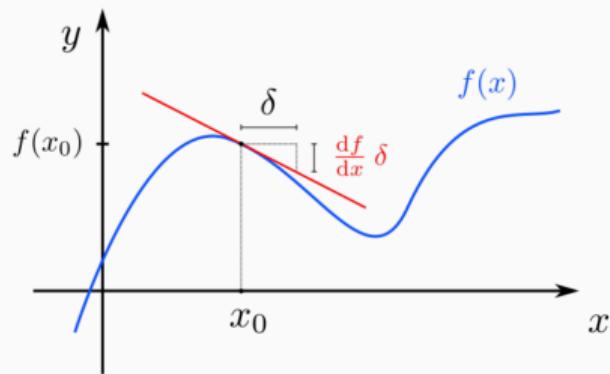
$$\frac{df}{dx}(x) = f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

Intuition

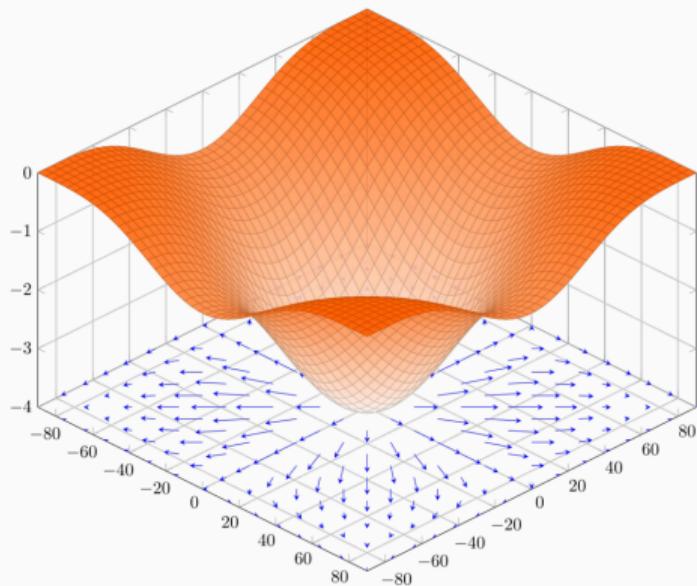
- I know $f(x)$ for some x
- I **slightly change the input** (from x to e.g. $x + 0.00001$)
- How will the **output** $f(x + 0.00001)$ change (in proportion to the change in input)?
- i.e., **how sensitive** is f to changes to its inputs (at particular points)?

$$\frac{df}{dx}(x) = f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

- “**Instantaneous**” rate of change
- $f'(x)$ is the **slope of the tangent line** going through x
- This is the best **local linear approximation**



- Let $f(x_1, x_2, \dots, x_D)$ be a **multi-variate scalar-valued** function ($f : \mathbb{R}^D \rightarrow \mathbb{R}$)
 - The **gradient** of f w.r.t. $\mathbf{x} = (x_1, \dots, x_D)^T$ at some point \mathbf{x} is defined as
- $$\nabla f(\mathbf{x}) = \begin{pmatrix} \frac{\partial f(\mathbf{x})}{\partial x_1} \\ \frac{\partial f(\mathbf{x})}{\partial x_2} \\ \vdots \\ \frac{\partial f(\mathbf{x})}{\partial x_D} \end{pmatrix}$$
- $\nabla f(\mathbf{x})$ points in the direction of steepest ascent



<https://commons.wikimedia.org/wiki/File:3d-gradient-cos.svg>

- Let $f(\mathbf{x})$ be a **vector-valued** function ($f : \mathbb{R}^D \rightarrow \mathbb{R}^M$), defined as

$$f(\mathbf{x}) = (f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_M(\mathbf{x}))^T$$

where $f_i(\mathbf{x}) : \mathbb{R}^D \rightarrow \mathbb{R}$.

- The **Jacobian matrix** of f at some point \mathbf{x} is

$$J_f(\mathbf{x}) = \begin{pmatrix} \nabla f_1(\mathbf{x})^T \\ \nabla f_2(\mathbf{x})^T \\ \vdots \\ \nabla f_M(\mathbf{x})^T \end{pmatrix} = \begin{pmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \frac{\partial f_1(\mathbf{x})}{\partial x_2} & \cdots & \frac{\partial f_1(\mathbf{x})}{\partial x_D} \\ \frac{\partial f_2(\mathbf{x})}{\partial x_1} & \frac{\partial f_2(\mathbf{x})}{\partial x_2} & \cdots & \frac{\partial f_2(\mathbf{x})}{\partial x_D} \\ \vdots & \vdots & \vdots & \vdots \\ \frac{\partial f_M(\mathbf{x})}{\partial x_1} & \frac{\partial f_M(\mathbf{x})}{\partial x_2} & \cdots & \frac{\partial f_M(\mathbf{x})}{\partial x_D} \end{pmatrix} \in \mathbb{R}^{M \times D}$$

Chain Rule in 1D

If $f : \mathbb{R} \rightarrow \mathbb{R}$ and $g : \mathbb{R} \rightarrow \mathbb{R}$, then

$$(f \circ g)'(x) = f'(g(x))g'(x)$$

Chain Rule in arbitrary dimensions

If $g : \mathbb{R}^D \rightarrow \mathbb{R}^E$ and $f : \mathbb{R}^E \rightarrow \mathbb{R}^F$, then

$$J_{f \circ g}(\mathbf{x}) = J_f(g(\mathbf{x})) \ J_g(\mathbf{x})$$

Chain Rule in 1D

If $f : \mathbb{R} \rightarrow \mathbb{R}$ and $g : \mathbb{R} \rightarrow \mathbb{R}$, then

$$(f \circ g)'(x) = f'(g(x))g'(x)$$

Chain Rule in arbitrary dimensions

If $g : \mathbb{R}^D \rightarrow \mathbb{R}^E$ and $f : \mathbb{R}^E \rightarrow \mathbb{R}^F$, then

$$J_{f \circ g}(\mathbf{x}) = J_f(g(\mathbf{x})) \ J_g(\mathbf{x})$$

Important Special Case (Vector-Jacobian Product)

If $g : \mathbb{R}^D \rightarrow \mathbb{R}^E$ and $f : \mathbb{R}^E \rightarrow \mathbb{R}$, then $\nabla_{\mathbf{x}}\{f(g(\mathbf{x}))\}^T = \nabla f(g(\mathbf{x}))^T \ J_g(\mathbf{x})$

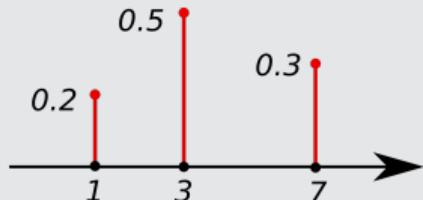
- Jacobians are **the most general form of “derivatives”**
 - Gradients are (transposed) Jacobians
 - Scalar derivatives are 1×1 Jacobians
- Training deep neural nets needs **gradients of loss function** w.r.t. model parameters
- **Backpropagation = efficiently applying the chain rule**
- **Autodiff (AD)**: PyTorch keeps track of Jacobians in the background
 - Systematic computation of the chain rule, directly on the numerical values
 - Maintains computational graph of function
 - Computes **vector-jacobian products** (reverse mode AD)
 - More on this later in the semester

Probability Theory

Random Variables

Discrete Random Variables

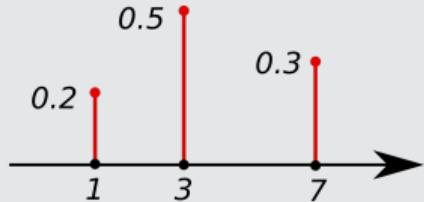
- X is a **discrete random variable** if the set of values \mathcal{X} it can assume is **countable**
 - e.g., $\mathcal{X} = \{1, 3, 7\}$ or $\mathcal{X} = \mathbb{N}$
- **Probability Mass Function:**
$$p(x) = \mathbb{P}(X = x)$$
 - $p : \mathcal{X} \rightarrow [0, 1]$
 - Probability of X assuming $x \in \mathcal{X}$



https://commons.wikimedia.org/wiki/File:Discrete_probability_distrib.svg

Discrete Random Variables

- X is a **discrete random variable** if the set of values \mathcal{X} it can assume is **countable**
 - e.g., $\mathcal{X} = \{1, 3, 7\}$ or $\mathcal{X} = \mathbb{N}$
- **Probability Mass Function:** $p(x) = \mathbb{P}(X = x)$
 - $p : \mathcal{X} \rightarrow [0, 1]$
 - Probability of X assuming $x \in \mathcal{X}$

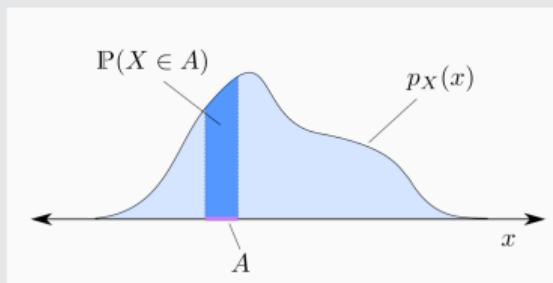


https://commons.wikimedia.org/wiki/File:Discrete_probability_distrib.svg

Continuous Random Variables

- Assume $\mathcal{X} = \mathbb{R}$
- **Probability Density Function** $p : \mathcal{X} \rightarrow [0, \infty)$ such that

$$\mathbb{P}(X \in A) = \int_A p(x) dx$$



- Let $p(x, y)$ be a **joint distribution (mass function)**
 - Assigns probability mass to all possible configurations $\in \mathcal{X} \times \mathcal{Y}$
- Then, we can obtain the **marginal distribution** as follows:

$$p(x) = \sum_{y \in \mathcal{Y}} p(x, y)$$

- If $p(x, y)$ is a joint **density**, we compute

$$p(x) = \int_{\mathcal{Y}} p(x, y) dy$$

- Let $p(x, y)$ be a mass function or density
- The **conditionals** $p(x | y)$ and $p(y | x)$ are given by

$$p(x | y) = \frac{p(x, y)}{p(y)}, \quad p(y | x) = \frac{p(x, y)}{p(x)}$$

- Which also shows how we can **factorize the joint**:

$$p(x | y)p(y) = p(x, y), \quad p(y | x)p(x) = p(x, y)$$

- The **relationship between** $p(y | x)$ and $p(x | y)$ is given by **Bayes' Law**

$$p(y | x) = \frac{p(x | y)p(y)}{p(x)}$$

Expectations

- Let $p(x)$ be a **probability mass function**, $x \in \mathcal{X}$:

$$\mathbb{E}_{x \sim p(x)} [f(x)] = \sum_{x \in \mathcal{X}} p(x)f(x)$$

- Let $p(x)$ be a **probability density function**, $x \in \mathcal{X} = \mathbb{R}$:

$$\mathbb{E}_{x \sim p(x)} [f(x)] = \int_{-\infty}^{\infty} p(x)f(x) dx$$

Monte Carlo Simulation

$$\mathbb{E}_{x \sim p(x)} [f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$

where $x_1, \dots, x_N \stackrel{\text{i.i.d.}}{\sim} p(x)$.

Expectations

Notation

$$\begin{aligned}\mathbb{E}_{x \sim p(x)}[f(x)] &= \int_{-\infty}^{\infty} p(x) f(x) dx \\ &= \mathbb{E}_{p(x)}[f(x)] \\ &= \mathbb{E}[f(x)]\end{aligned}$$

Example: For a joint distribution $p(x, y)$, the extended notation clarifies which variable we integrate over:

$$\mathbb{E}_{(x,y) \sim p^*(x,y)} [\mathcal{L}(y, f(\mathbf{x}))] = \mathbb{E}_{\mathbf{x} \sim p^*(\mathbf{x})} [\mathbb{E}_{y \sim p^*(y | \mathbf{x})} [\mathcal{L}(y, f(\mathbf{x}))]]$$

Properties of Expectations

- **Linearity:**

$$\mathbb{E}_{p(x,y)}[aX + bY] = a\mathbb{E}_{p(x)}[X] + b\mathbb{E}_{p(y)}[Y]$$

with $a, b \in \mathbb{R}$.

- **Expectations of constants:**

$$\mathbb{E}[v] = v$$

with $v \in \mathbb{R}$.

Maximum Likelihood Estimation

- Given a **dataset** = $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$
 - Measurement data $\mathbf{x}^{(i)} \in \mathbb{R}^D$
 - Images $\mathbf{x}^{(i)} \in \mathbb{R}^{H \times W \times C}$
 - Text, molecules, proteins ...
- We assume the **data is drawn i.i.d. from some data generating distribution** $p^*(\mathbf{x})$

- Given a **dataset** = $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$
 - Measurement data $\mathbf{x}^{(i)} \in \mathbb{R}^D$
 - Images $\mathbf{x}^{(i)} \in \mathbb{R}^{H \times W \times C}$
 - Text, molecules, proteins ...
- We assume the **data is drawn i.i.d. from some data generating distribution** $p^*(\mathbf{x})$

Goal (informal)

Learn model $p_\theta(\mathbf{x})$ with parameters θ such that

$$p_\theta \approx p^*$$

Idea

Find model p_θ such that the **observed data** is as **likely as possible** (under p_θ).

Idea

Find model p_θ such that the **observed data** is as **likely as possible** (under p_θ).

- The **Maximum Likelihood Estimate** (MLE) is thus given by

$$\theta_{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} p_\theta(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})$$

- Since $\mathbf{x}^{(i)} \stackrel{\text{i.i.d.}}{\sim} p^*(\mathbf{x})$, we have

$$p_\theta(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}) = \prod_{i=1}^N p_\theta(\mathbf{x}^{(i)})$$

- Hence,

$$\theta_{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^N p_\theta(\mathbf{x}^{(i)})$$

- Since $\log(\cdot)$ is a monotonically increasing function:

$$\operatorname{argmax}_{\theta} \prod_{i=1}^N p_{\theta}(\mathbf{x}^{(i)}) = \operatorname{argmax}_{\theta} \log \left(\prod_{i=1}^N p_{\theta}(\mathbf{x}^{(i)}) \right) = \operatorname{argmax}_{\theta} \sum_{i=1}^N \log \left(p_{\theta}(\mathbf{x}^{(i)}) \right)$$

- Equivalently, we can minimize the *negative log-likelihood* (NLL):

$$\operatorname{argmax}_{\theta} \sum_{i=1}^N \log \left(p_{\theta}(\mathbf{x}^{(i)}) \right) = \operatorname{argmin}_{\theta} \sum_{i=1}^N -\log \left(p_{\theta}(\mathbf{x}^{(i)}) \right)$$

- In most scenarios, our goal is to learn additional information associated with the data, such as the objects present in an image.
- Consider $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$, with $\mathbf{x}^{(i)} \in \mathbb{R}^D, y^{(i)} \in \mathcal{C}$
 - e.g. binary classification, $\mathcal{C} = \{0, 1\}$
- **Assumption:** $(\mathbf{x}^{(i)}, y^{(i)}) \stackrel{\text{i.i.d.}}{\sim} p^*(\mathbf{x}, y)$
- **Discriminative classifier** tries to learn $p_\theta(y | \mathbf{x})$
 - i.e., a distribution **for each \mathbf{x}**

Goal (informal)

Learn model $p_\theta(y | \mathbf{x})$ with parameters θ such that

$$p_\theta(y | \mathbf{x}) \approx p^*(y | \mathbf{x})$$

- **Conditional Maximum Likelihood:**

$$\theta_{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} \ p_{\theta}(y^{(1)}, \dots, y^{(N)} | \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})$$

- We do not care about modeling $p^*(\mathbf{x})$

Notation

$$p_{\text{model}}(y \mid x; \theta) = p_\theta(y \mid x)$$

Both notations describe the **conditional likelihood** of the target y given input x under model parameters θ , informally described as:

$p_\theta(y \mid x)$ = “probability of y given x , under parameters θ ”.

Motivation

What if we already have **some prior knowledge** about what the parameters θ should look like?

Idea

We treat θ as a **random variable** and apply Bayes' rule:

$$p(\theta | \mathbf{X}) = \frac{p(\mathbf{X} | \theta) p(\theta)}{p(\mathbf{X})}.$$

The prior $p(\theta)$ encodes our beliefs about likely parameter values, while the likelihood $p(\mathbf{X} | \theta)$ comes from the data.

Since $p(\mathbf{X})$ does not depend on the parameters θ , we can disregard it when maximizing the posterior.

We therefore choose the parameter value that maximizes the posterior:

$$\theta_{\text{MAP}} = \arg \max_{\theta} p_{\text{model}}(\theta \mid \mathbf{X}) = \arg \max_{\theta} p(\mathbf{X} \mid \theta) p(\theta).$$

Decision Theory

- Supervised learning: $= \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$, with $\mathbf{x}^{(i)} \in \mathbb{R}^D, y^{(i)} \in \mathcal{C}$
- Assume we have the **true posterior** $p^*(y | \mathbf{x})$
- We have perfect knowledge about the **uncertainty in the system**
- What if we **have to make a hard decision** about the value of y

- Supervised learning: $= \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$, with $\mathbf{x}^{(i)} \in \mathbb{R}^D, y^{(i)} \in \mathcal{C}$
- Assume we have the **true posterior** $p^*(y | \mathbf{x})$
- We have perfect knowledge about the **uncertainty in the system**
- What if we **have to make a hard decision** about the value of y

Unfair Coin Toss Example

- Let $\mathcal{C} = \{0, 1\}$ (heads and tails)
- For a fixed \mathbf{x} , assume $p^*(y = 0 | \mathbf{x}) = 0.7$, and $p^*(y = 1 | \mathbf{x}) = 0.3$

If you want to be **correct as often as possible**,
what would be your prediction \hat{y} before flipping the coin?

- You would **always predict** $\hat{y} = 0$, since $p^*(y = 0 | \mathbf{x}) > p^*(y = 1 | \mathbf{x})$

- What if we **alter the game**:
- When you predict correctly, you lose nothing
- When you predict $\hat{y} = 0$, but the coin lands on $y = 1$, you lose 100\$
- When you predict $\hat{y} = 1$, but the coin lands on $y = 0$, you lose 1\$

Is it still a good idea to always predict $\hat{y} = 0$?

- What if we **alter the game**:
- When you predict correctly, you lose nothing
- When you predict $\hat{y} = 0$, but the coin lands on $y = 1$, you lose 100\$
- When you predict $\hat{y} = 1$, but the coin lands on $y = 0$, you lose 1\$

Is it still a good idea to always predict $\hat{y} = 0$?

- **No, surely not**
- We don't care about being right **often**, we care about **losing as little as possible**
 - One type of error is much more painful (false negative)

- We want to find a prediction that minimizes the **expected loss**

$$\operatorname{argmin}_{\hat{y} \in \{0,1\}} \mathbb{E}_{y \sim p^*(y | x)} [\mathcal{L}(y, \hat{y})]$$

with $\mathcal{L}(0, 0) = \mathcal{L}(1, 1) = 0$, $\mathcal{L}(0, 1) = 1$, and $\mathcal{L}(1, 0) = 100$.

- With $\hat{y} = 0$ we have

$$\mathbb{E}_{y \sim p^*(y | x)} [\mathcal{L}(y, \hat{y})] = 0.7 \cdot 0\$ + 0.3 \cdot 100\$ = 30\$$$

- And with $\hat{y} = 1$ we have

$$\mathbb{E}_{y \sim p^*(y | x)} [\mathcal{L}(y, \hat{y})] = 0.7 \cdot 1\$ + 0.3 \cdot 0\$ = 0.7\$$$

- We should predict $\hat{y} = 1$ in this game

- We just focused on **one particular** conditional (**fixed x**)
- To globally act optimally, we can **act optimally for each x**
- For a different x , **the posterior $p(y | x)$ changes**
 - i.e., the coin changes
- We want a **decision function** $f : \mathbb{R}^D \rightarrow \{0, 1\}$ that minimizes

$$\mathbb{E}_{y \sim p^*(y | x)} [\mathcal{L}(y, f(x))]$$

for each x

- This f will then also minimize

$$\mathbb{E}_{(x,y) \sim p^*(x,y)} [\mathcal{L}(y, f(x))] = \mathbb{E}_{x \sim p^*(x)} [\mathbb{E}_{y \sim p^*(y | x)} [\mathcal{L}(y, f(x))]]$$