

AUTOMATIC DIFFERENTIATION

DEEP LEARNING KU (DAT.C302UF)

Simon Hitzinger

Oct 29, 2025

Institute of Machine Learning and Neural Computation
Graz University of Technology, Austria

THE GOAL OF SUPERVISED LEARNING

Training Data

- Given training dataset $\mathcal{D}_{\text{train}} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$
- Features $\mathbf{x} \in \mathbb{R}^d$

Training Data

- Given training dataset $\mathcal{D}_{\text{train}} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$
- Features $\mathbf{x} \in \mathbb{R}^d$

Classification

$$y \in \mathcal{C}$$

e.g., $\mathcal{C} = \{0, 1, \dots, 9\}$

Regression

$$y \in \mathbb{R} \quad \text{or} \quad y \in \mathbb{R}^K$$

Model

- We wish to find a **model** f_{θ} that predicts y from x
 - $\theta \in \mathbb{R}^n$ are the model parameters
- Can be a **point estimate** (e.g. if $y \in \mathbb{R}$):

$$f_{\theta}(x) = \hat{y}$$

- Or a **distribution** over outputs (e.g. if $y \in \mathcal{C} = \{c_1, \dots, c_k\}$):

$$f_{\theta}(x) = (p_{\theta}(y = c_1 | x), \dots, p_{\theta}(y = c_k | x))^{\top}$$

Model

- We wish to find a **model** f_{θ} that predicts y from x
 - $\theta \in \mathbb{R}^n$ are the model parameters
- Can be a **point estimate** (e.g. if $y \in \mathbb{R}$):

$$f_{\theta}(x) = \hat{y}$$

- Or a **distribution** over outputs (e.g. if $y \in \mathcal{C} = \{c_1, \dots, c_k\}$):

$$f_{\theta}(x) = (p_{\theta}(y = c_1 | x), \dots, p_{\theta}(y = c_k | x))^{\top}$$

Loss Function

- We specify a **loss function**

$$\ell(f_{\theta}(x), y) \in \mathbb{R}$$

that measures **deviation** between the prediction $f_{\theta}(x)$ and the “true” solution y

Training Loss 💪

- We can now compute the **average loss** over all training examples:

$$\mathcal{L}_{\text{train}}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N \ell(f_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}), y^{(i)})$$

- Notice that $\mathcal{L}_{\text{train}} : \mathbb{R}^n \rightarrow \mathbb{R}$!

Question 🤔

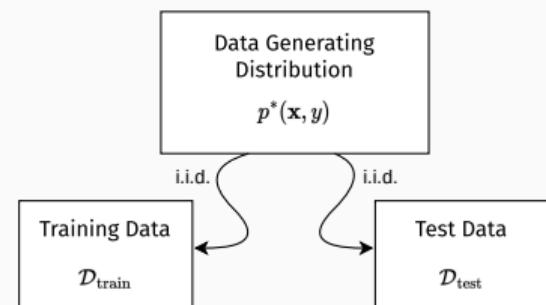
Is the final goal of supervised learning to minimize the average training loss? That is, finding

$$\boldsymbol{\theta}^* = \operatorname{argmin}_{\boldsymbol{\theta}} \mathcal{L}_{\text{train}}(\boldsymbol{\theta})$$

- Statistical Learning Theory can answer this!
- If we had access to $p^*(x, y)$, we want to minimize the generalization error:

$$\mathcal{L}^*(\theta) = \mathbb{E}_{p^*(x,y)} [\ell(f_\theta(x), y)]$$

- But since we don't know p^* , we minimize $\mathcal{L}_{\text{train}}$
- However, we **truly care about the average test loss** $\mathcal{L}_{\text{test}}$, since this is an unbiased estimate of the generalization error



MINIMIZING TRAINING LOSS

- Let's minimize the training loss

$$\operatorname{argmin}_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(f_{\theta}(x^{(i)}), y^{(i)})$$

- In *Deep Learning*, there is generally no closed form solution for this 😞

- Let's minimize the training loss

$$\operatorname{argmin}_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(f_{\theta}(x^{(i)}), y^{(i)})$$

- In *Deep Learning*, there is generally no closed form solution for this 😞
- Hence, we rely on iterative optimization algorithms that use gradient information:

$$\nabla_{\theta} \mathcal{L}_{\text{train}}(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \ell(f_{\theta}(x^{(i)}), y^{(i)})$$

- Is there a problem with this ?

- Let's minimize the training loss

$$\operatorname{argmin}_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(f_{\theta}(x^{(i)}), y^{(i)})$$

- In *Deep Learning*, there is generally no closed form solution for this 😞
- Hence, we rely on iterative optimization algorithms that use gradient information:

$$\nabla_{\theta} \mathcal{L}_{\text{train}}(\theta) = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \ell(f_{\theta}(x^{(i)}), y^{(i)})$$

- Is there a problem with this ?
- For large N , this is computationally expensive. Instead, let's approximate this gradient with a Mini-Batch Gradient (with $B \ll N$):

$$\nabla_{\theta} \mathcal{L}_{\text{train}}(\theta) \approx \nabla_{\theta} \mathcal{L}_B(\theta) = \frac{1}{B} \sum_{i=1}^B \nabla_{\theta} \ell(f_{\theta}(x^{(i)}), y^{(i)})$$

- Thankfully, we don't need to compute the gradients by hand 🙏
- Deep Learning Frameworks offer Automatic Differentiation (autodiff) 😍

AUTOMATIC DIFFERENTIATION

CREDIT: MATT JOHNSON'S GREAT TALK ON AUTODIFF

[JOHNSON, 2017]

$$F : \mathbb{R}^n \rightarrow \mathbb{R}$$

$$F : \begin{array}{c} \text{[vertical bar]} \\ \mapsto \\ \text{[square]} \\ y \in \mathbb{R} \end{array}$$

$x \in \mathbb{R}^n$

$$F : \mathbb{R}^n \rightarrow \mathbb{R}$$

$$F : \begin{array}{c} \text{[vertical bar]} \\ \mapsto \\ \text{[square]} \end{array} \quad \begin{matrix} & y \in \mathbb{R} \\ \text{[blue x]} & \in \mathbb{R}^n \end{matrix}$$

$$F = D \circ C \circ B \circ A$$

$$\textcolor{red}{y} = F(\textcolor{blue}{x}) = D(C(B(A(\textcolor{blue}{x}))))$$

$$F : \mathbb{R}^n \rightarrow \mathbb{R}$$

$$F : \begin{array}{c} \text{[vertical bar]} \\ \mapsto \\ \text{[square]} \end{array} \quad \begin{matrix} & y \in \mathbb{R} \\ \textcolor{blue}{x} \in \mathbb{R}^n \end{matrix}$$

$$F = D \circ C \circ B \circ A \quad \quad \quad \textcolor{red}{y} = F(\textcolor{blue}{x}) = D(C(B(A(\textcolor{blue}{x}))))$$

$$\textcolor{red}{y} = D(\textcolor{blue}{c}), \quad \textcolor{blue}{c} = C(\textcolor{blue}{b}), \quad \textcolor{blue}{b} = B(\textcolor{blue}{a}), \quad \textcolor{blue}{a} = A(\textcolor{blue}{x})$$

$$\textcolor{red}{y} = D(\mathbf{c}), \quad \mathbf{c} = C(\mathbf{b}), \quad \mathbf{b} = B(\mathbf{a}), \quad \mathbf{a} = A(\textcolor{blue}{x})$$

$$F'(\textcolor{blue}{x}) = \frac{\partial \textcolor{red}{y}}{\partial \textcolor{blue}{x}} = \begin{bmatrix} \frac{\partial \textcolor{red}{y}}{\partial \textcolor{blue}{x}_1} & \cdots & \frac{\partial \textcolor{red}{y}}{\partial \textcolor{blue}{x}_n} \end{bmatrix}$$

CHAIN RULE

$$\textcolor{red}{y} = D(\mathbf{c}), \quad \mathbf{c} = C(\mathbf{b}), \quad \mathbf{b} = B(\mathbf{a}), \quad \mathbf{a} = A(\textcolor{blue}{x})$$

$$F'(\textcolor{blue}{x}) = \frac{\partial \textcolor{red}{y}}{\partial \textcolor{blue}{x}} = \begin{bmatrix} \frac{\partial \textcolor{red}{y}}{\partial \textcolor{blue}{x}_1} & \cdots & \frac{\partial \textcolor{red}{y}}{\partial \textcolor{blue}{x}_n} \end{bmatrix}$$

$$F'(\textcolor{blue}{x}) = \begin{bmatrix} \frac{\partial \textcolor{red}{y}}{\partial \mathbf{c}} & \frac{\partial \mathbf{c}}{\partial \mathbf{b}} & \frac{\partial \mathbf{b}}{\partial \mathbf{a}} & \frac{\partial \mathbf{a}}{\partial \textcolor{blue}{x}} \end{bmatrix}$$

CHAIN RULE

$$\textcolor{red}{y} = D(\mathbf{c}), \quad \mathbf{c} = C(\mathbf{b}), \quad \mathbf{b} = B(\mathbf{a}), \quad \mathbf{a} = A(\mathbf{x})$$

$$F'(\mathbf{x}) = \frac{\partial \textcolor{red}{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial \textcolor{red}{y}}{\partial \mathbf{x}_1} & \cdots & \frac{\partial \textcolor{red}{y}}{\partial \mathbf{x}_n} \end{bmatrix}$$

$$F'(\mathbf{x}) = \frac{\partial \textcolor{red}{y}}{\partial \mathbf{c}} \quad \frac{\partial \mathbf{c}}{\partial \mathbf{b}} \quad \frac{\partial \mathbf{b}}{\partial \mathbf{a}} \quad \frac{\partial \mathbf{a}}{\partial \mathbf{x}}$$
$$\frac{\partial \textcolor{red}{y}}{\partial \mathbf{c}} = D'(\mathbf{c}) \quad \frac{\partial \mathbf{c}}{\partial \mathbf{b}} = C'(\mathbf{b}) \quad \frac{\partial \mathbf{b}}{\partial \mathbf{a}} = B'(\mathbf{a}) \quad \frac{\partial \mathbf{a}}{\partial \mathbf{x}} = A'(\mathbf{x})$$



FORWARD VS. REVERSE AUTODIFF

$$F'(\mathbf{x}) = \frac{\partial \mathbf{y}}{\partial \mathbf{c}} \underbrace{\left(\frac{\partial \mathbf{c}}{\partial \mathbf{b}} \begin{pmatrix} \frac{\partial \mathbf{b}}{\partial \mathbf{a}} & \frac{\partial \mathbf{a}}{\partial \mathbf{x}} \end{pmatrix} \right)}_{\frac{\partial \mathbf{b}}{\partial \mathbf{x}}} \quad \text{Forward accumulation}$$
$$\frac{\partial \mathbf{b}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial b_1}{\partial x_1} & \dots & \frac{\partial b_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial b_m}{\partial x_1} & \dots & \frac{\partial b_m}{\partial x_n} \end{bmatrix}$$

FORWARD VS. REVERSE AUTODIFF

$$F'(\mathbf{x}) = \underbrace{\frac{\partial \mathbf{y}}{\partial \mathbf{c}} \left(\frac{\partial \mathbf{c}}{\partial \mathbf{b}} \left(\frac{\partial \mathbf{b}}{\partial \mathbf{a}} \frac{\partial \mathbf{a}}{\partial \mathbf{x}} \right) \right)}_{\frac{\partial \mathbf{b}}{\partial \mathbf{x}}} \quad \text{Forward accumulation}$$

$$F'(\mathbf{x}) = \underbrace{\left(\left(\frac{\partial \mathbf{y}}{\partial \mathbf{c}} \frac{\partial \mathbf{c}}{\partial \mathbf{b}} \right) \frac{\partial \mathbf{b}}{\partial \mathbf{a}} \right) \frac{\partial \mathbf{a}}{\partial \mathbf{x}}}_{\frac{\partial \mathbf{y}}{\partial \mathbf{b}}} \quad \text{Reverse accumulation}$$

Reverse-Mode Autodiff (“Backpropagation”)

- `loss.backward()` does reverse-mode autodiff !
 - Hence this will fail if `loss` is not a scalar
- Implemented using only Vector-Jacobian products $v^T J$
- Backward pass needs \approx same amount of compute than forward pass 😍
- But we need to store all activations in the forward pass 😓

Reverse-Mode Autodiff (“Backpropagation”)

- `loss.backward()` does reverse-mode autodiff !
 - Hence this will fail if `loss` is not a scalar
- Implemented using only Vector-Jacobian products $v^T J$
- Backward pass needs \approx same amount of compute than forward pass 😍
- But we need to store all activations in the forward pass 😭

Forward-Mode Autodiff

- Constant memory overhead 😊
- Compute inefficient if we compute gradients of $\mathcal{L} : \mathbb{R}^n \rightarrow \mathbb{R}$ 😫
- In typical Deep Learning scenarios, you rarely care about forward-mode AD

AUTODIFF DEMO IN PyTorch

REFERENCES I

-  Johnson, M. J. (2017).
Automatic differentiation.
Presented at the Deep Learning Summer School 2017 in Montreal.