

# Deep Learning: Convolutional Neural Networks

**Robert Legenstein**

Institute of Machine Learning and Neural Computation

[Robert.legenstein@tugraz.at](mailto:Robert.legenstein@tugraz.at)

Deep Learning VO - WS 25/26

Lecture 7

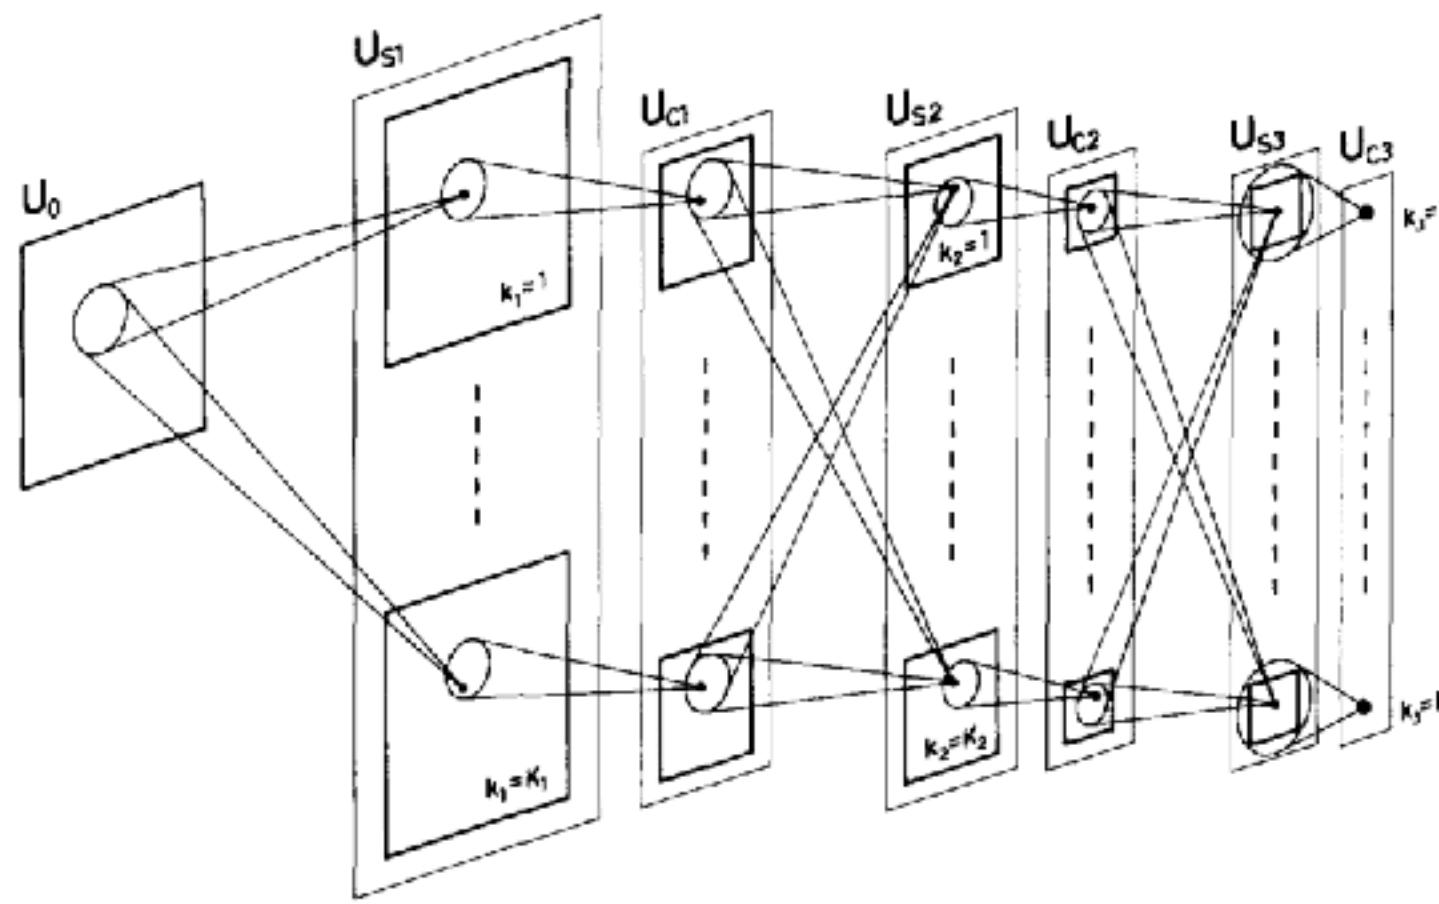
# Today

---

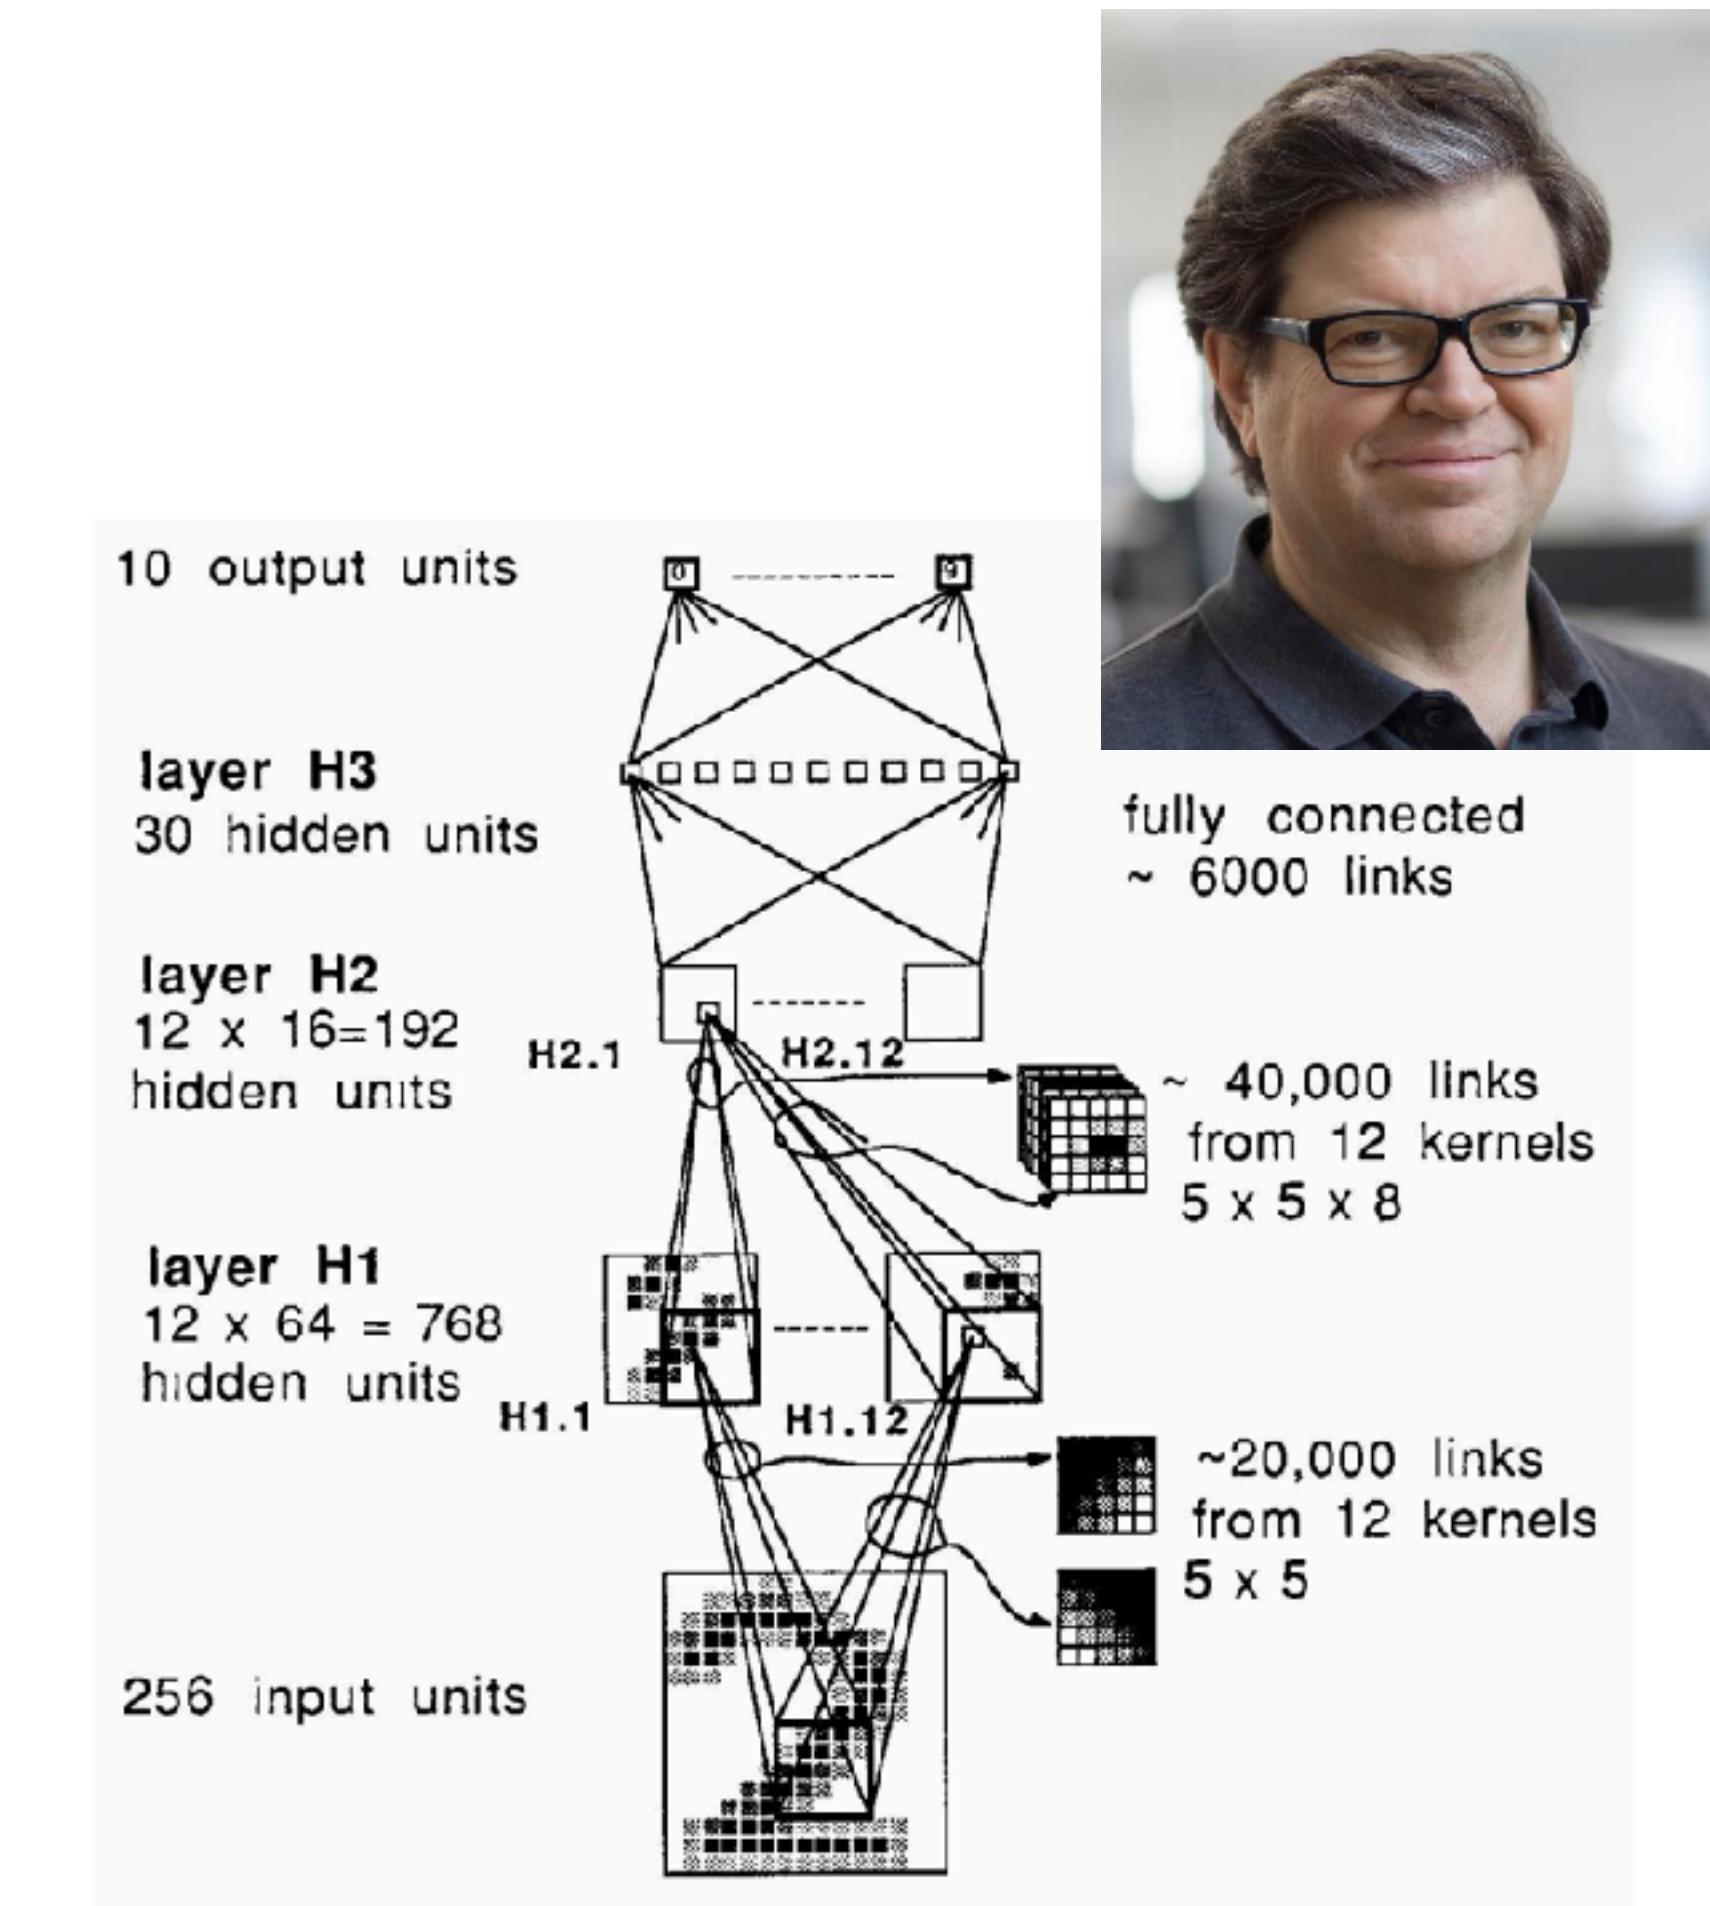
- ❑ Convolutional Neural Networks
  - ❑ Background
  - ❑ CNN Architecture
  - ❑ Going Deeper with CNNs
  - ❑ Residual Networks (ResNets) & Extensions

# Convolutional Neural Networks

- Convolutional neural networks (CNNs) were introduced already in 1989 by Yann LeCun.
- Prior work by Kunihiko Fukushima: "Neocognitron"
- Inspired by the architecture of the visual system of vertebrates.



[K. Fukushima (1980): "Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position", *Biological Cybernetics*.]



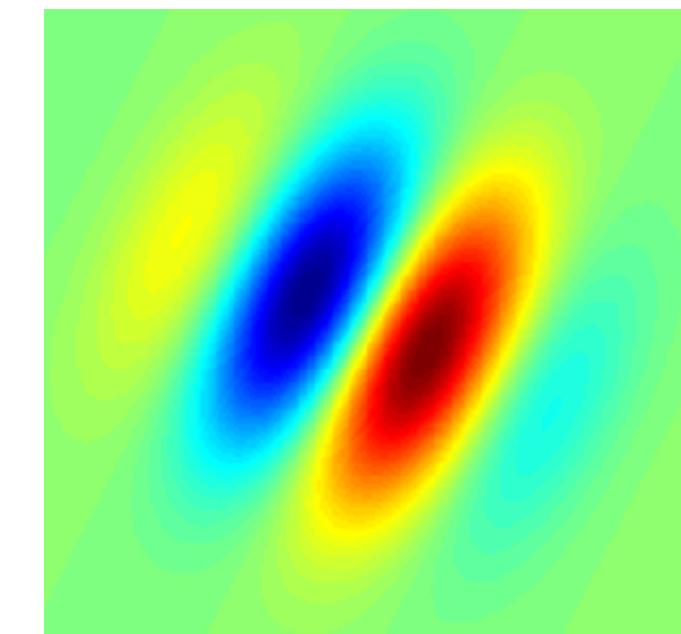
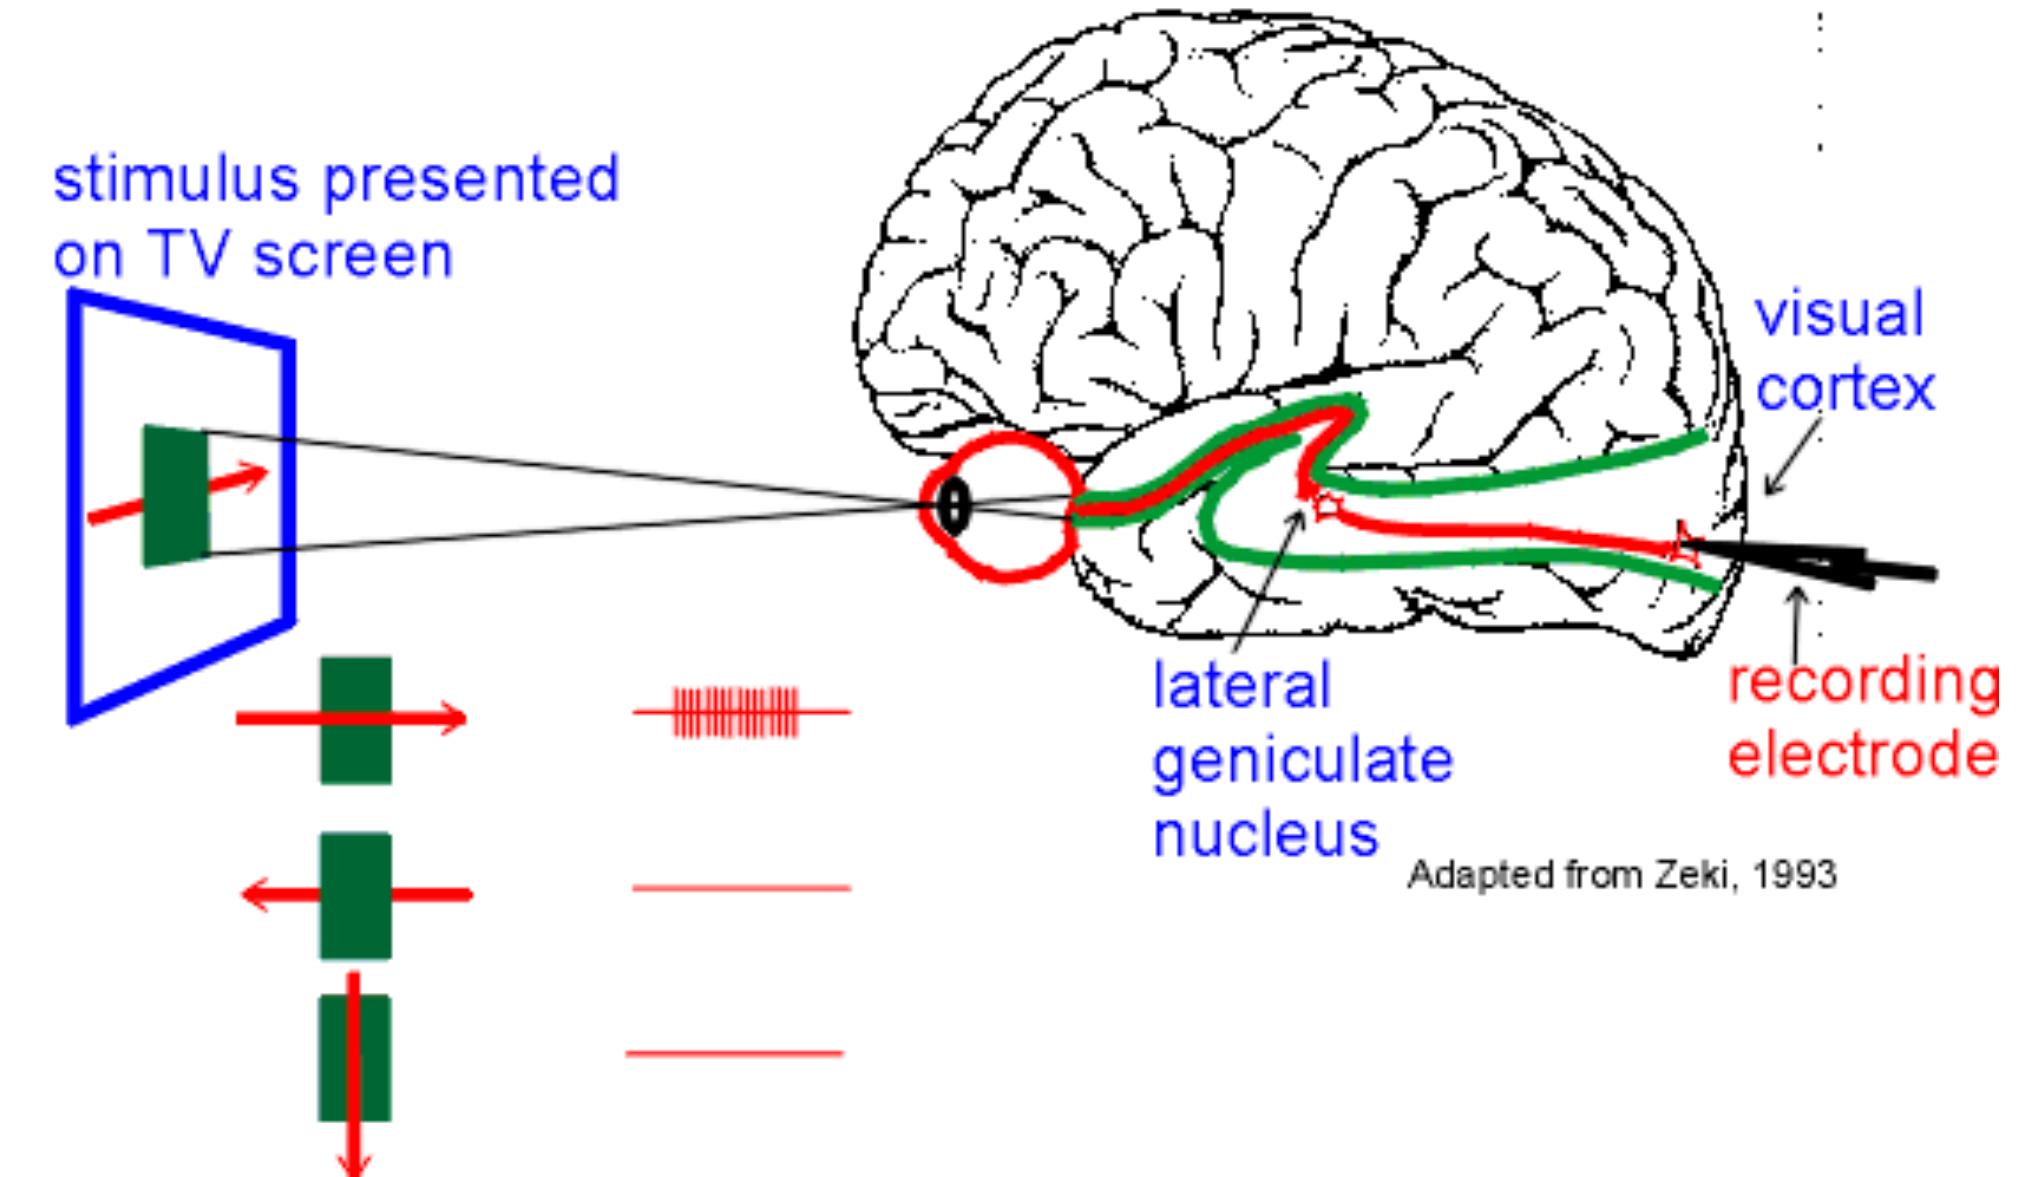
[Y. LeCun, et al. (1989): "Backpropagation applied to handwritten zip code recognition", *Neural Computation*.]  
[Y. LeCun (1989): "Generalization and network design strategies", TechRep CRG-TR-89-4 Univ. of Toronto.]



# Background on Vision

## Hubel and Wiesel (1950s & 60s):

- Neurons in cat and monkey visual cortices respond to small regions of the visual fields.
- Neighboring cells have similar and overlapping receptive fields.
- Simple cells: Output is maximized by straight edges having particular orientations.
- Complex cells: Output is insensitive to the exact position of the edges in the field.

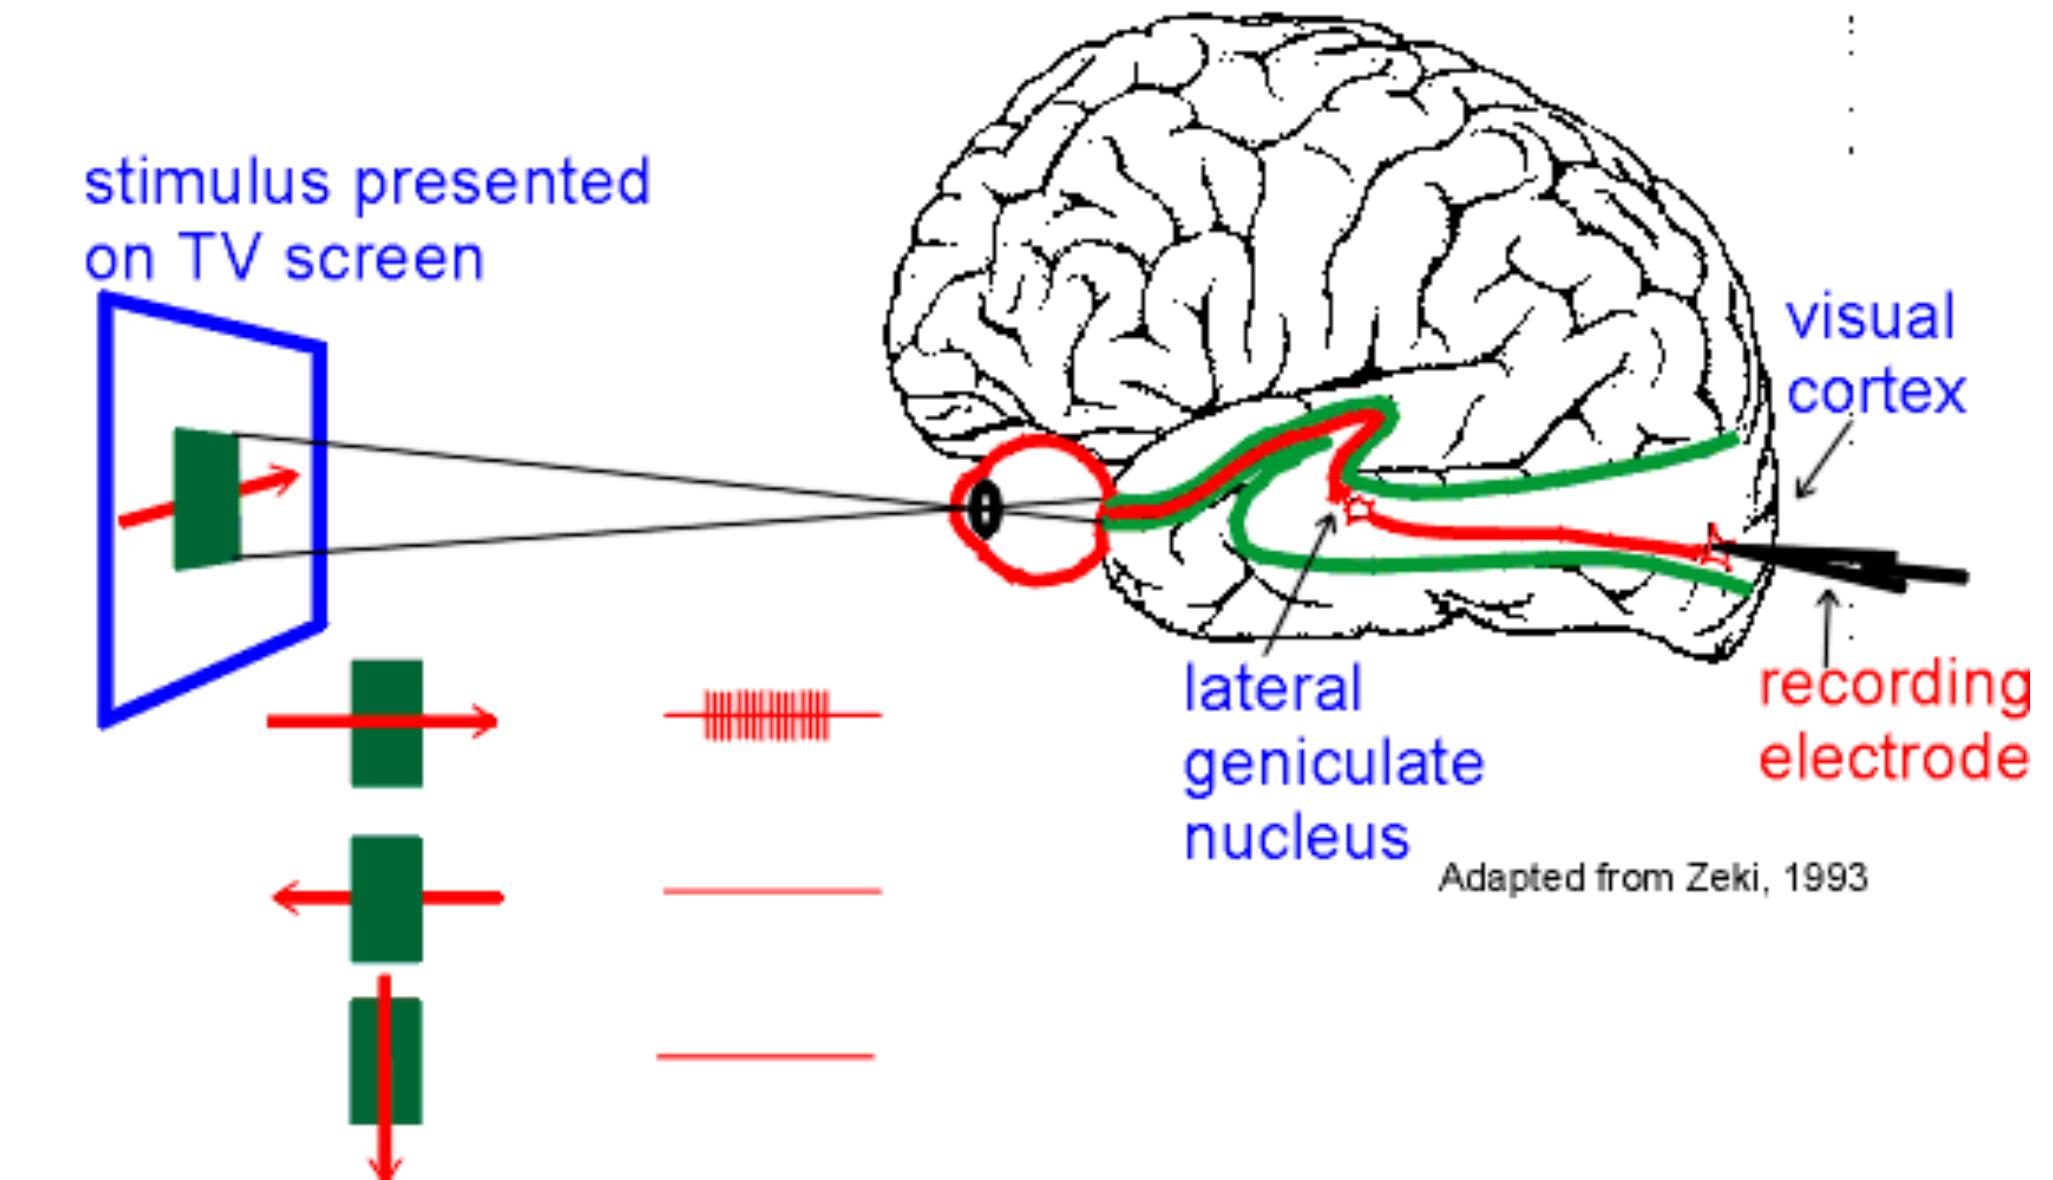


a Gabor filter

# Background on Vision

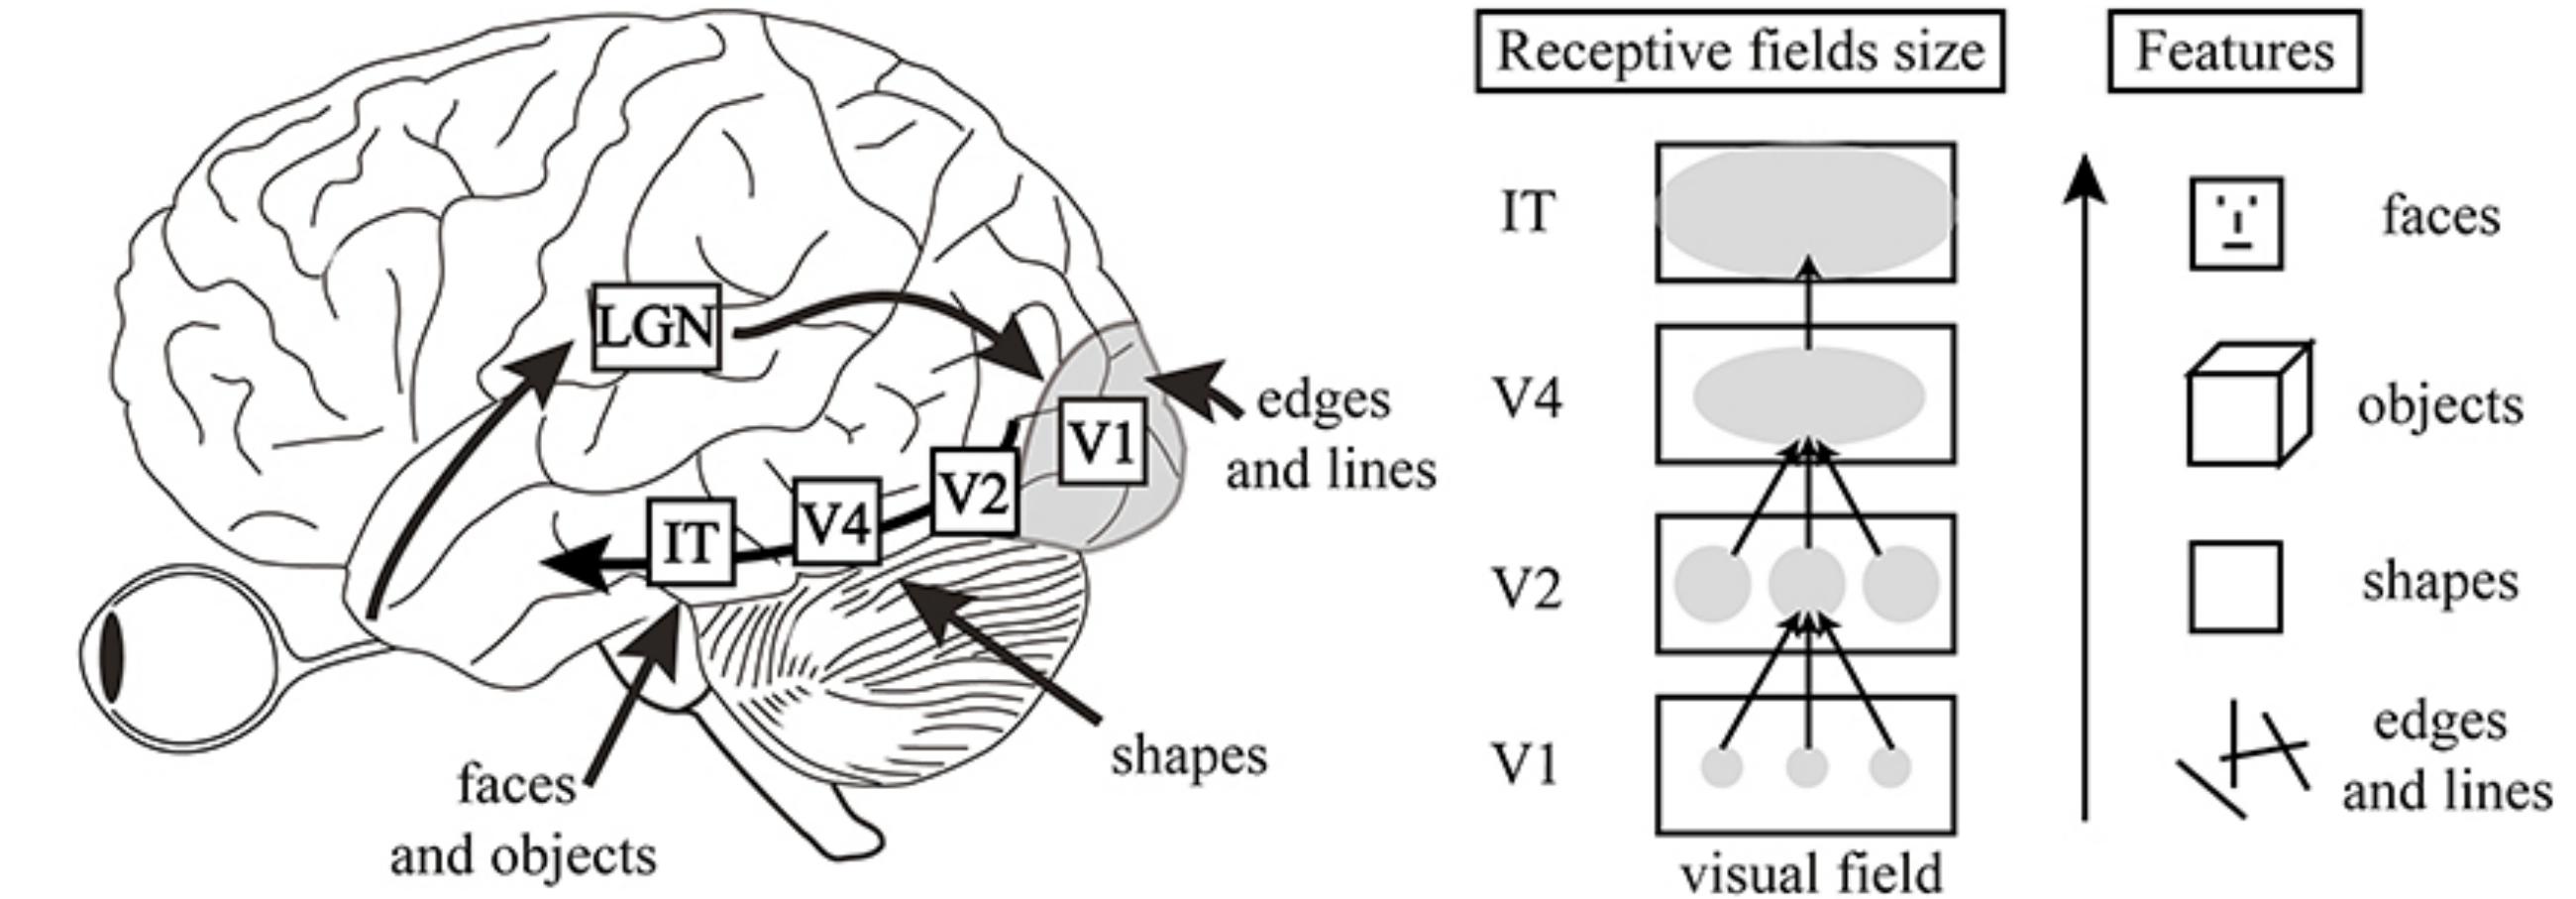
## Hubel and Wiesel (1950s & 60s):

- Neurons in cat and monkey visual cortices respond to small regions of the visual fields.
- Neighboring cells have similar and overlapping receptive fields.



## The visual system:

- Hierarchical organization of areas with increasing receptive field sizes and increasingly complex features.



[Michael H. Herzog and Aaron M. Clarke. (2014): "Why vision is not both hierarchical and feedforward"]

# Convolutional Networks

---

## Induced invariances by network structure (Inductive Bias).

Convolutional neural networks (CNNs) are:

- Feed-forward neural networks.
- Typically applied on inputs with local structure (images, video, audio).
- Trained end-to-end with stochastic gradient descent (or a variants like Adam).

## Statistical insights for image classification problems:

- Images have local features.
  - Close pixels are more strongly correlated than distant ones.
- Local features can appear at different locations.
  - e.g., edges, or objects at different positions.
- Output should be invariant to local distortions or position shifts.

# Today

---

- Convolutional Neural Networks
- Background
- CNN Architecture
- Going Deeper with CNNs
- Residual Networks (ResNets) & Extensions

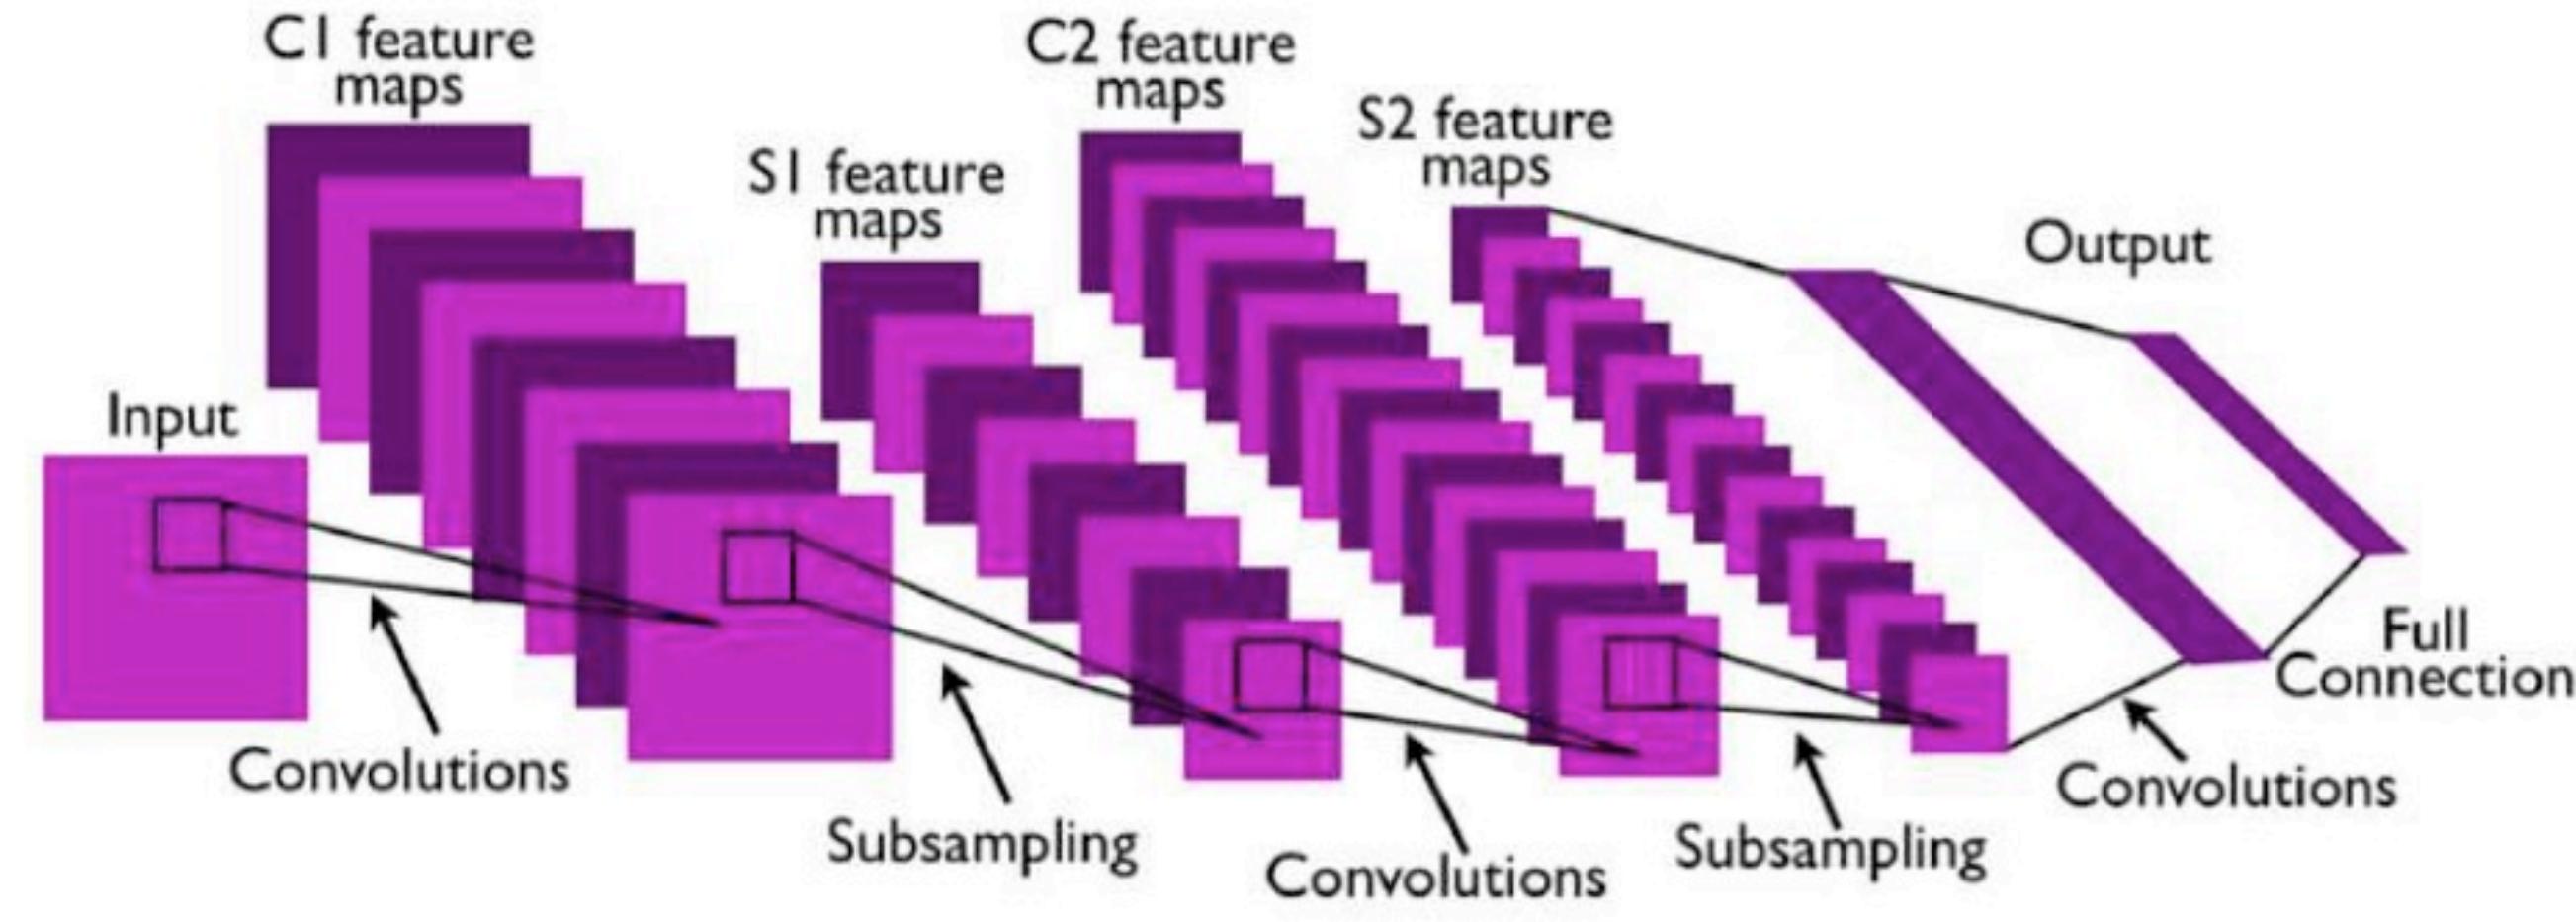
# CNN Architecture

## (1) Convolutional feature maps:

- Extract features at all possible locations.

## (2) Pooling (subsampling):

- Invariance to local distortions.
- Reduce feature map size.



LeNet-5 Architecture (1995)

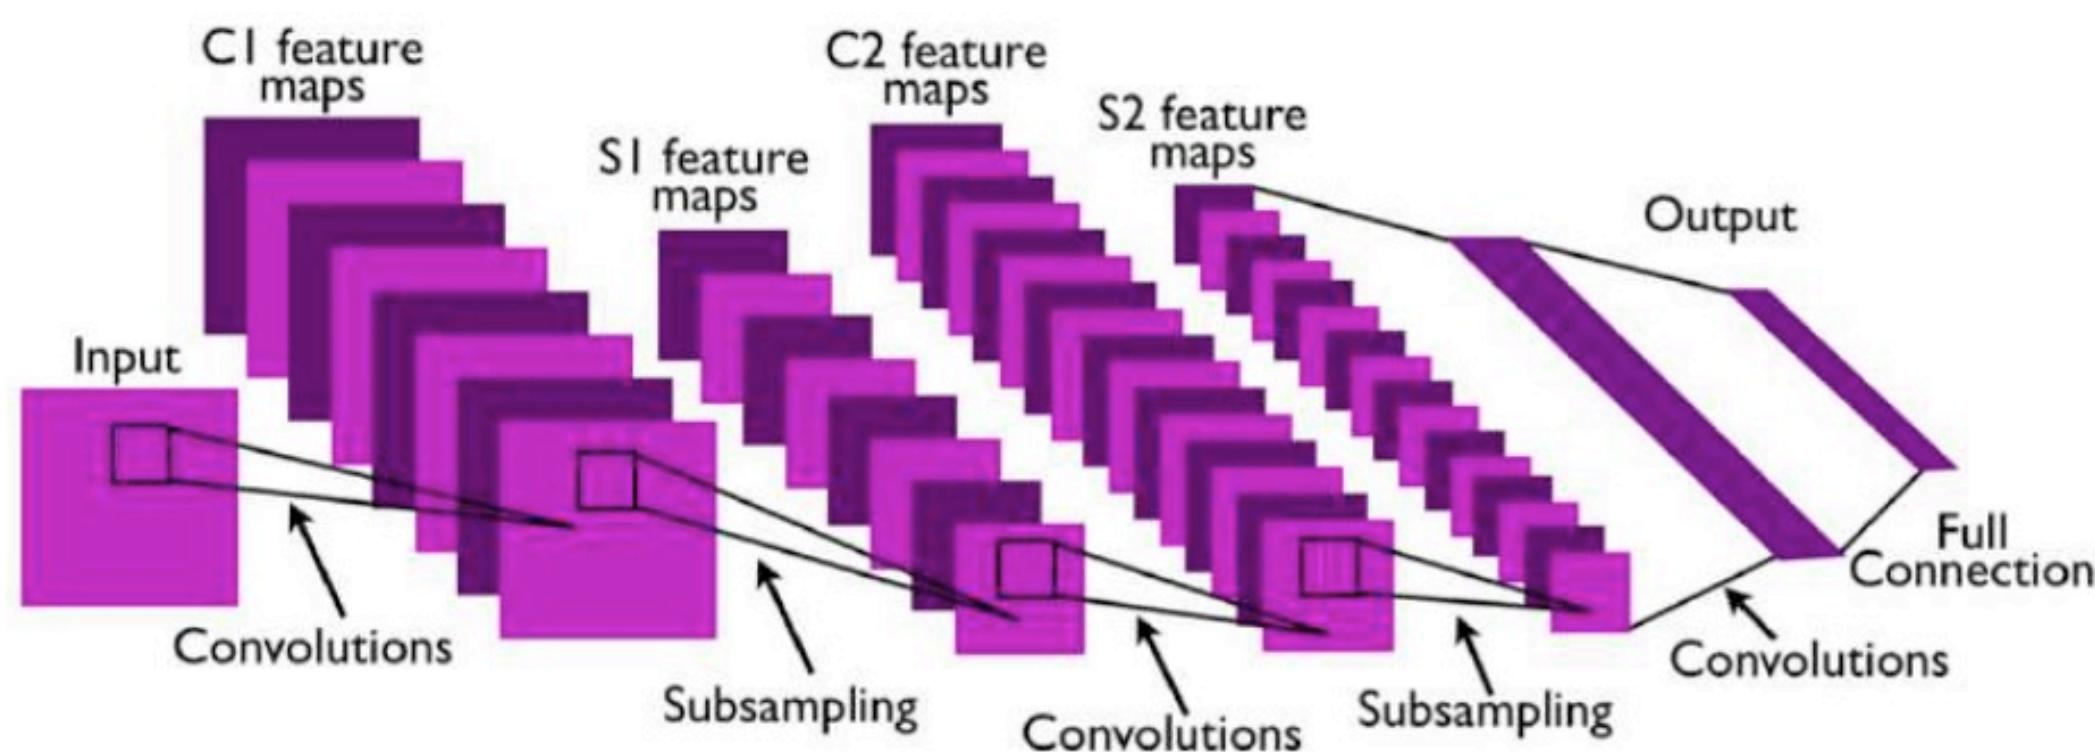
# CNN Architecture

## (1) Convolutional feature maps:

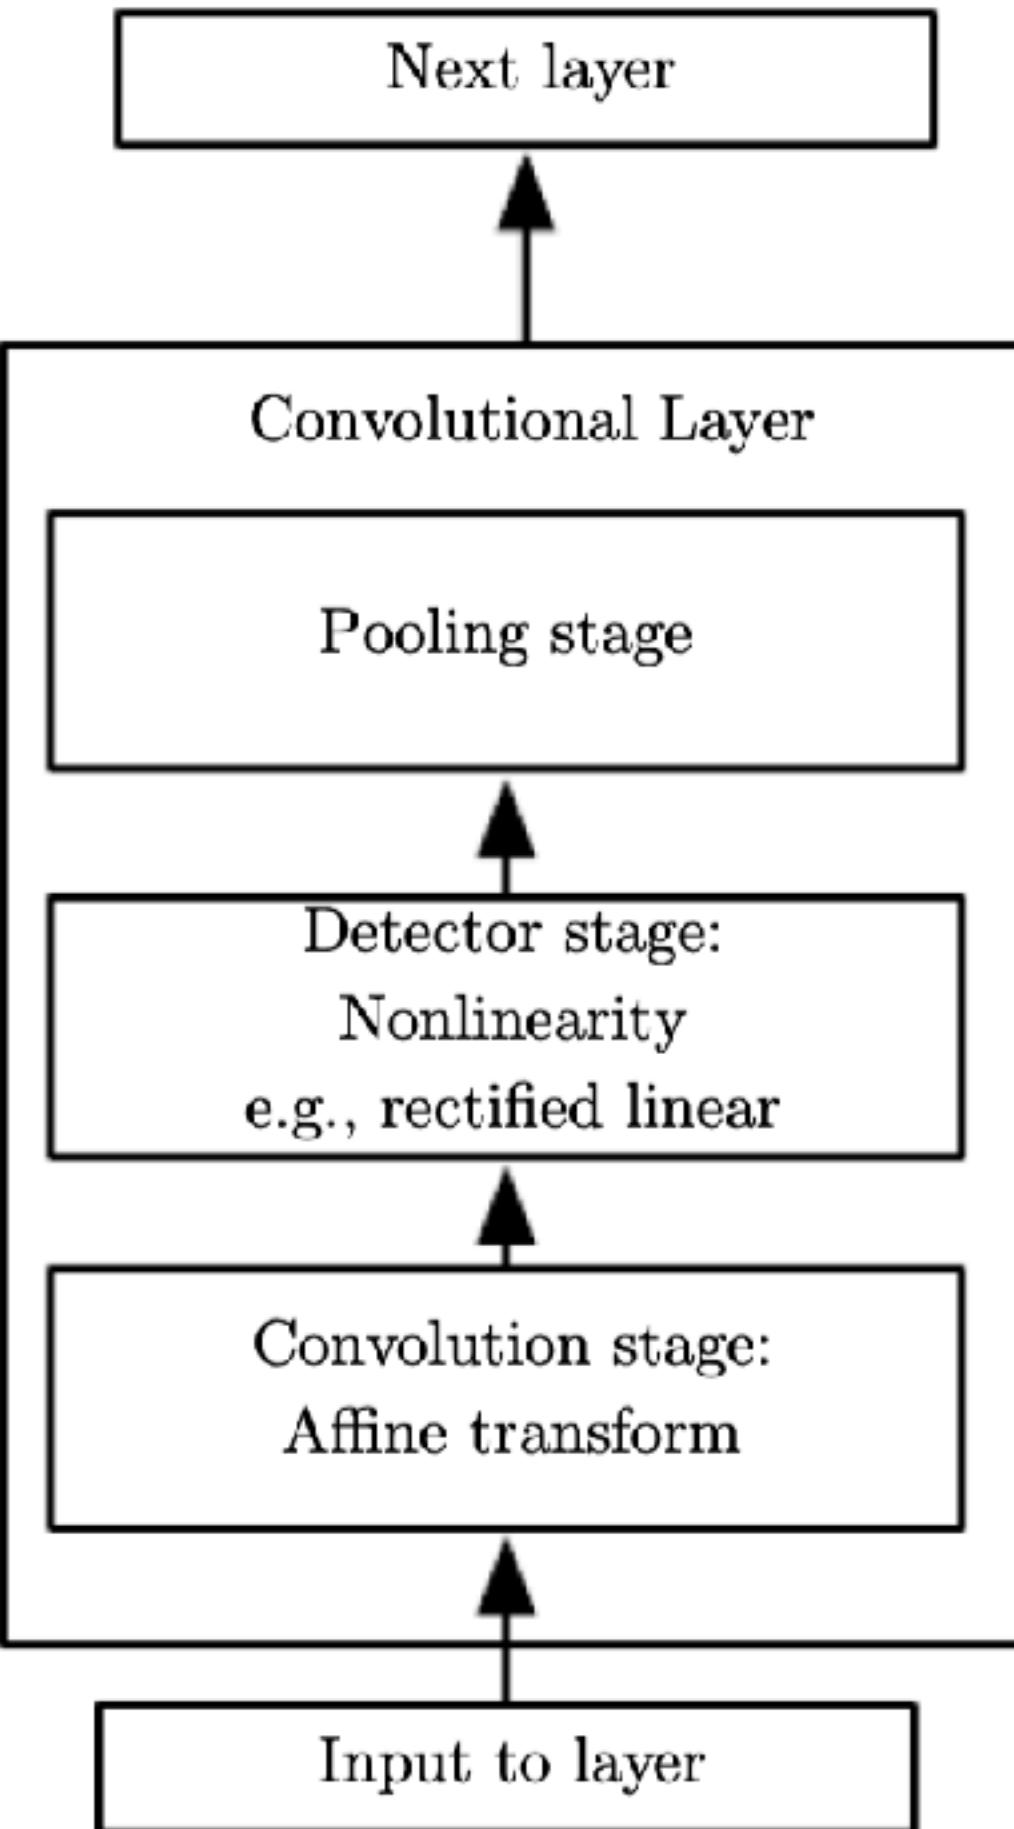
- Extract features at all possible locations.

## (2) Pooling (subsampling):

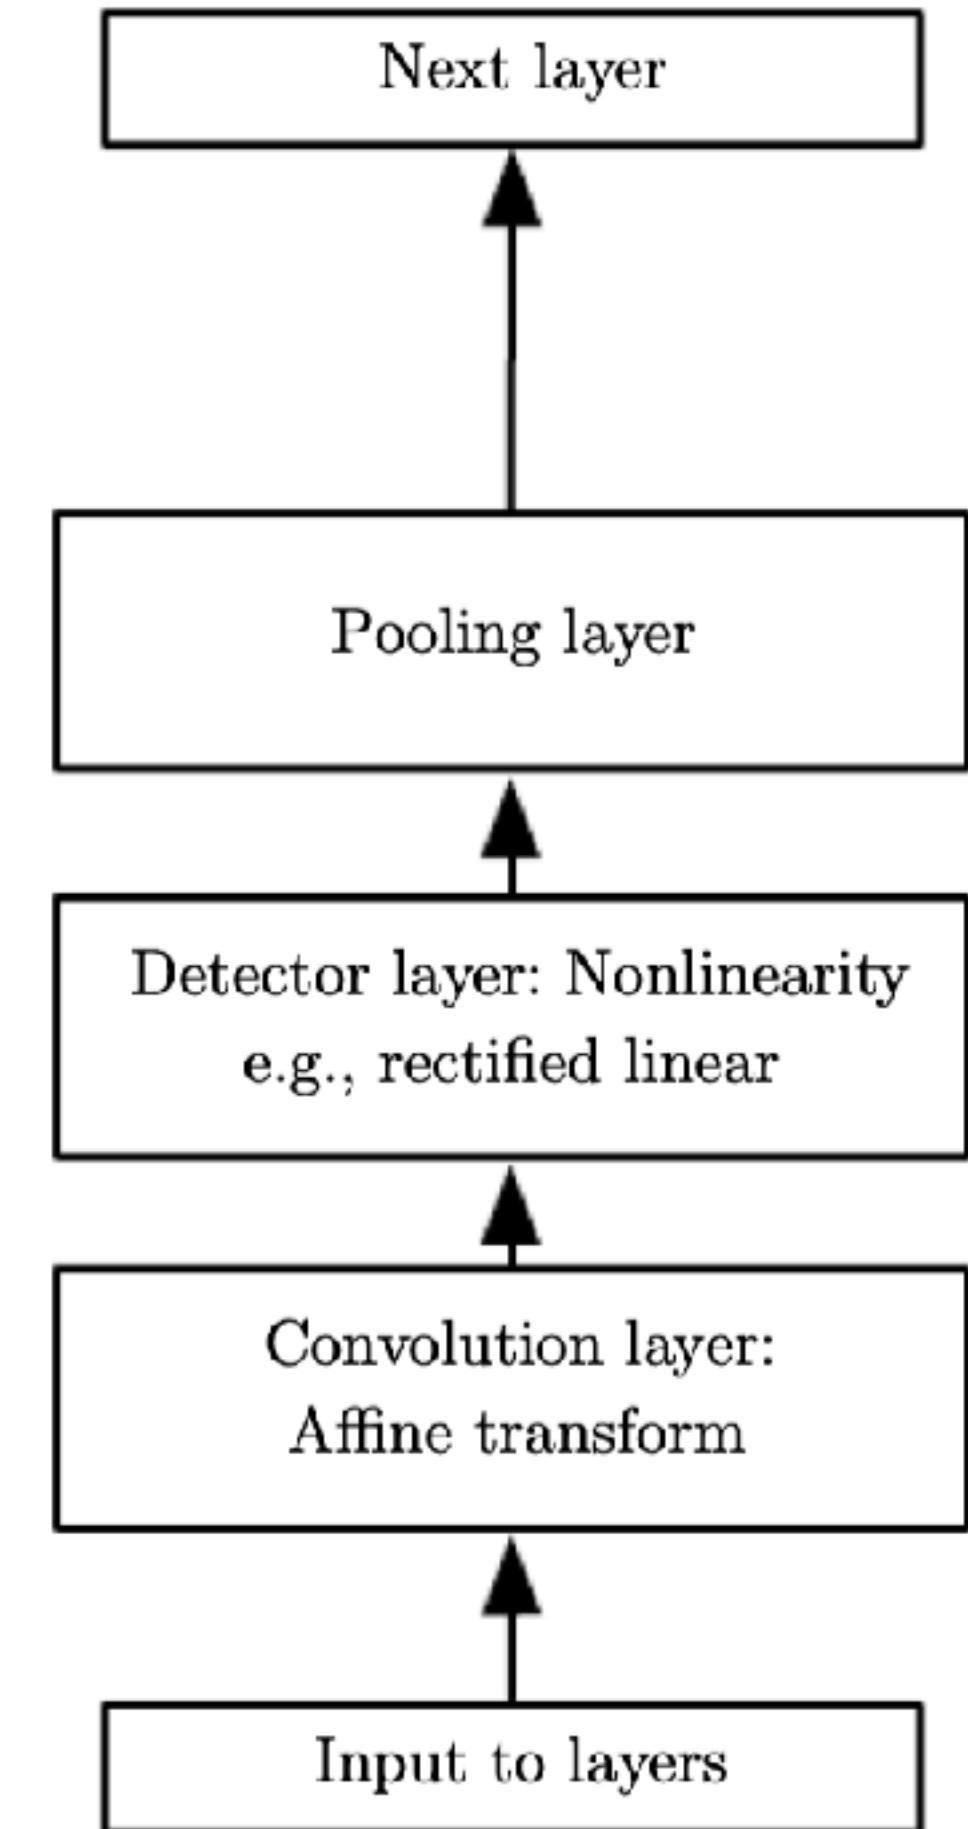
- Invariance to local distortions.
- Reduce feature map size.



Complex layer terminology

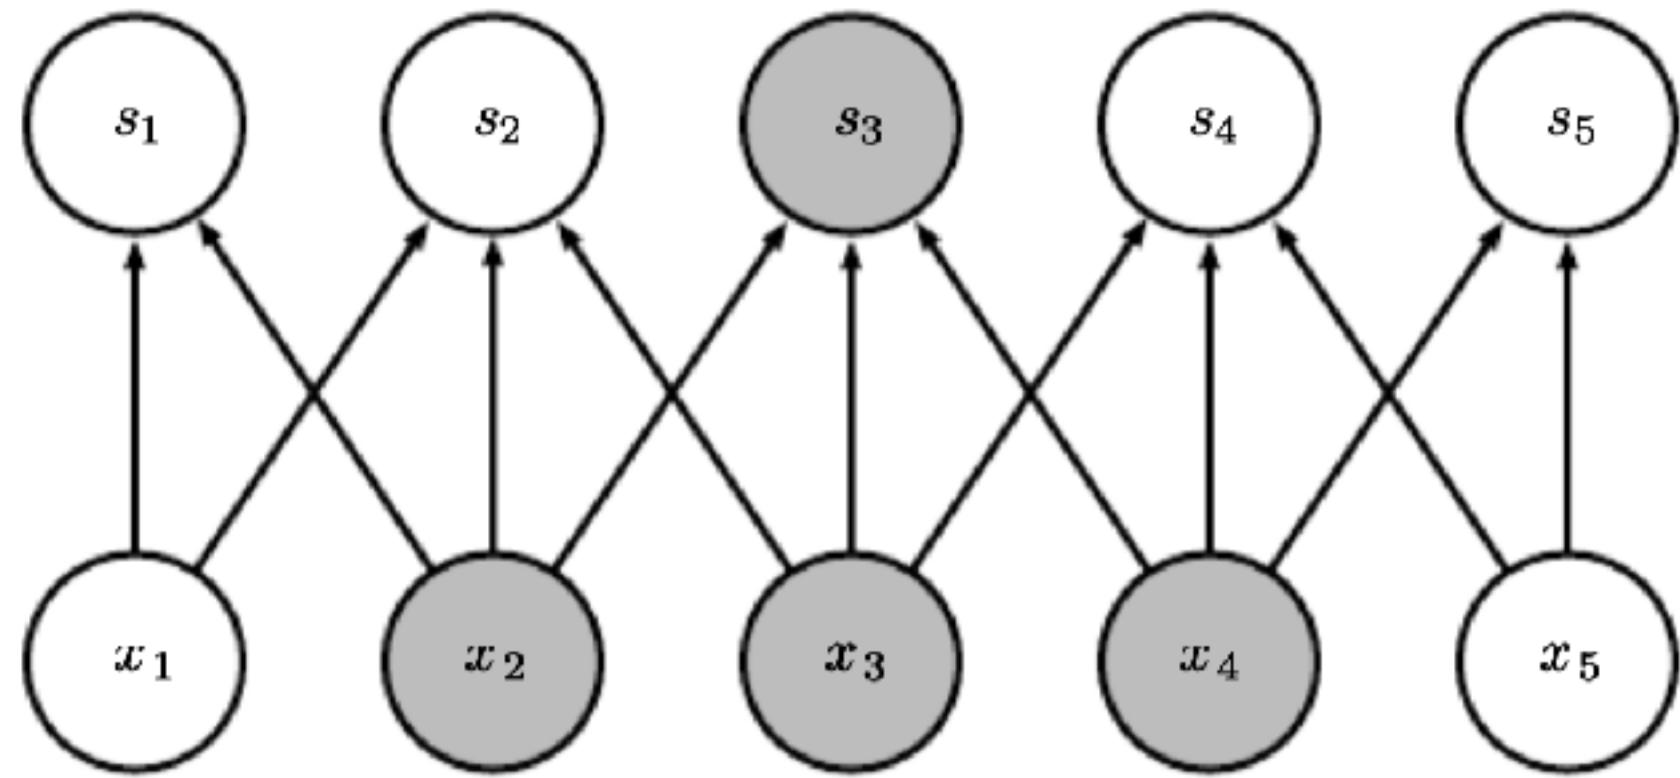


Simple layer terminology

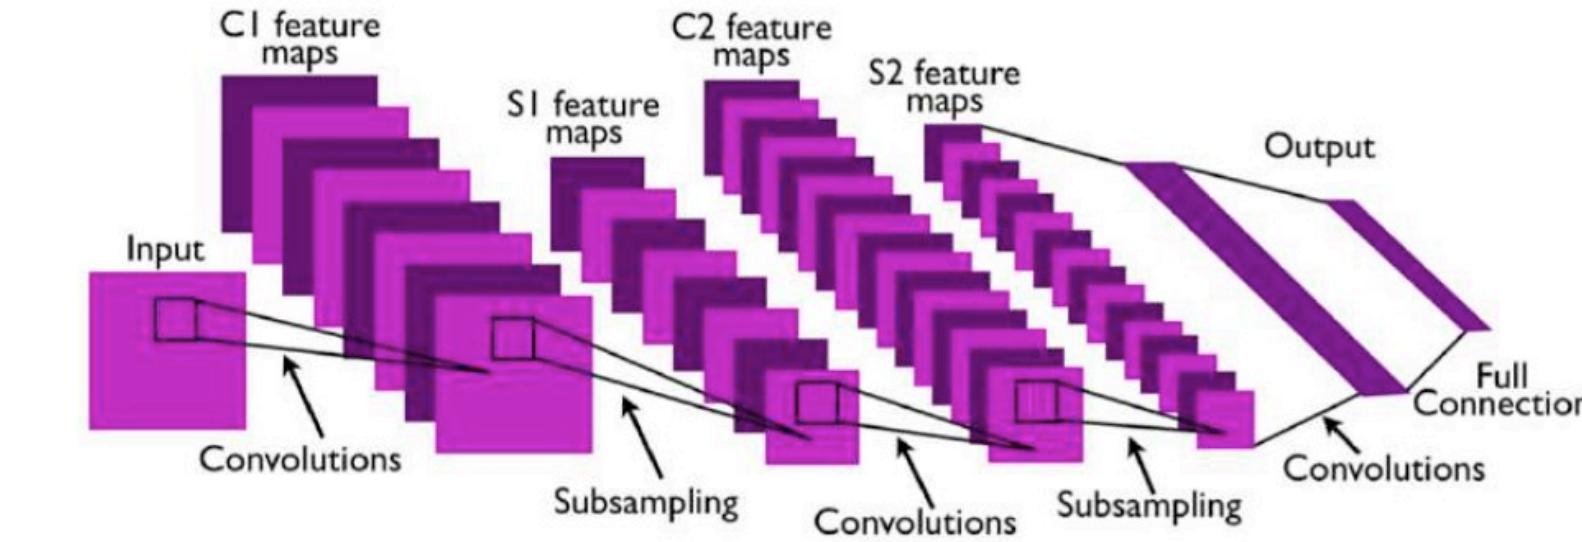
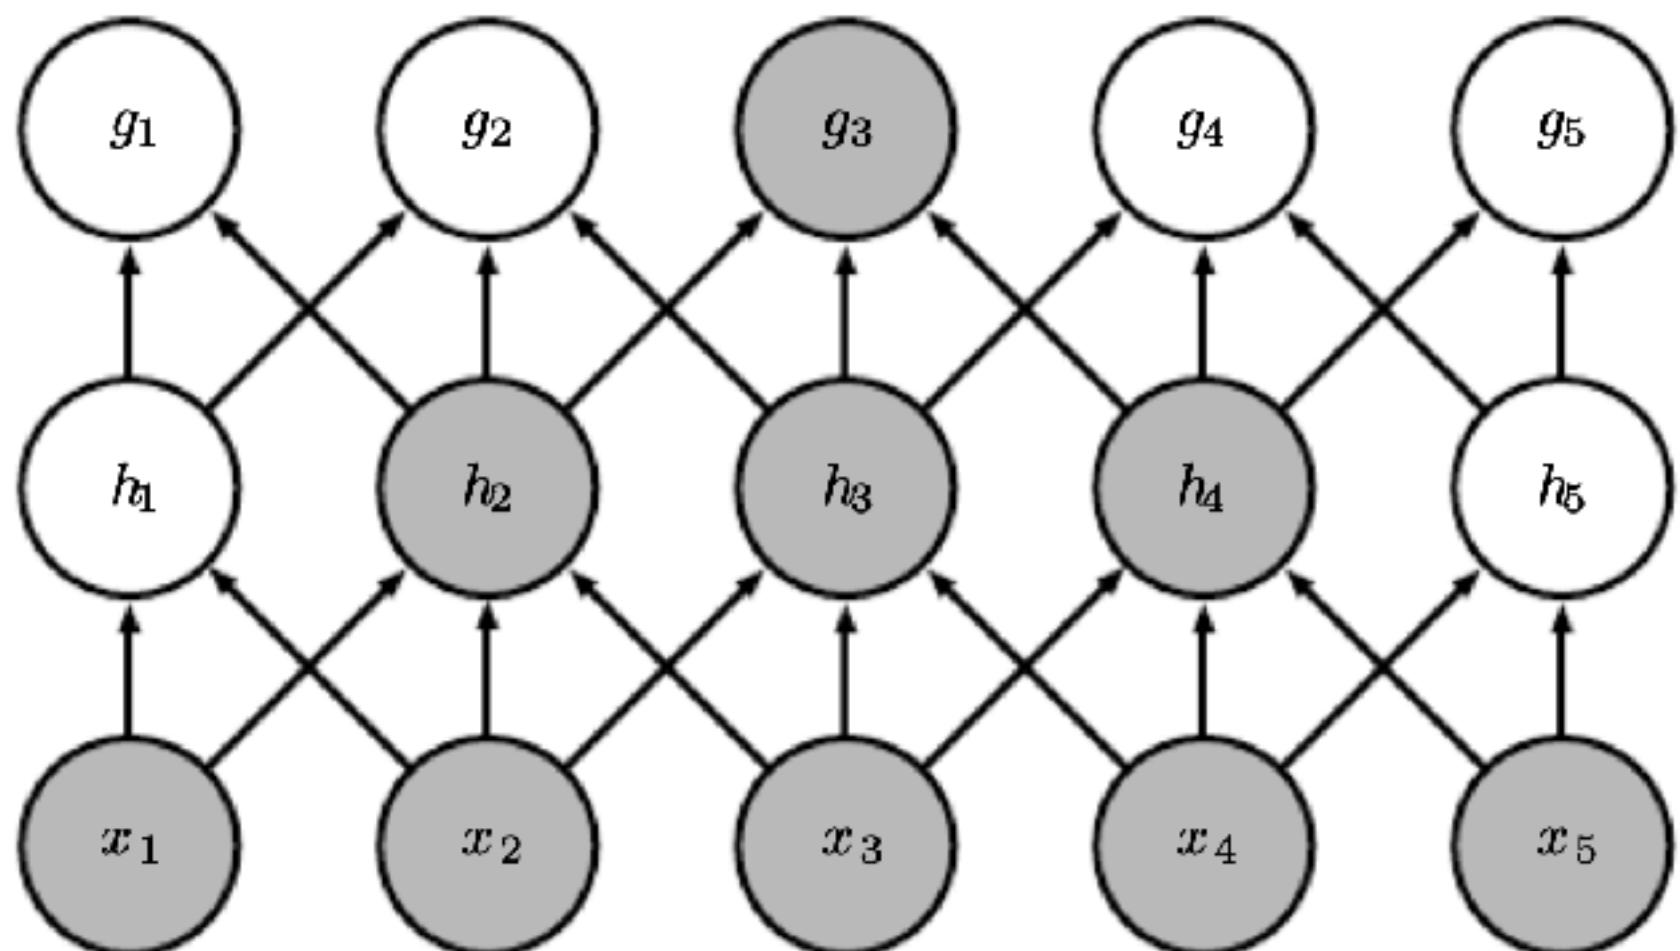


# Local Connectivity / Convolution

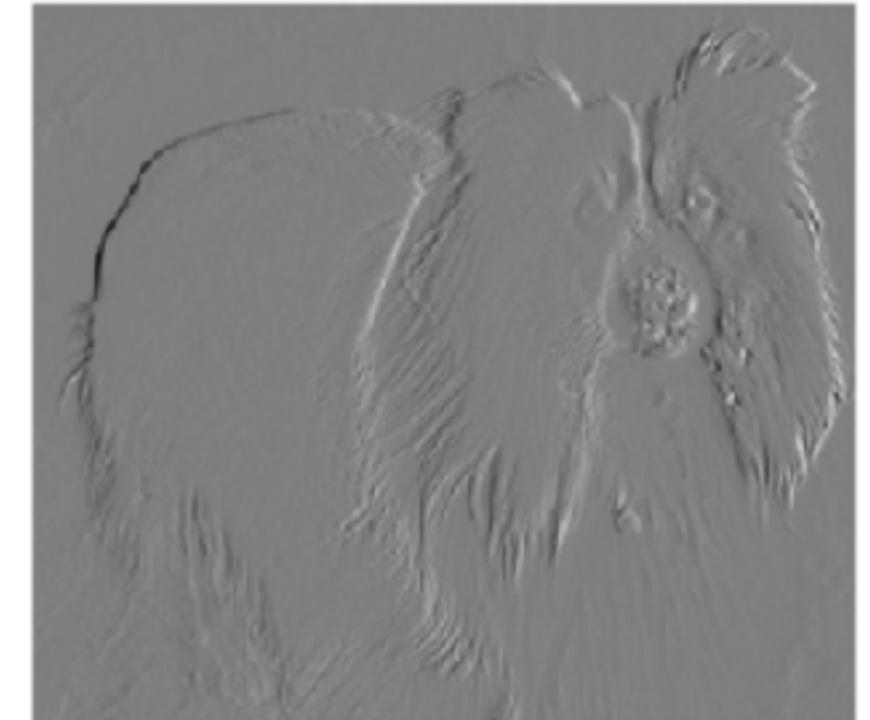
Each 1<sup>st</sup> layer neuron gets input from a local patch in the input image:



Higher layer neurons have larger “receptive fields”:



**Weight sharing:** All neurons in a feature map share the same weight values.



Example: Edge detection feature

# Local Connectivity / Convolution

**Weight sharing:** All neurons in a feature map share the same weight values.

This is equivalent to a 2D convolution:

$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i - m, j - n) K(m, n)$$

$S(i, j)$  : Output of convolution at position  $(i, j)$

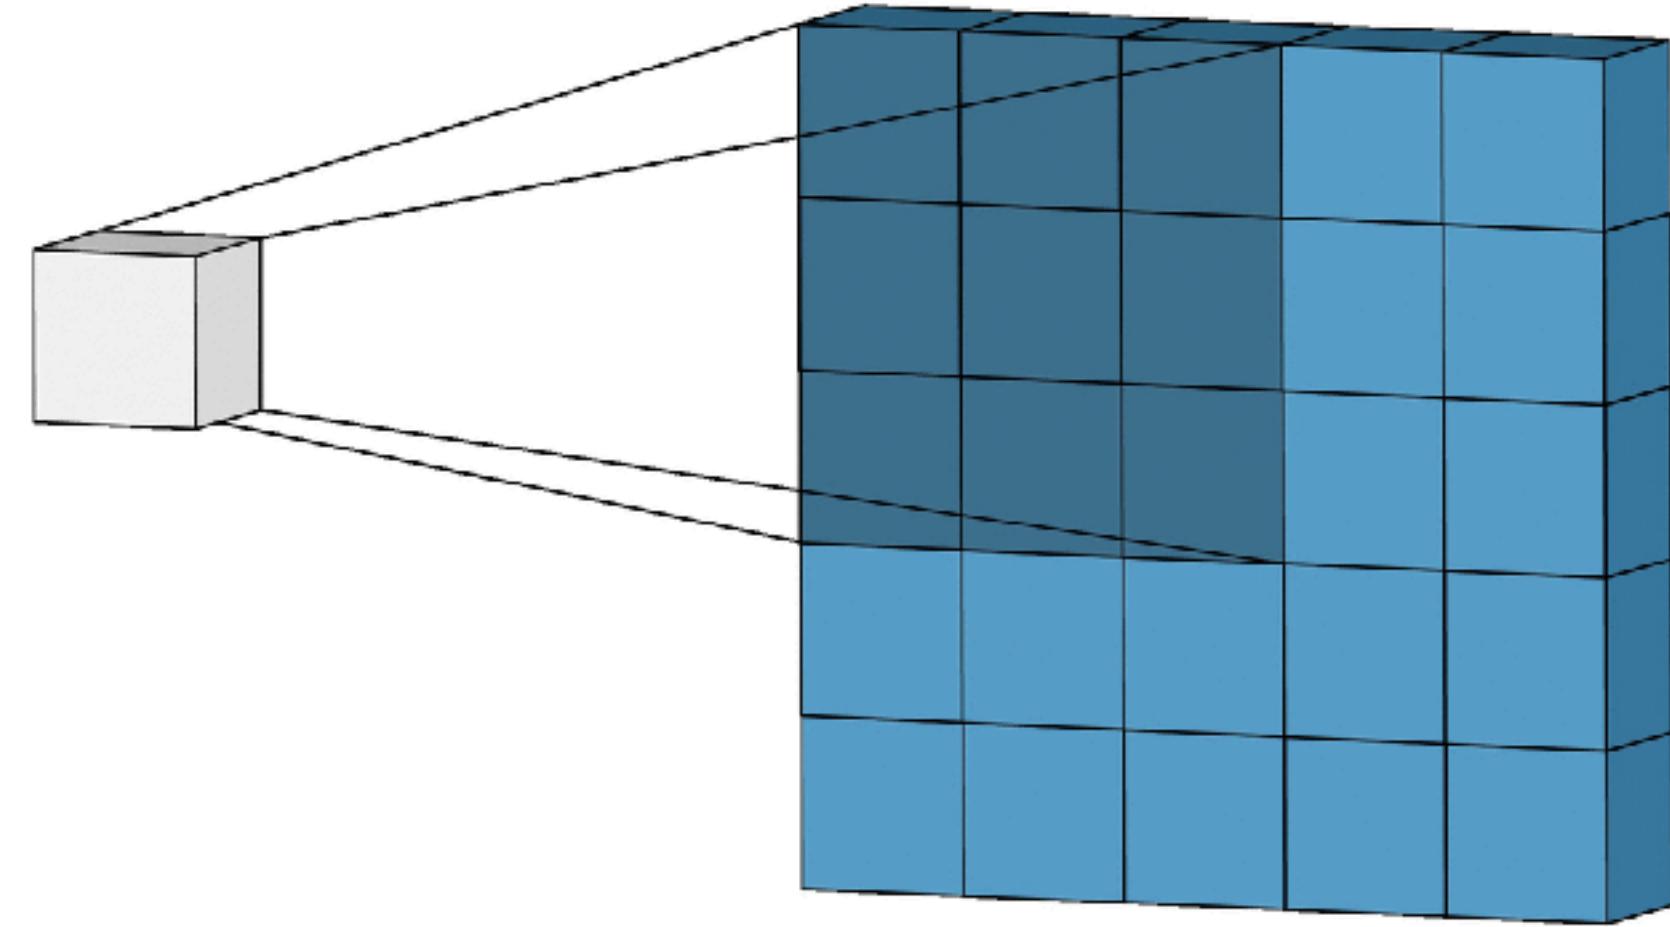
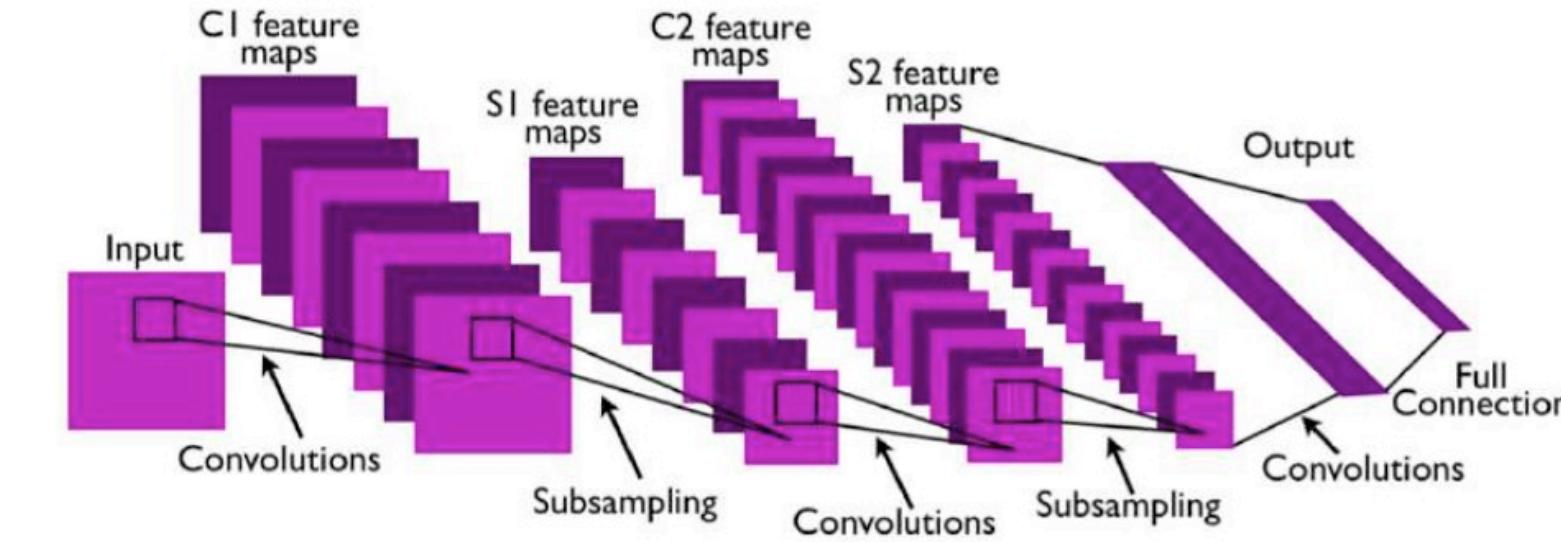
[activation of neuron in feature map]

$I(x, y)$  : Pixel-value of image at position  $(x, y)$

$K(m, n)$  : Kernel-value at kernel-position  $(m, n)$  [K defined by weight vector]

Equivalently (by mirroring the kernel  $K$ ):

$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$

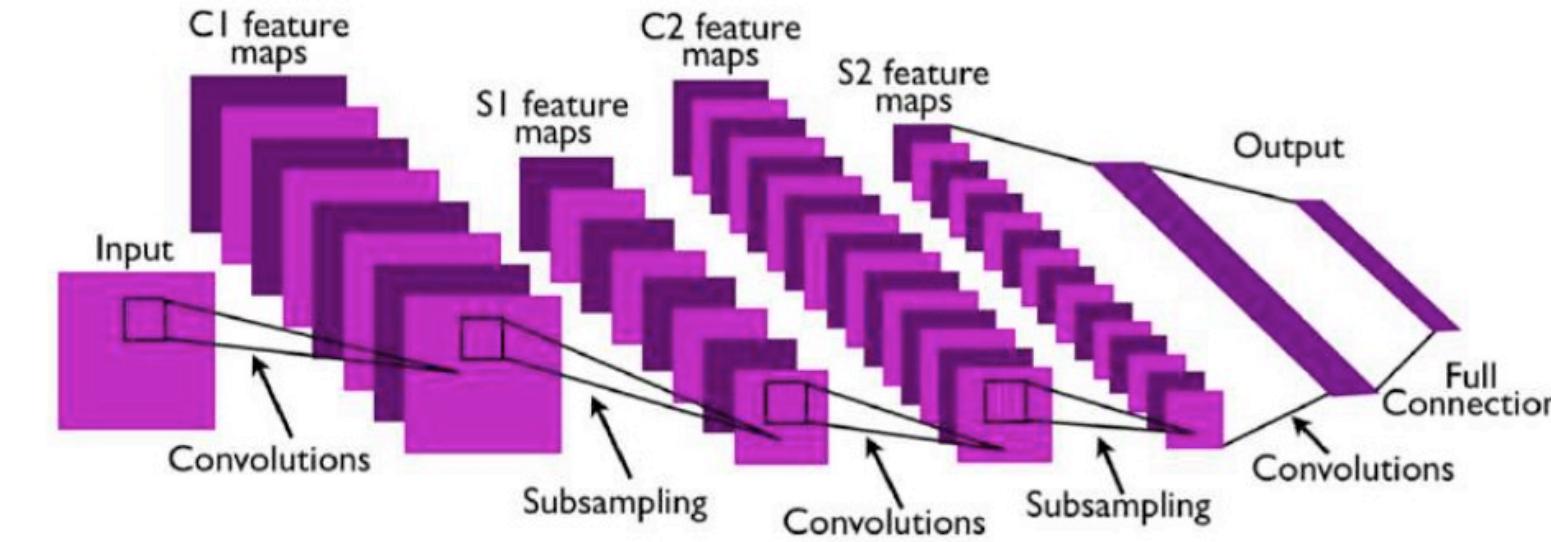
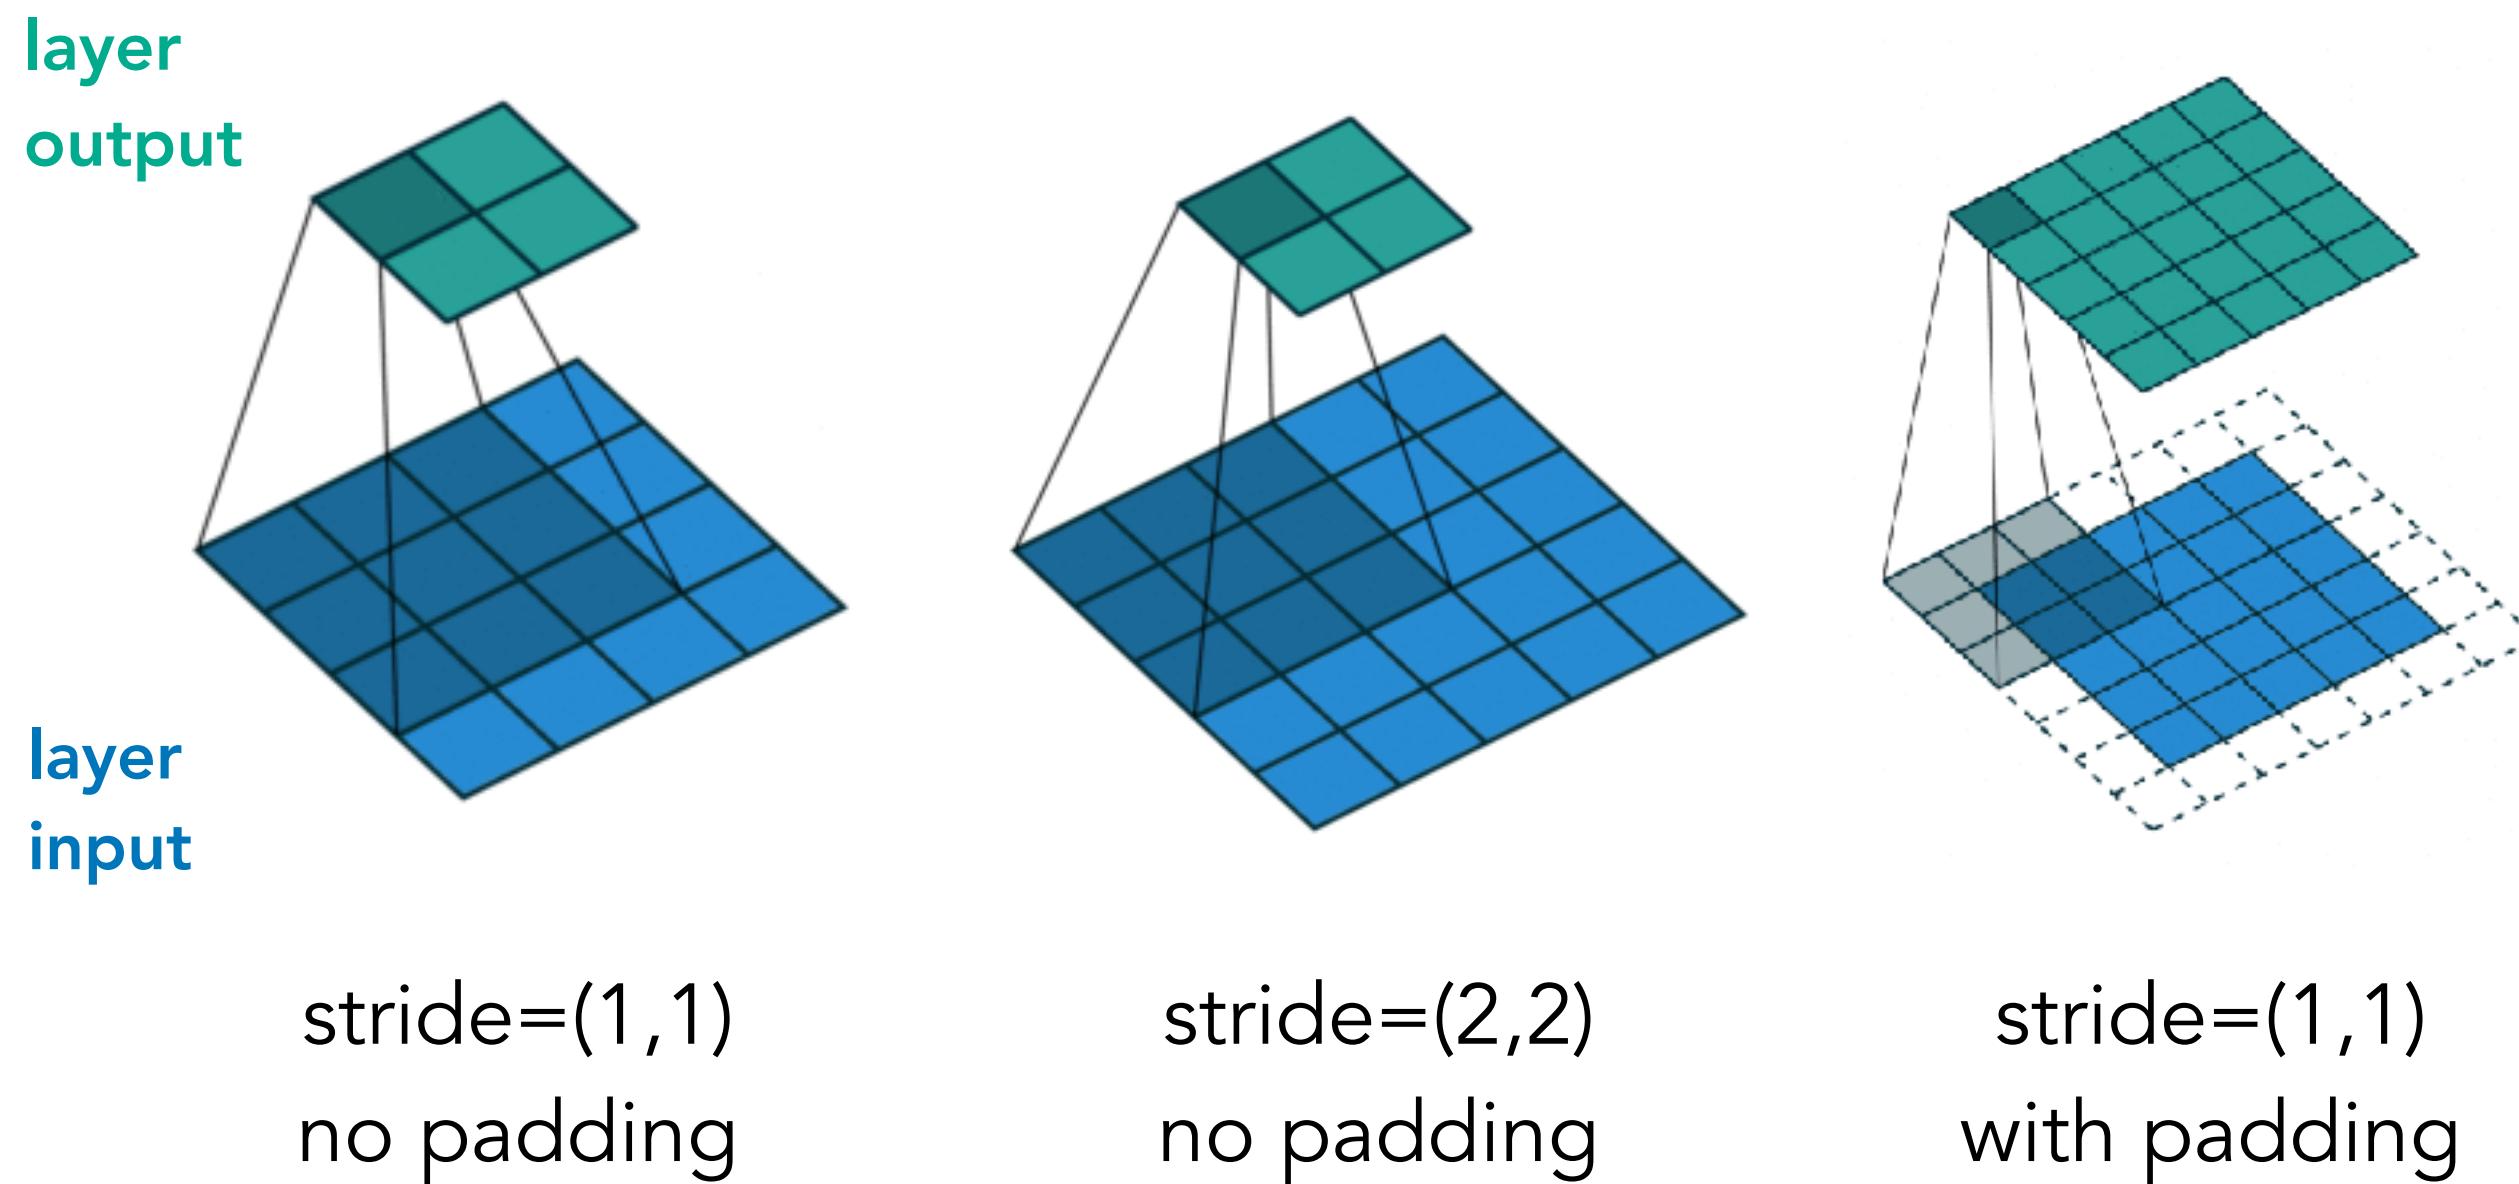


|                |                |                |   |   |
|----------------|----------------|----------------|---|---|
| 3 <sub>0</sub> | 3 <sub>1</sub> | 2 <sub>2</sub> | 1 | 0 |
| 0 <sub>2</sub> | 0 <sub>2</sub> | 1 <sub>0</sub> | 3 | 1 |
| 3 <sub>0</sub> | 1 <sub>1</sub> | 2 <sub>2</sub> | 2 | 3 |
| 2              | 0              | 0              | 2 | 2 |
| 2              | 0              | 0              | 0 | 1 |

|      |      |      |
|------|------|------|
| 12.0 | 12.0 | 17.0 |
| 10.0 | 17.0 | 19.0 |
| 9.0  | 6.0  | 14.0 |

# Convolution Operation in 2D

Examples with a **3x3 convolution kernel**:

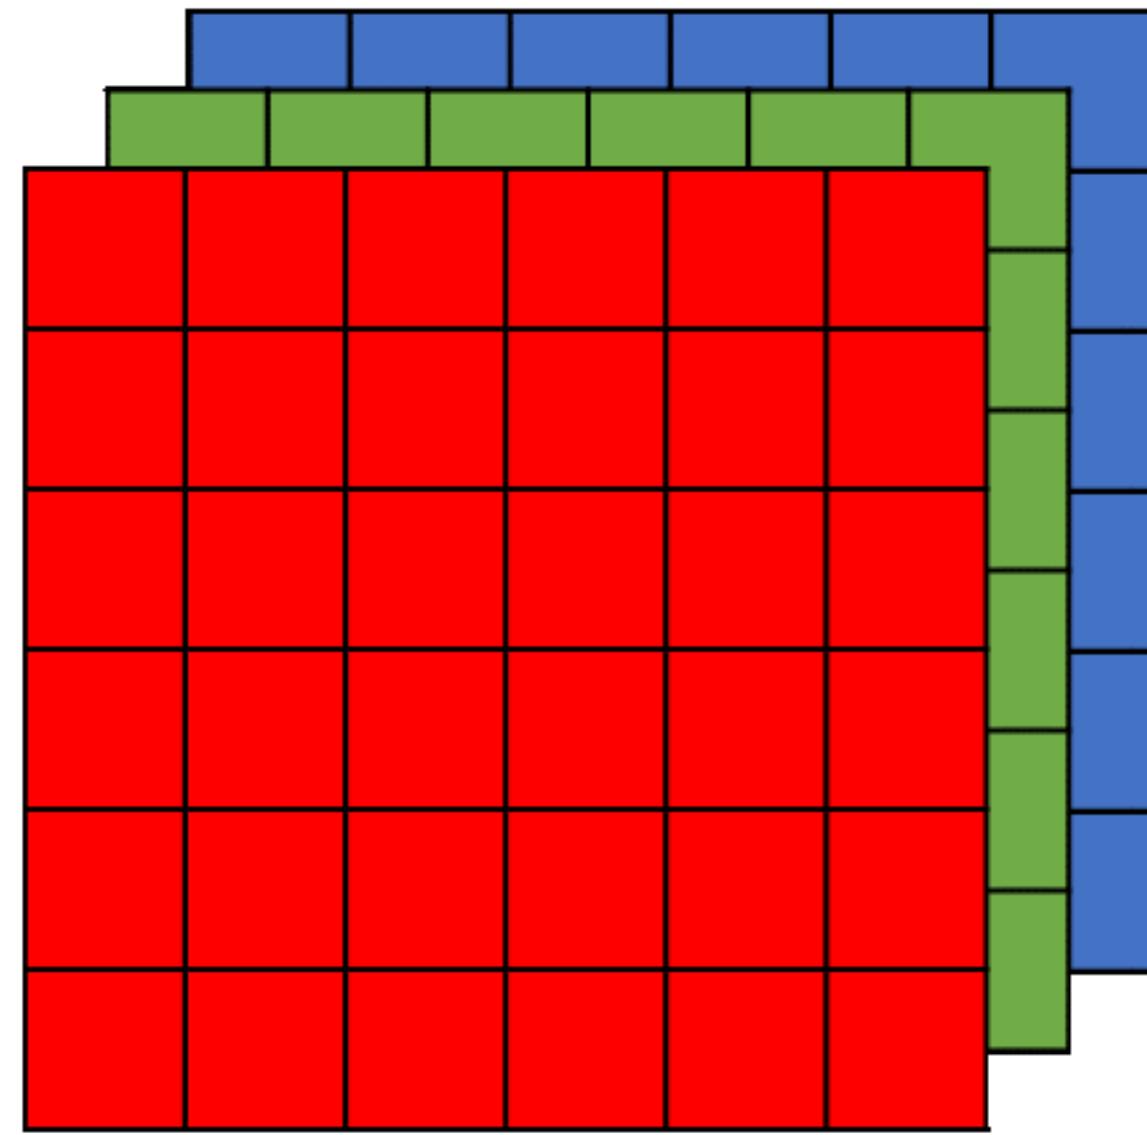


Animations from: Dumoulin & Visin (2018), "A guide to convolution arithmetic for deep learning".

# Convolution Operation in 2D

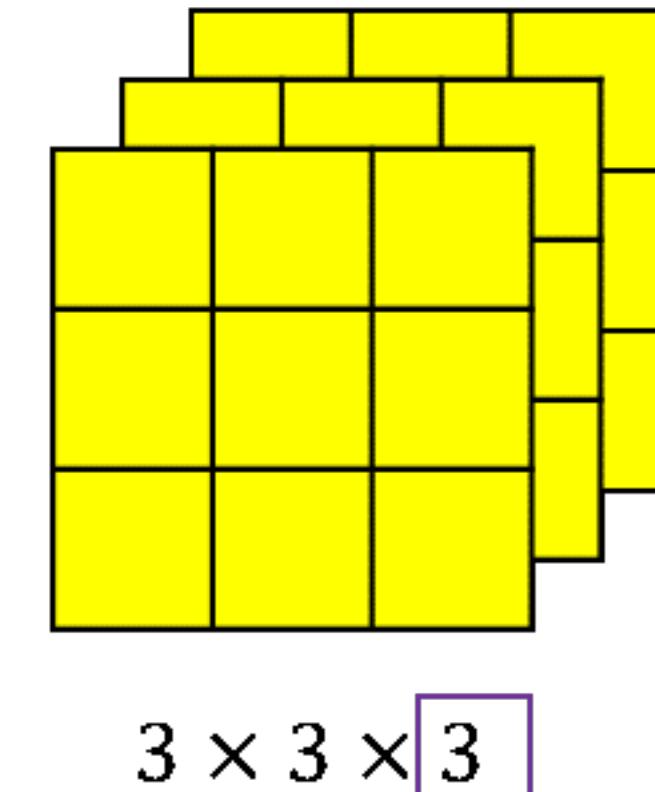
**Convolution over several channels:**

- e.g. RGB input
- Convolutional filter consists of one kernel per channel
- Sum over all channels



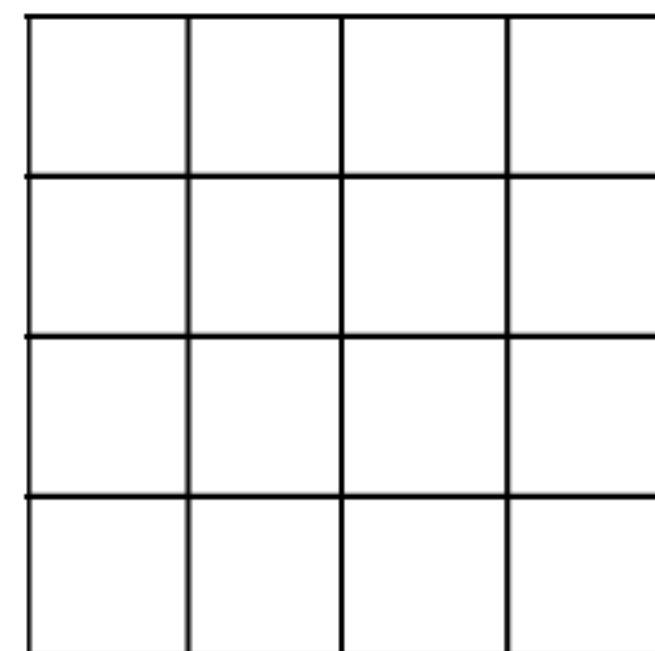
$6 \times 6 \times 3$

\*

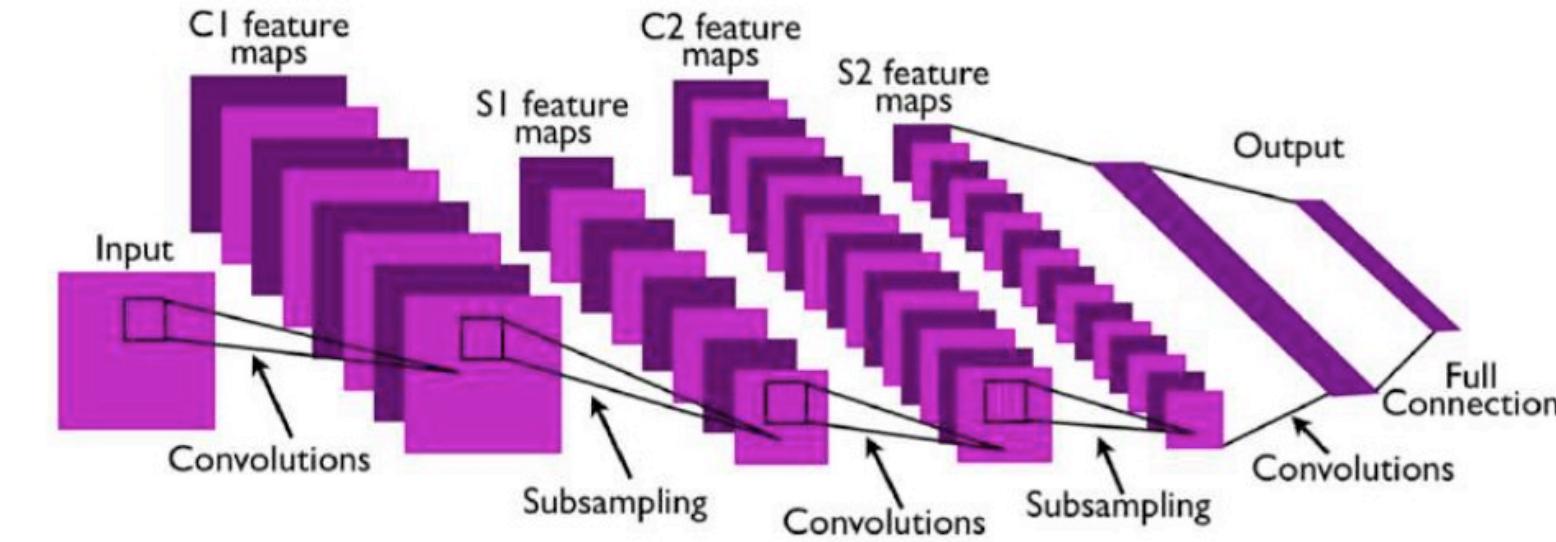


$3 \times 3 \times 3$

=



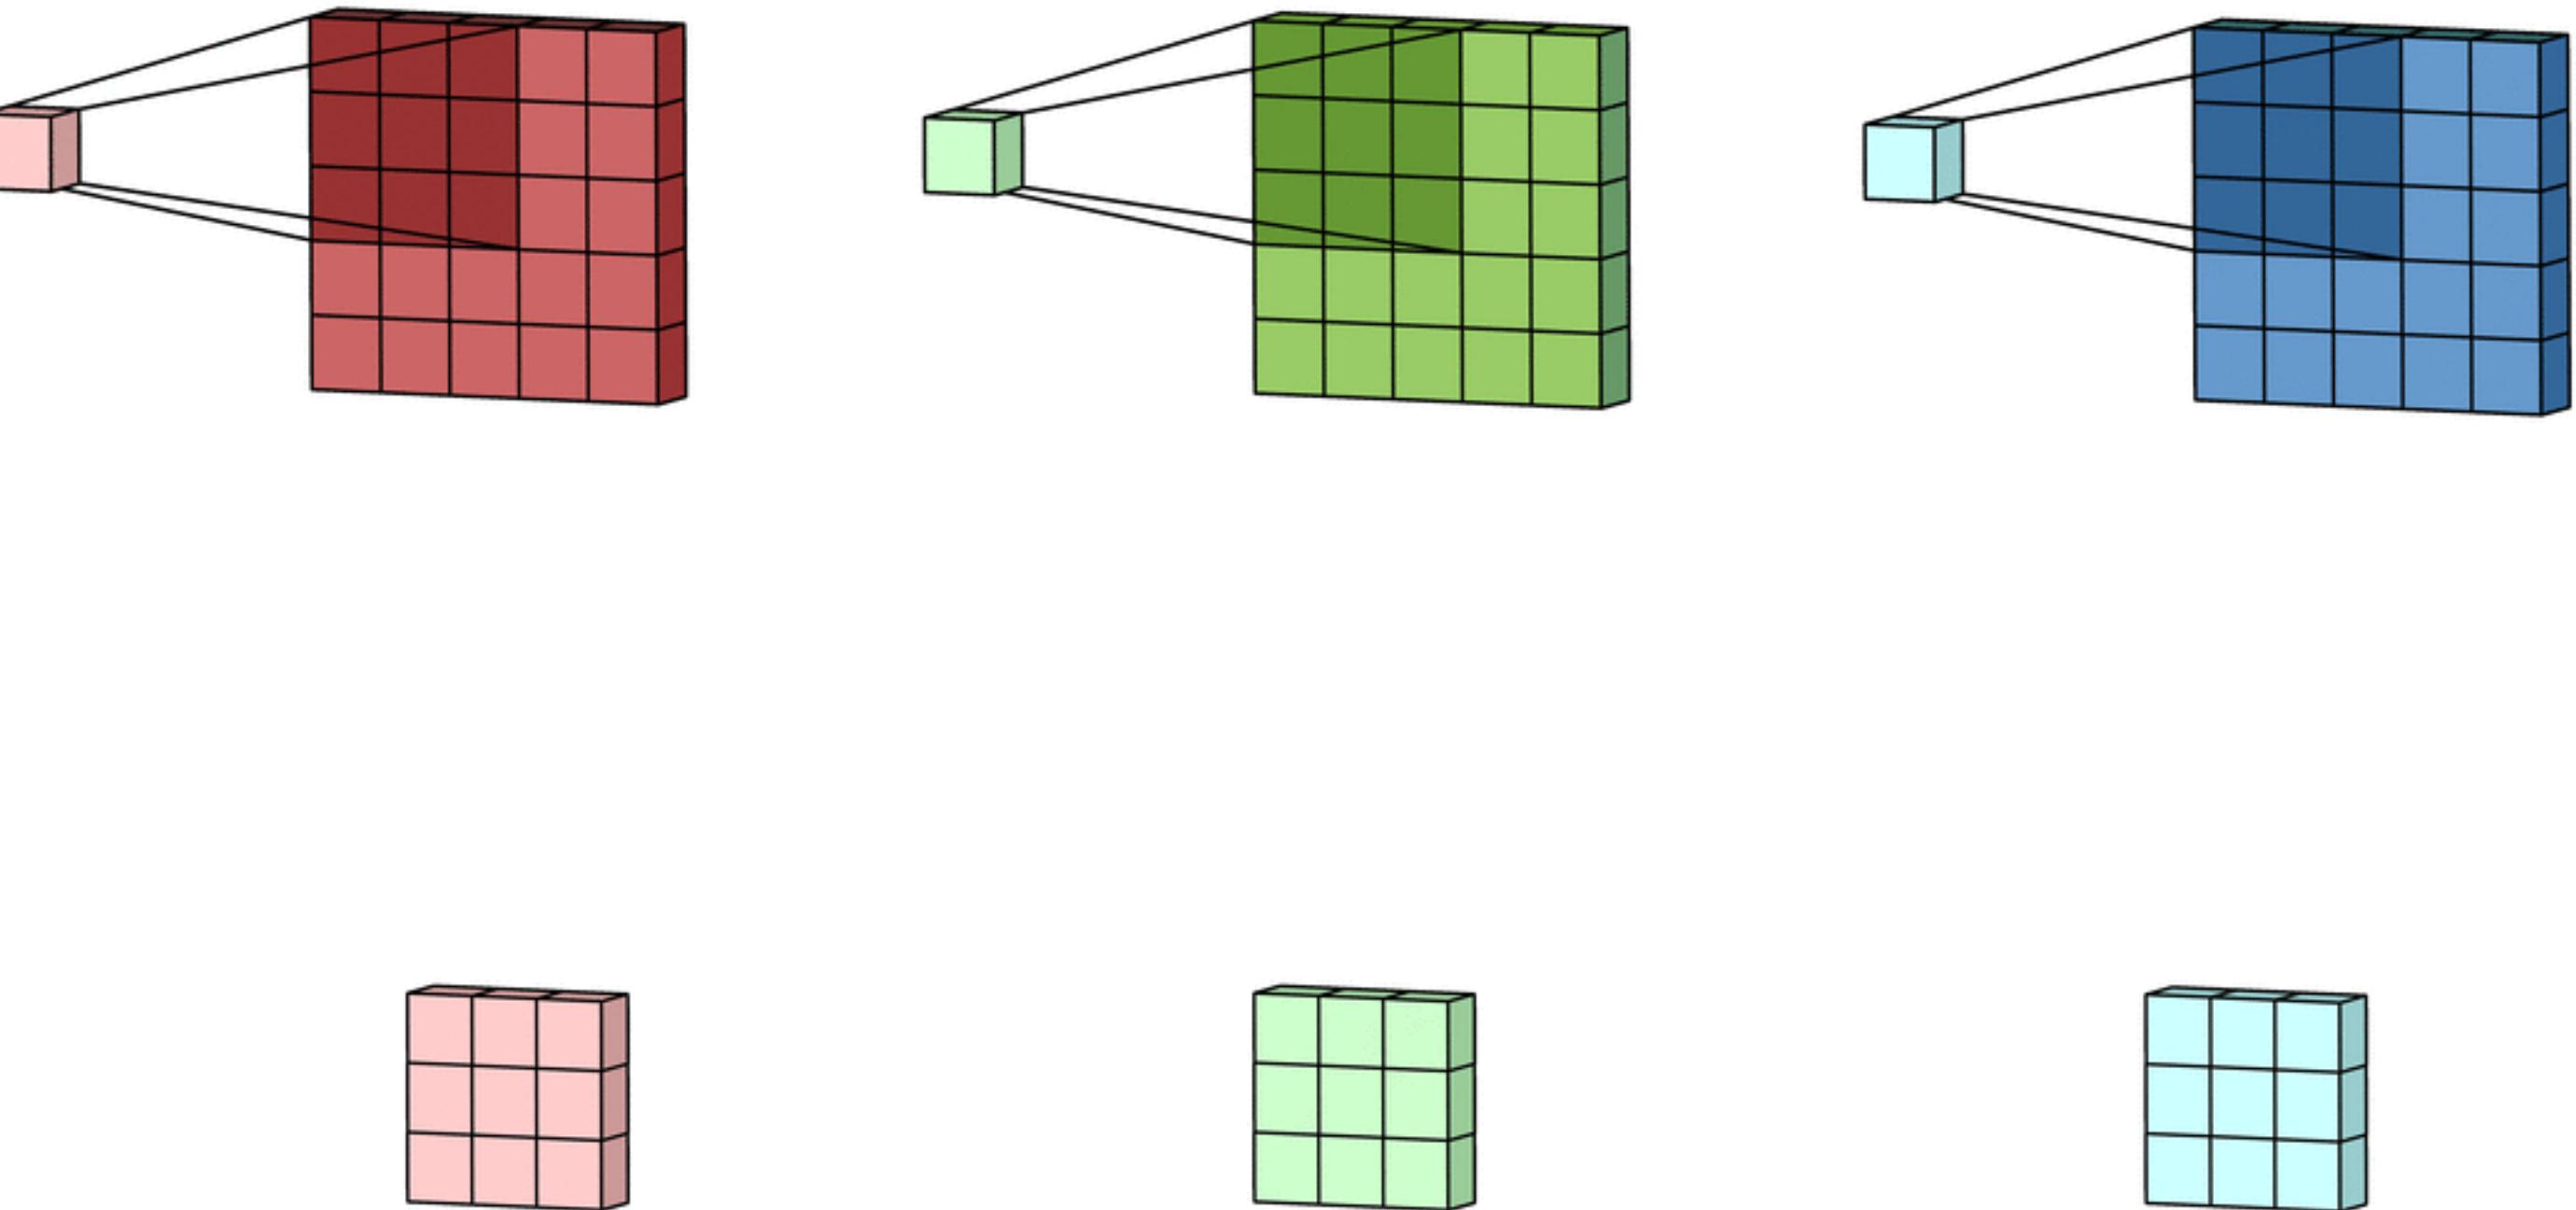
$4 \times 4$



# Convolution Operation in 2D

## Convolution over several channels:

- e.g. RGB input
- Convolutional filter consists of one kernel per channel
- Sum over all channels



# Pooling

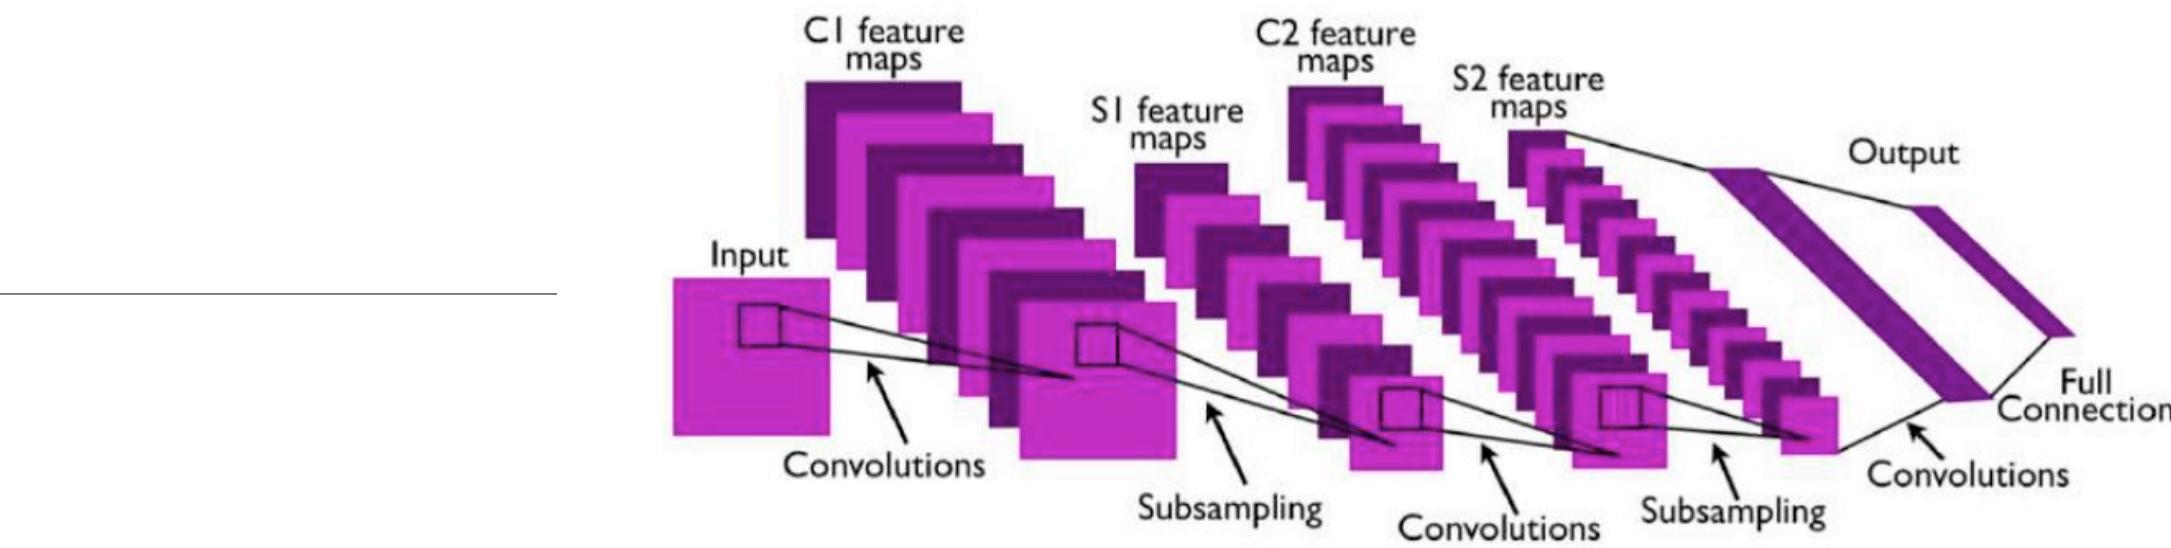
Compute simple statistics over a local neighborhood.

Some pooling functions:

- **Max pooling:** output the max of the inputs.
- **Average pooling:** output the mean of the inputs.

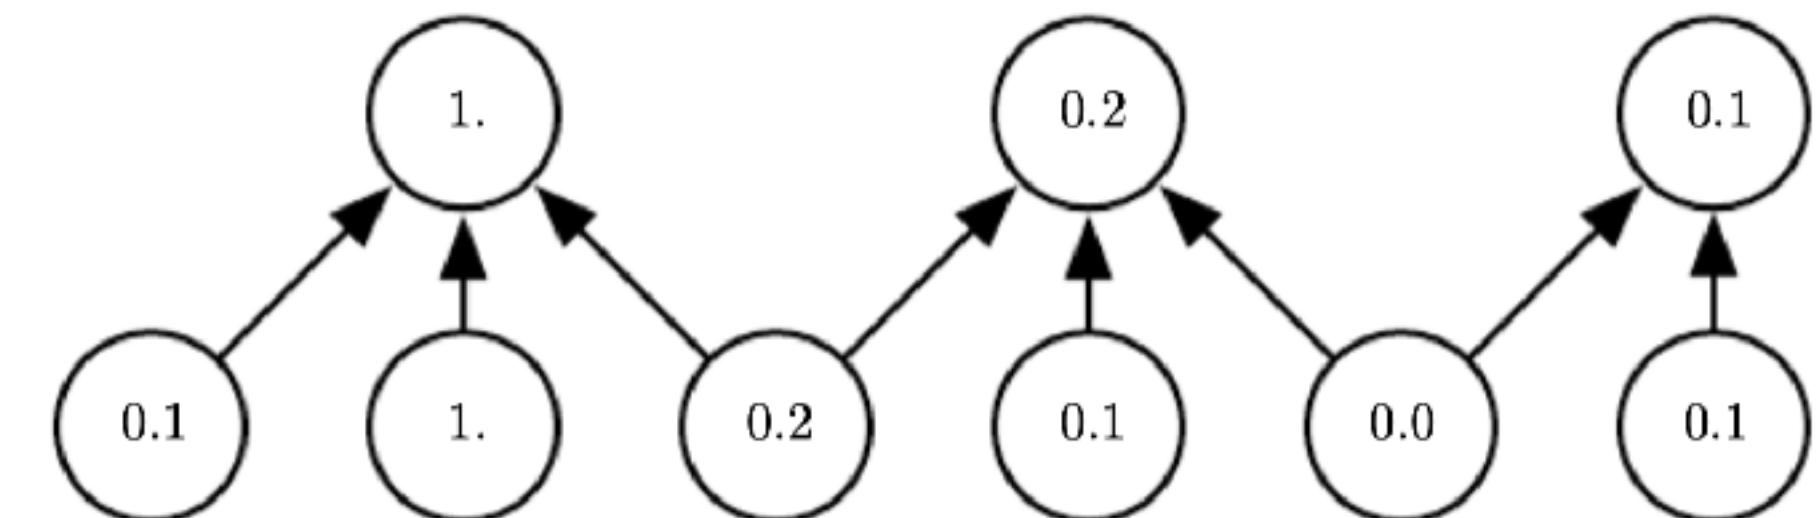
**Application of pooling:**

- One pooling layer per feature map.
- Stride of  $k$ : Move pooling window by  $k$  pixels when “moved” over the image.
- Leads to a subsampling of the feature map (reduction of feature map size)

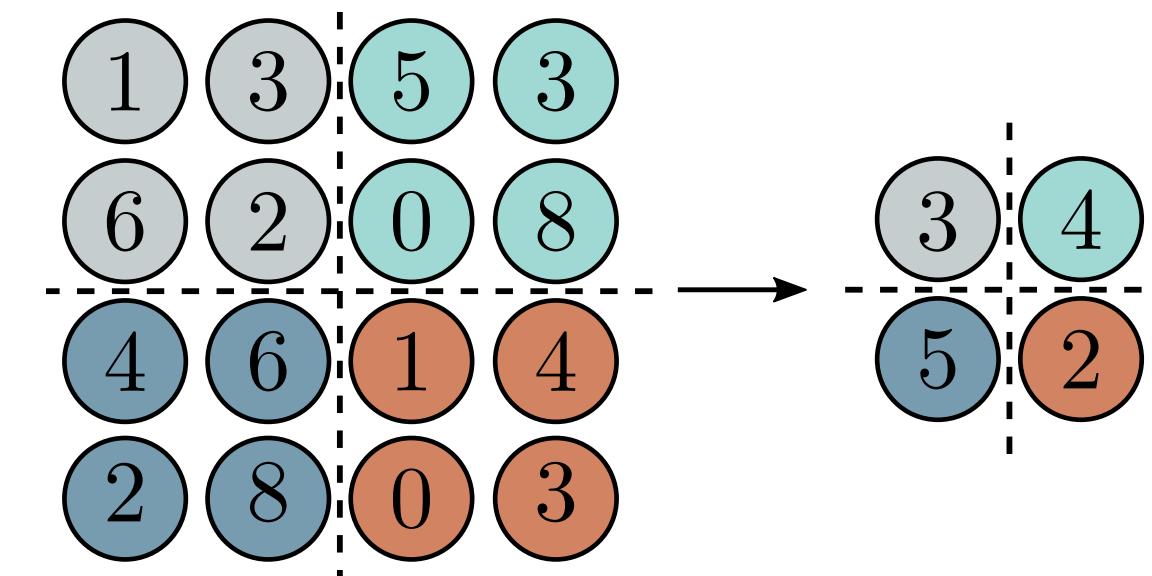


**Examples:**

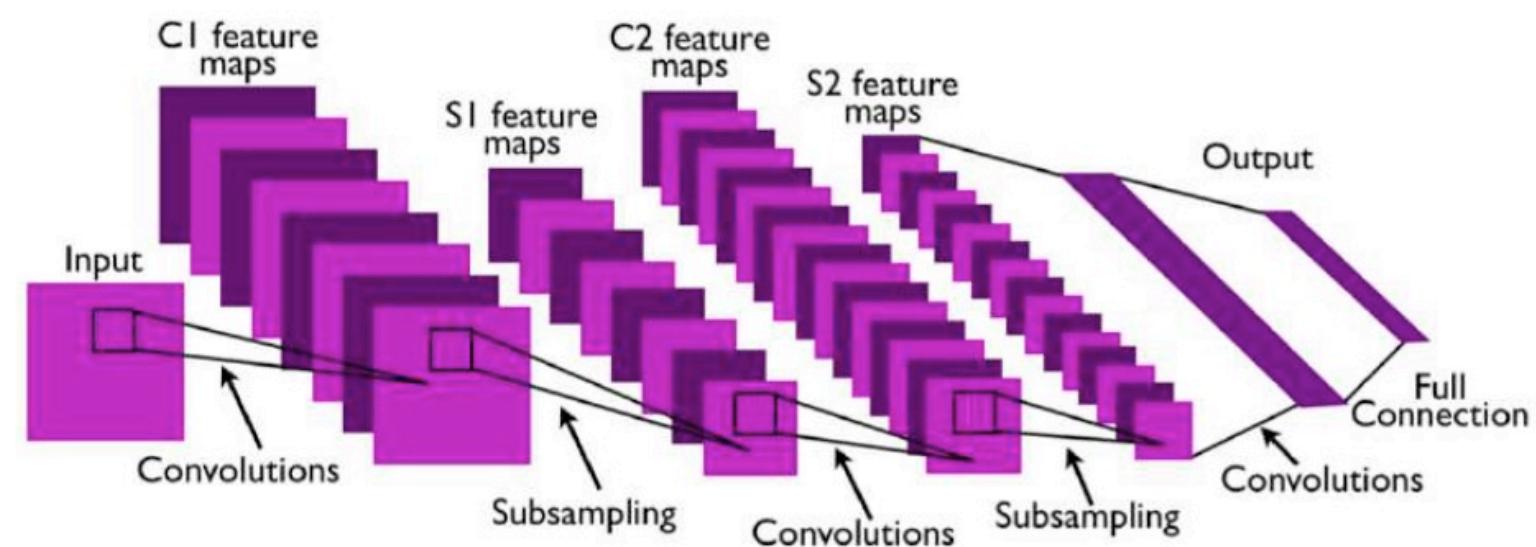
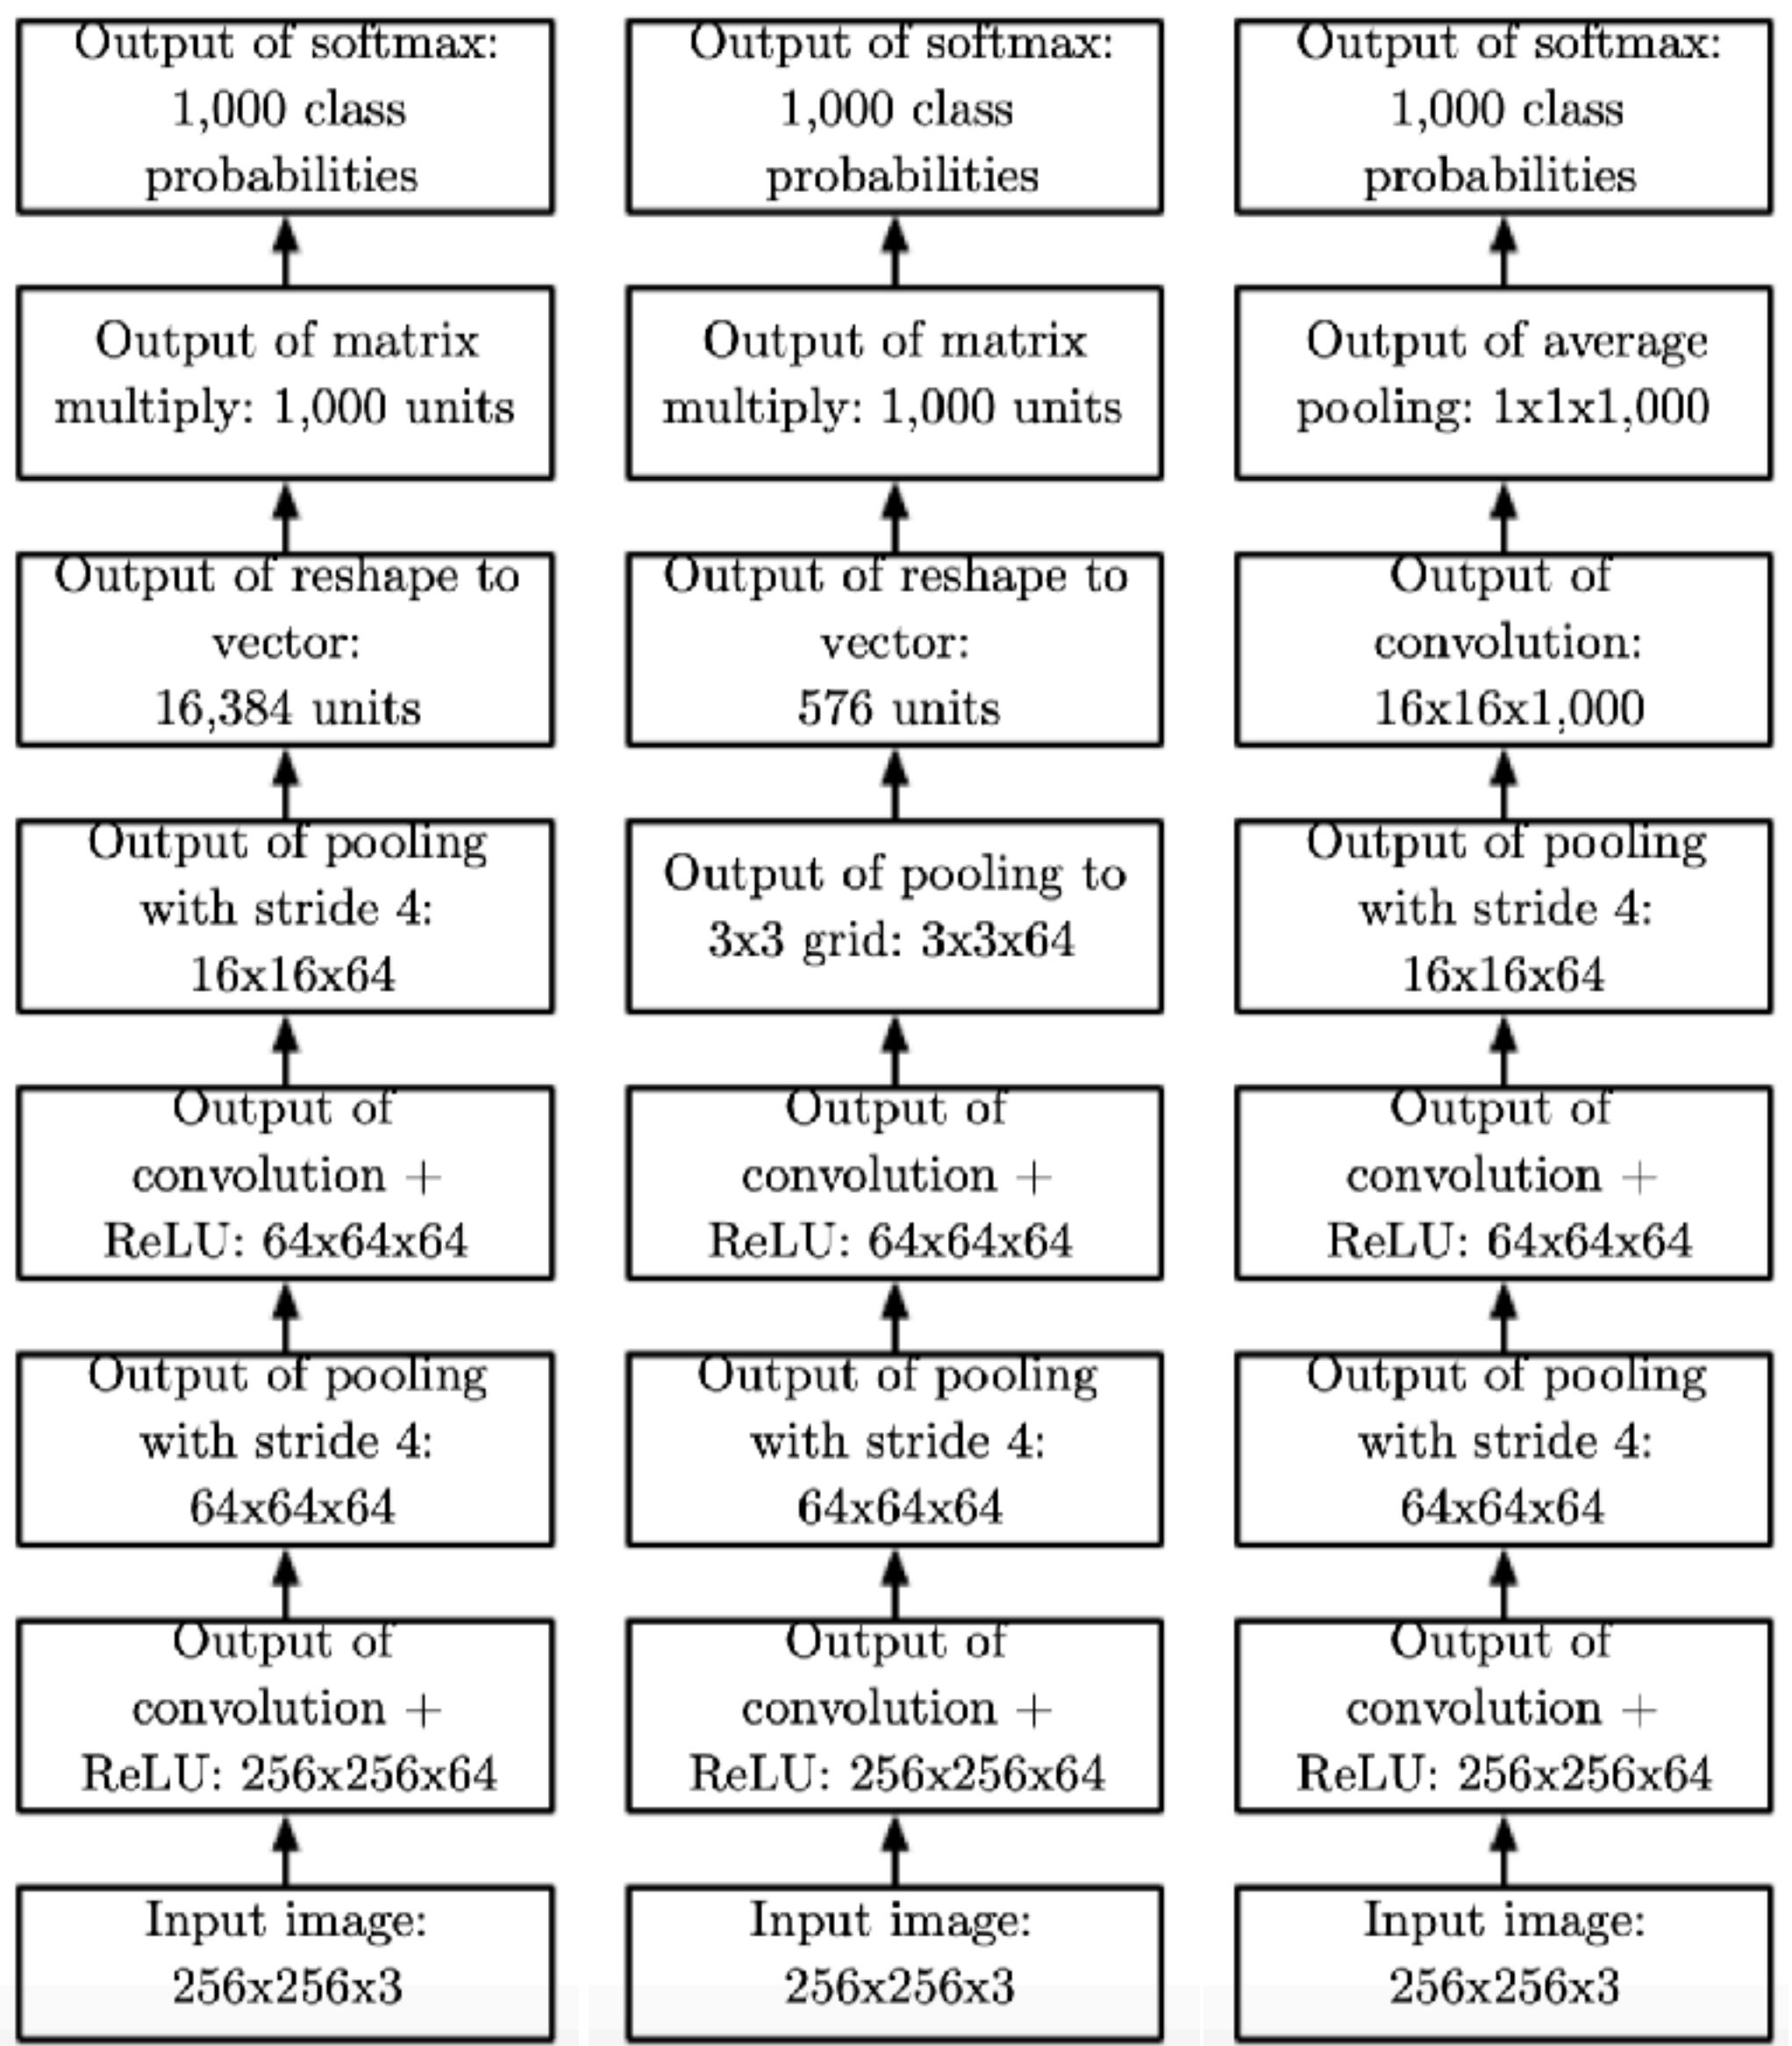
- Max-pooling (1D):  
pooling width: 3, stride: 2



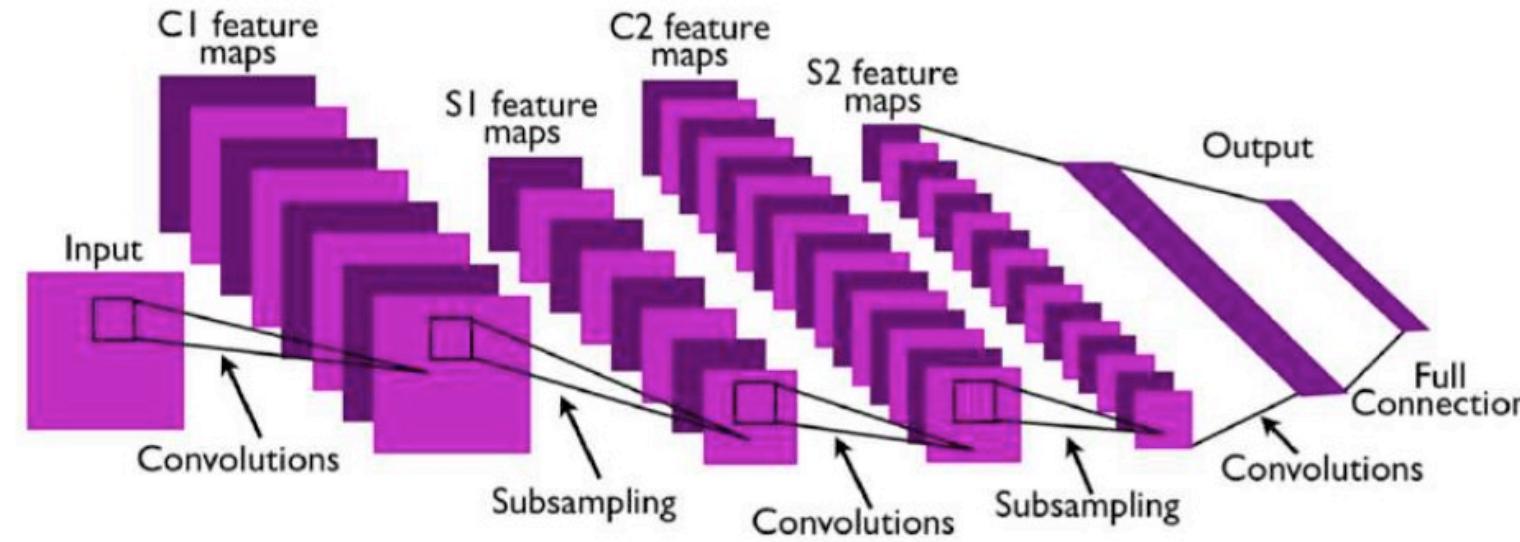
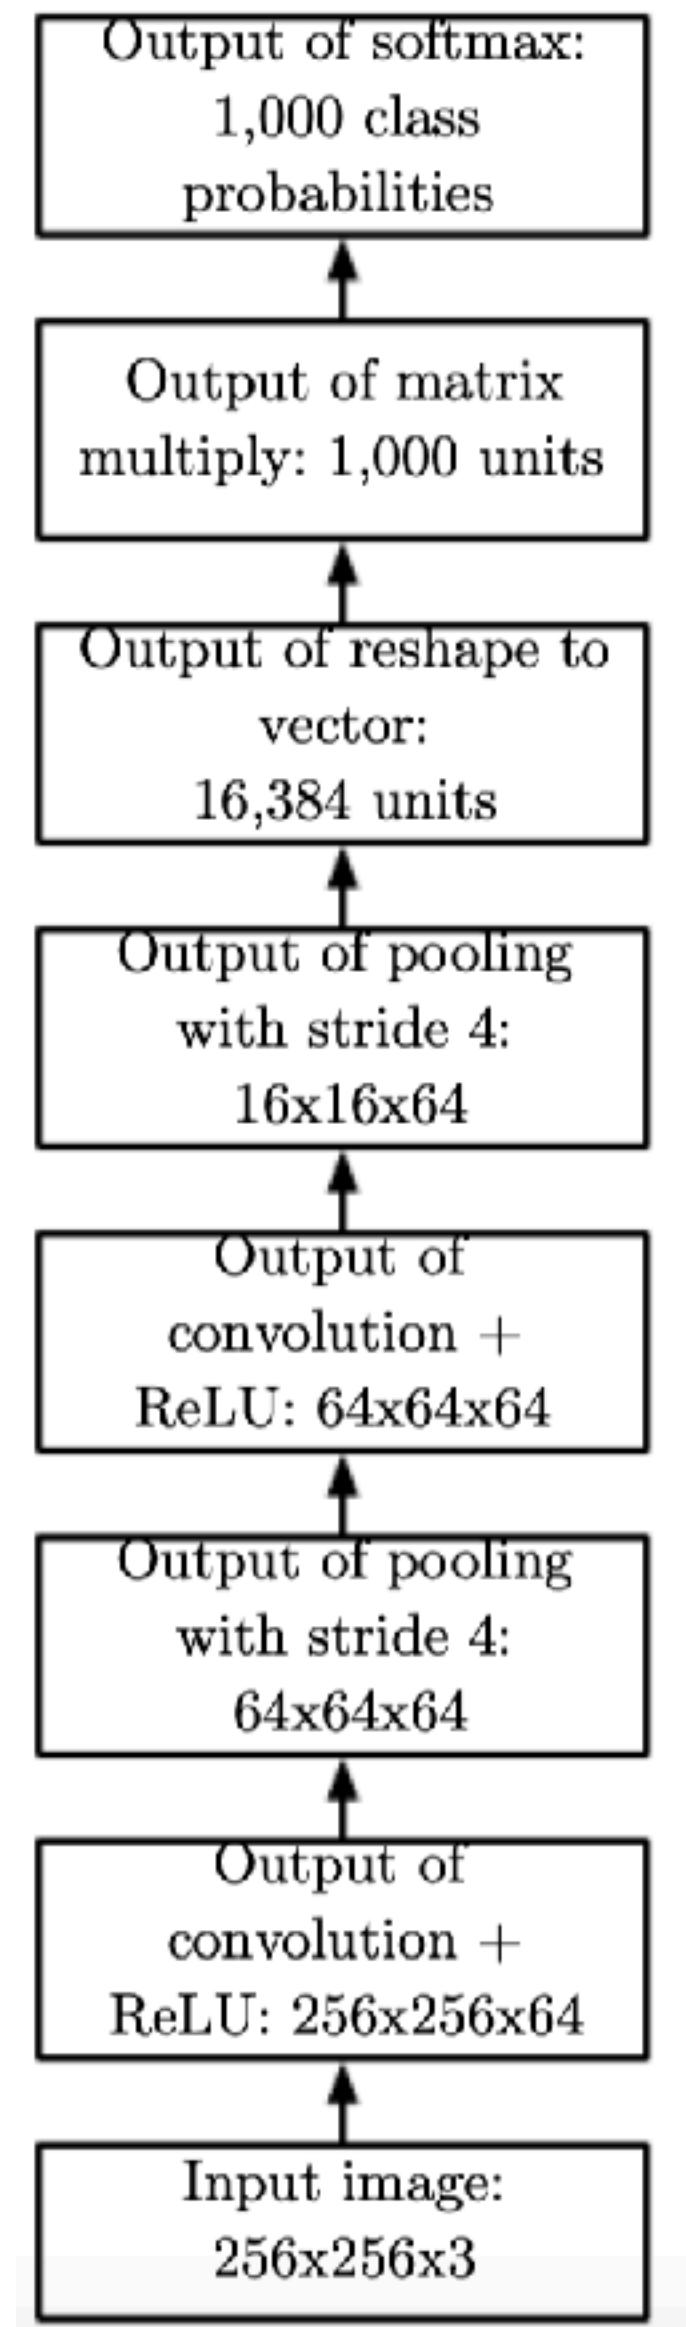
- Avg-pooling (2D):  
pooling size: (2,2), stride: (2,2)



# Illustrative Example Architectures



# Illustrative Example Architecture



## Conclusions from CNNs:

### Advantages:

- ▶ Effective in avoiding overfitting.
- ▶ # of parameters drastically decreased in comparison to fully-connected nets.
- ▶ State-of-the-art in image classification.

*(coming up next!)*

### Disadvantages:

- ▶ Only applicable to data with invariance properties (e.g., image data).

# Today

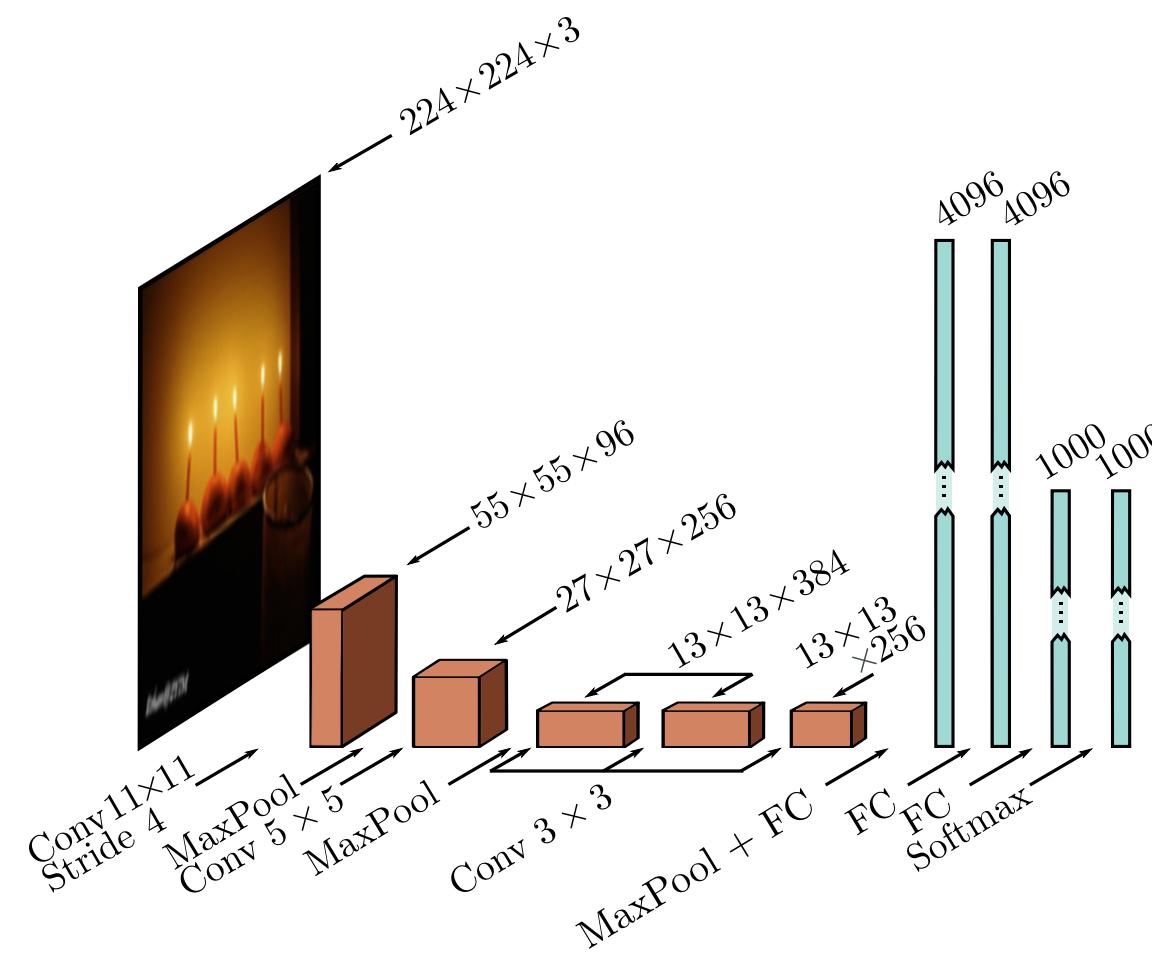
---

- Convolutional Neural Networks
- Background
- CNN Architecture
- Going Deeper with CNNs
- Residual Networks (ResNets)

# Going deeper with neural networks



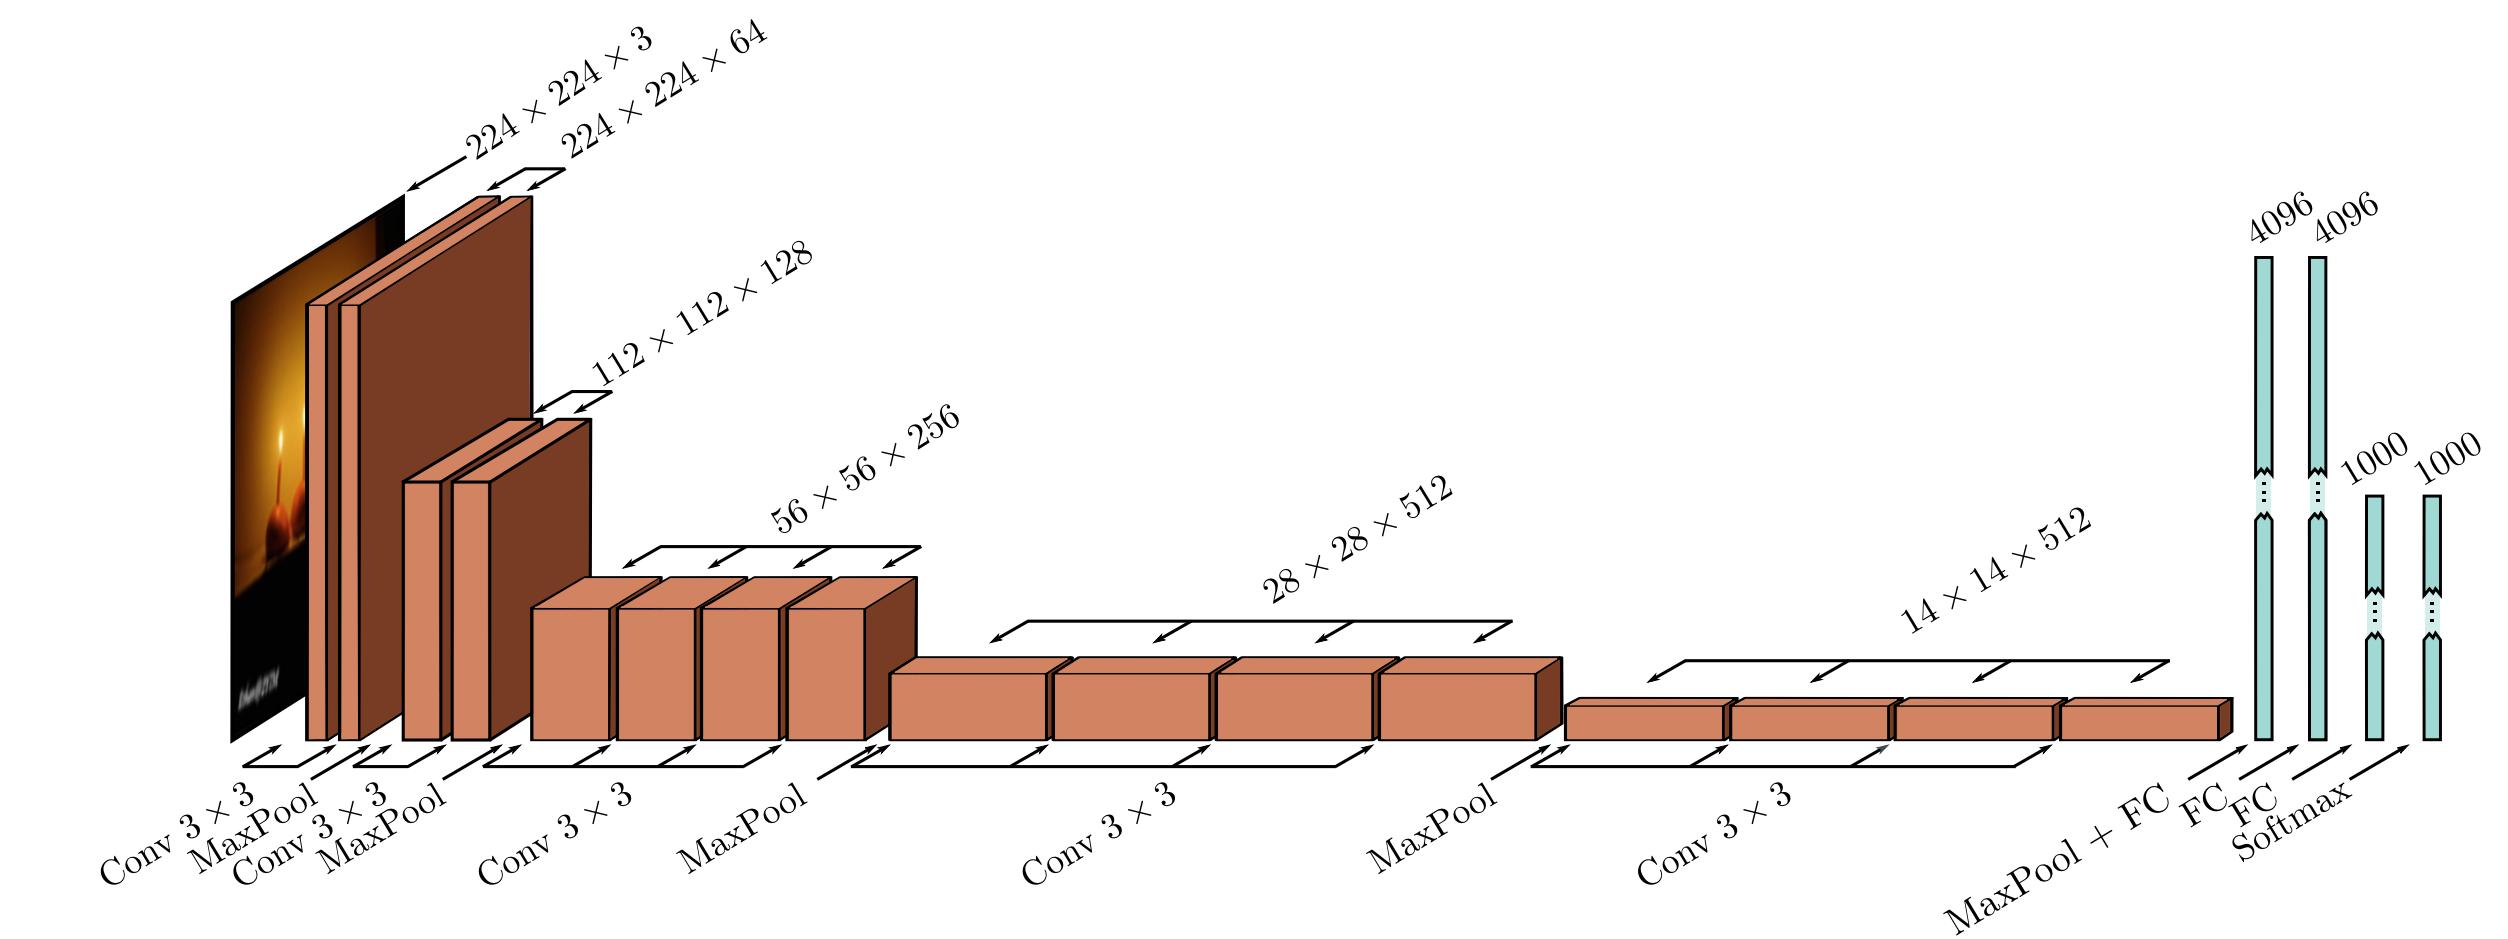
**AlexNet:** CNN with 5 convolutional layers.



[A. Krizhevsky et al. (2012): "ImageNet Classification with Deep Convolutional Neural Networks"]



**VGG:** CNN with 19 layers.



[K. Simonyan, and A. Zisserman (2014): "Very deep convolutional networks for large-scale image recognition"]

# Going deeper with neural networks

---

## Observations:

- Very deep networks often show worse performance, even in terms of the **training error**.
- This is surprising, since the top layers could simply compute an identity mapping.
  - Hence, performance should be at least as good as a shallower network.

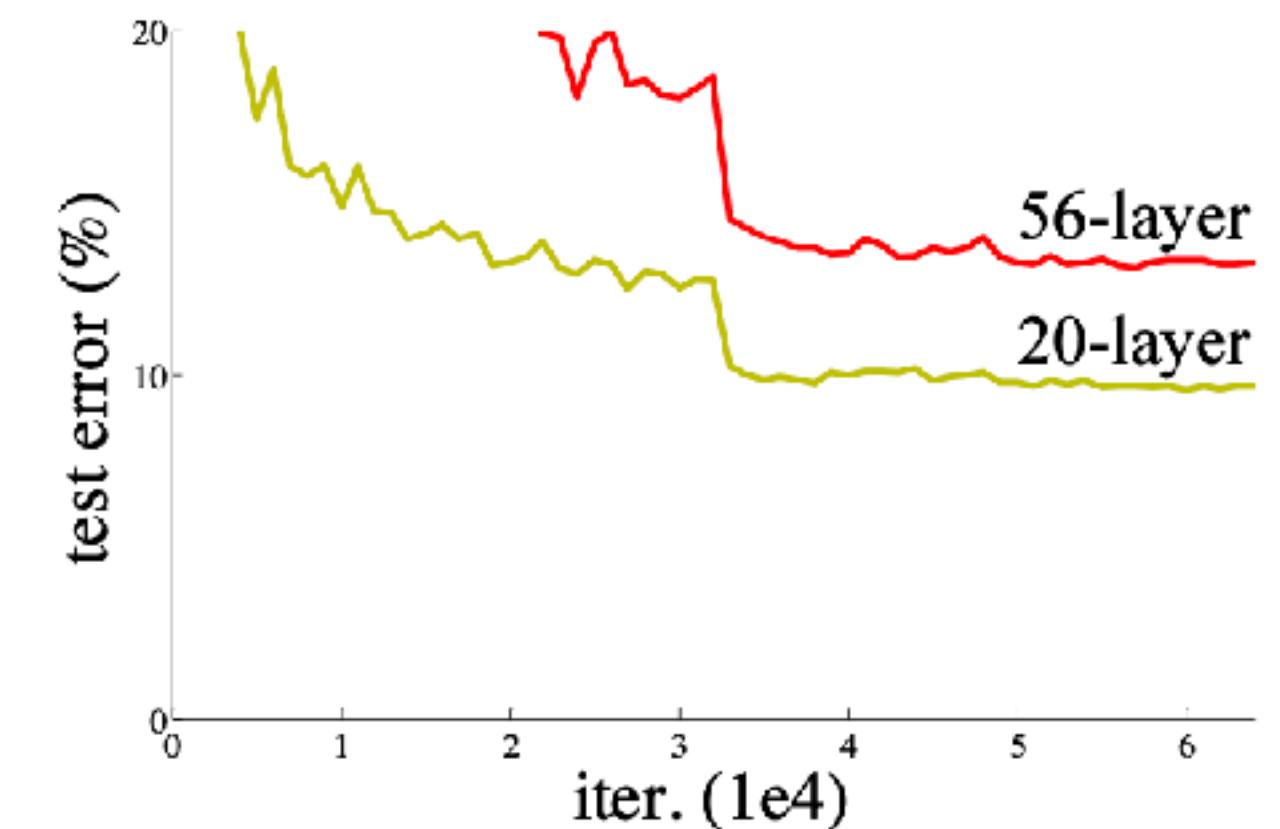
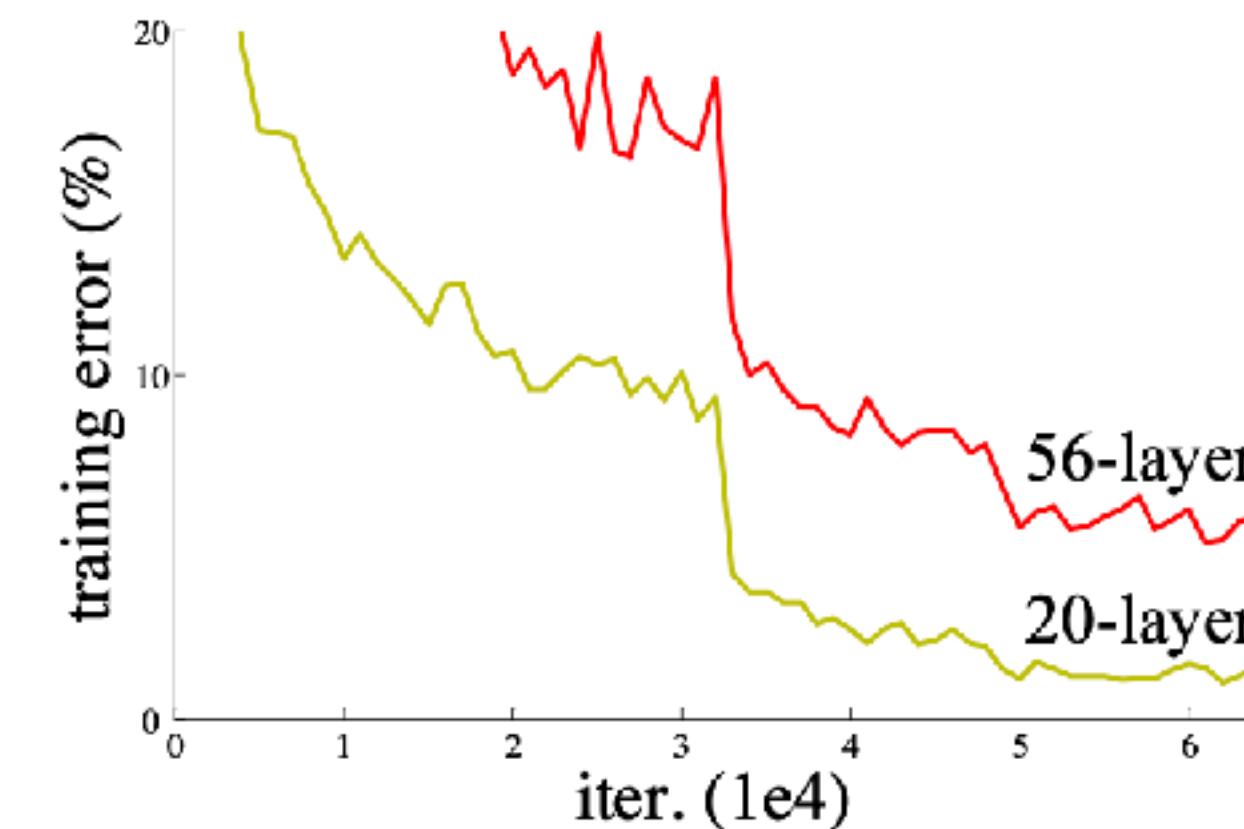
# Deep networks are hard to optimize

## Observations:

- Very deep networks often show worse performance, even in terms of the training error.
- This is surprising, since the top layers could simply compute an identity mapping.
  - Hence, performance should be at least as good as a shallower network.

## Problem:

- Deeper networks are harder to train.
- “Vanishing gradient” problem.



## Previous solution:

- Auxiliary loss in middle layers (e.g., GoogLeNet)

[Kaiming He, et al. “Deep residual learning for image recognition”, CVPR 2016]

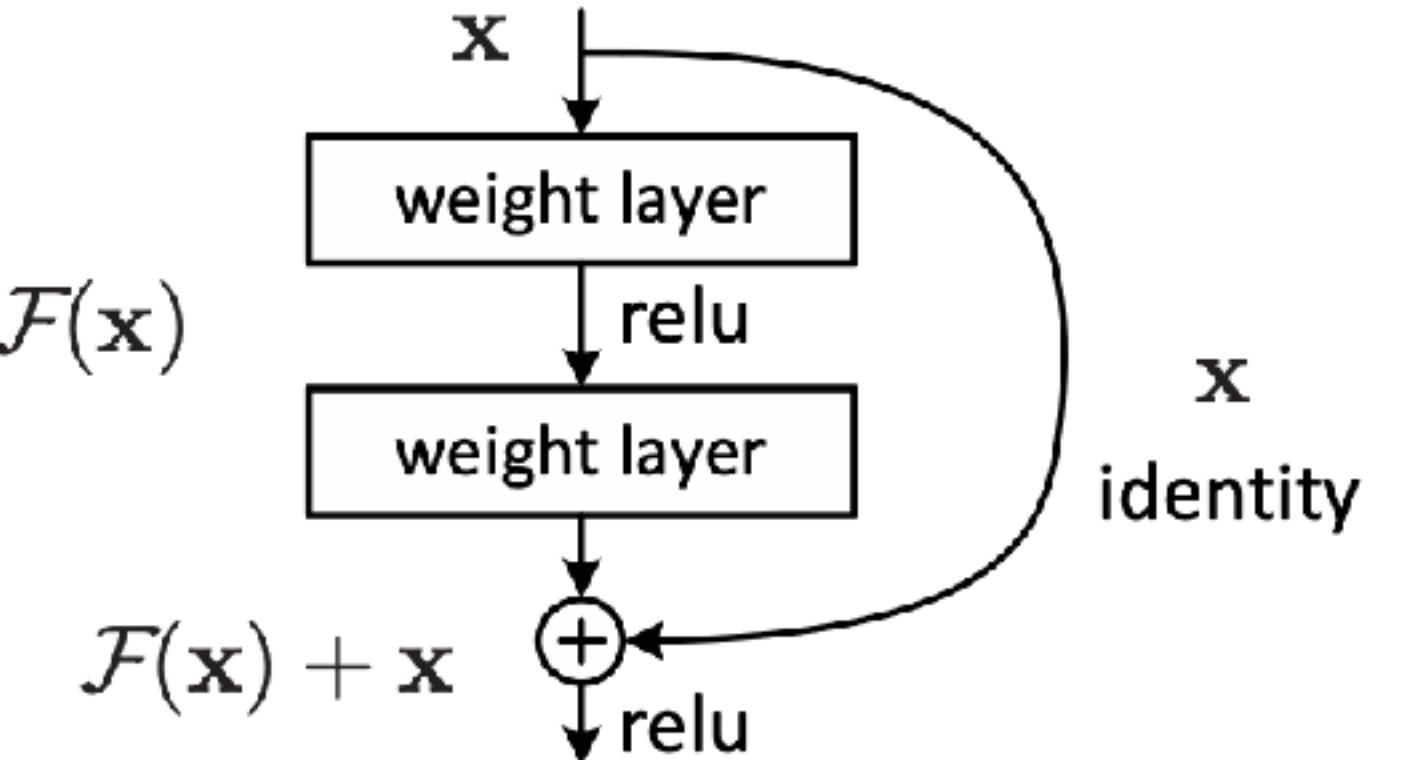
# Today

---

- Convolutional Neural Networks
  - Background
  - CNN Architecture
  - Going Deeper with CNNs
- Residual Networks (ResNets) & Extensions

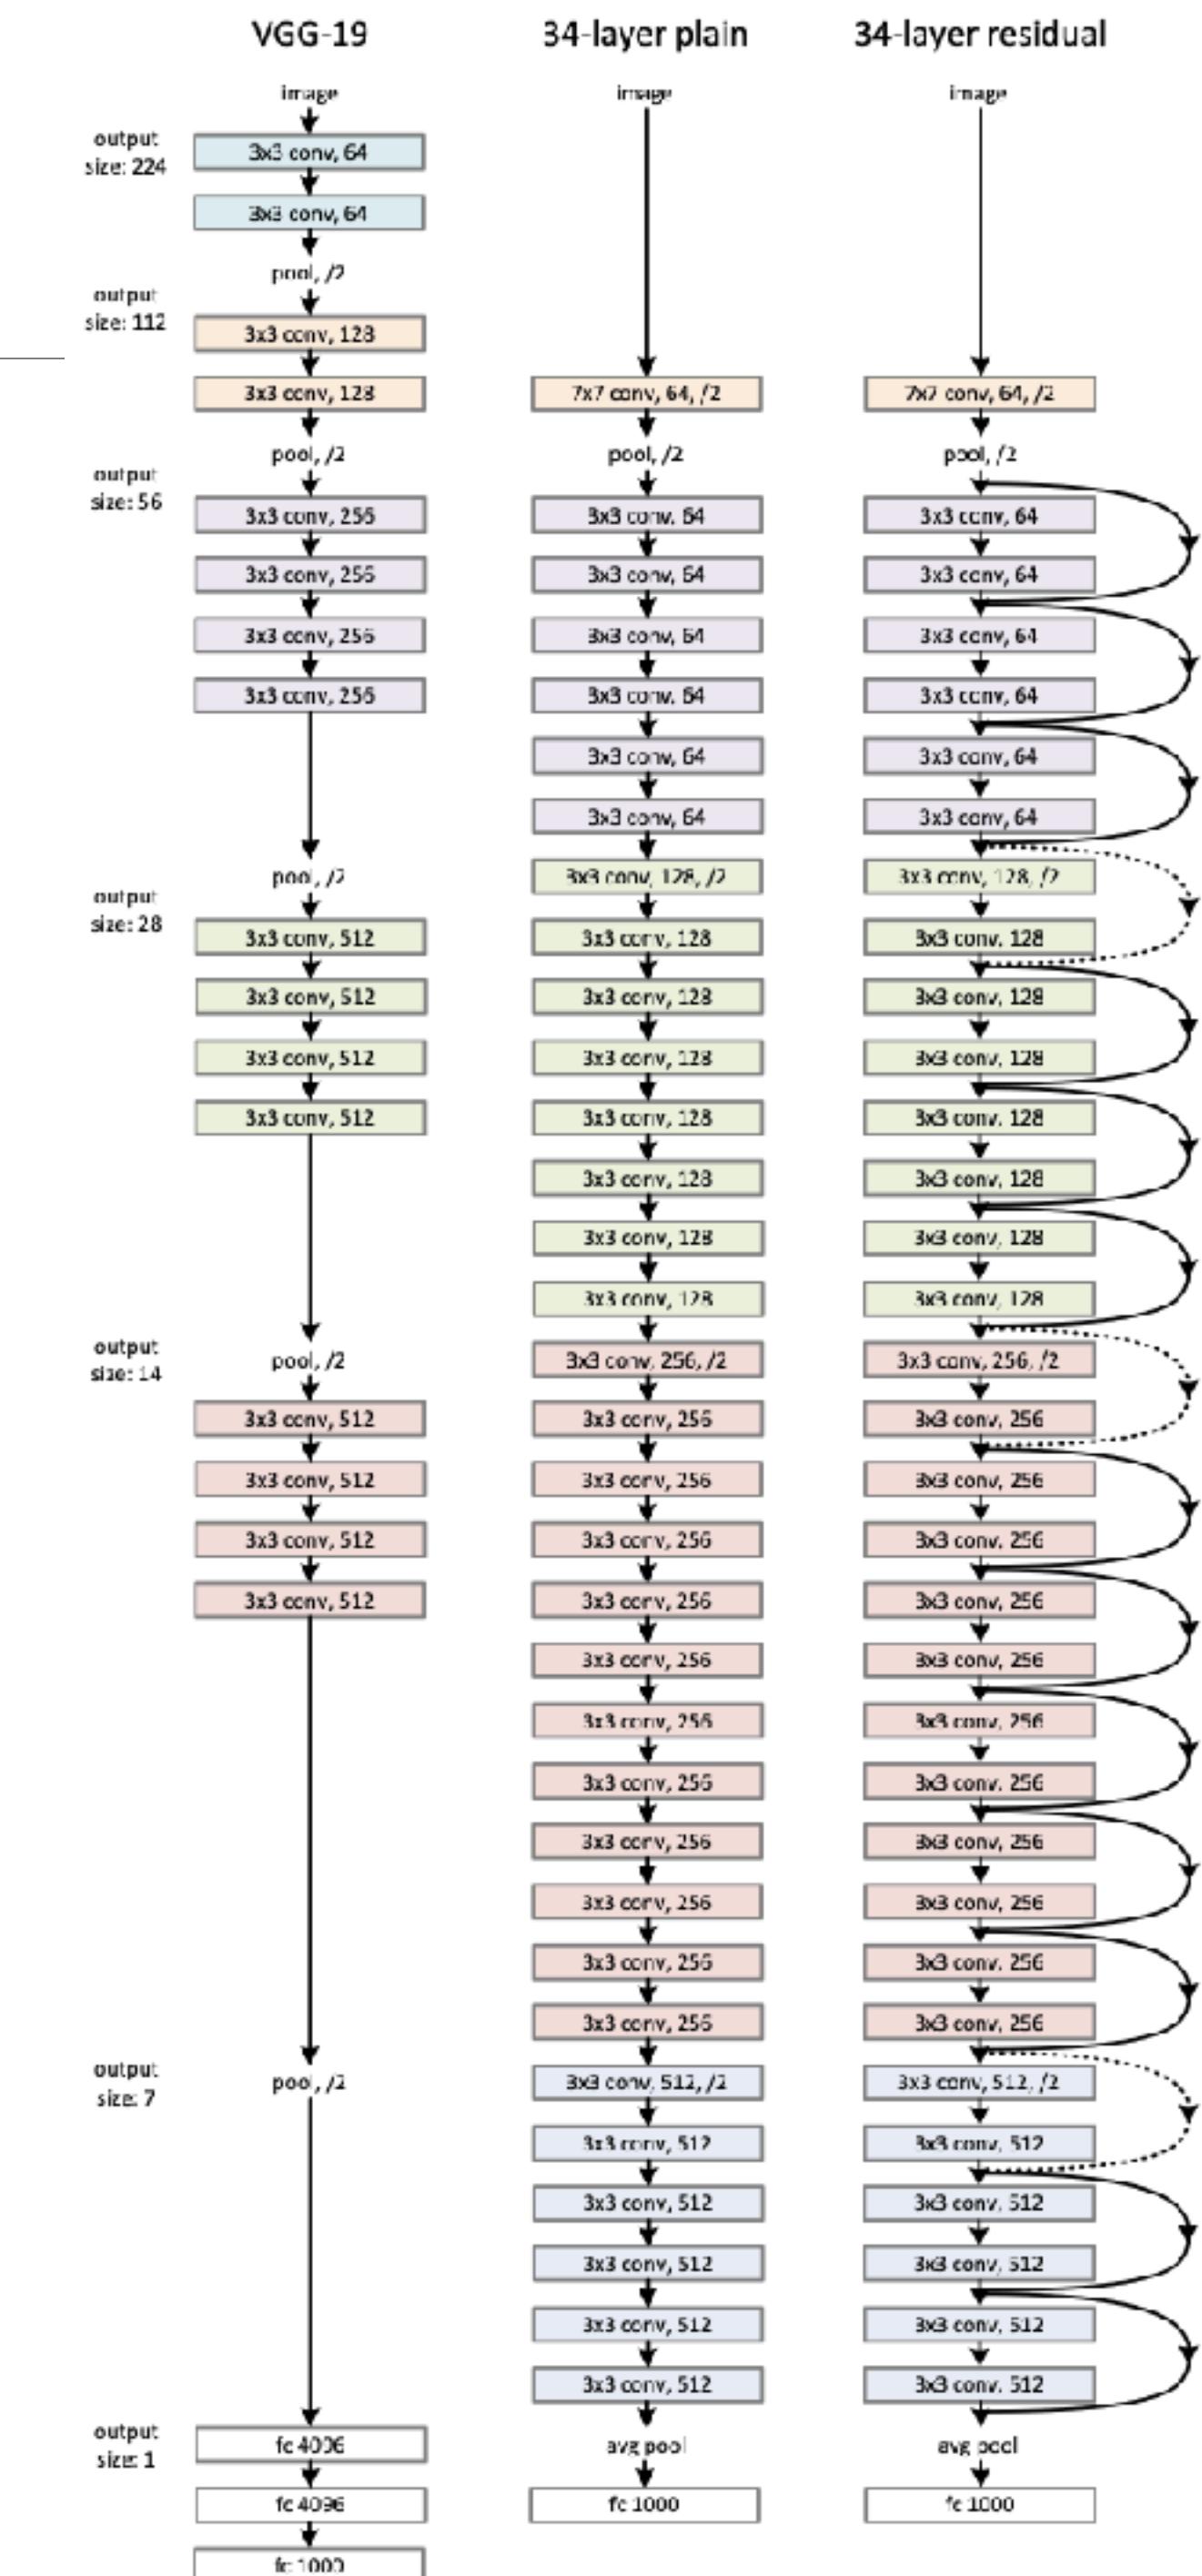
# Residual Networks (ResNets)

**Core idea:** Introduce “identity shortcut connections”



- Any additional “residual layer” should not hurt performance, as the network can simply use the identity.
  - Network can however improve by adding something to the identity.
- ▶ ResNets with hundreds or thousands of layers can be trained.
- ▶ Best-performing model on many datasets (including ResNet variants).

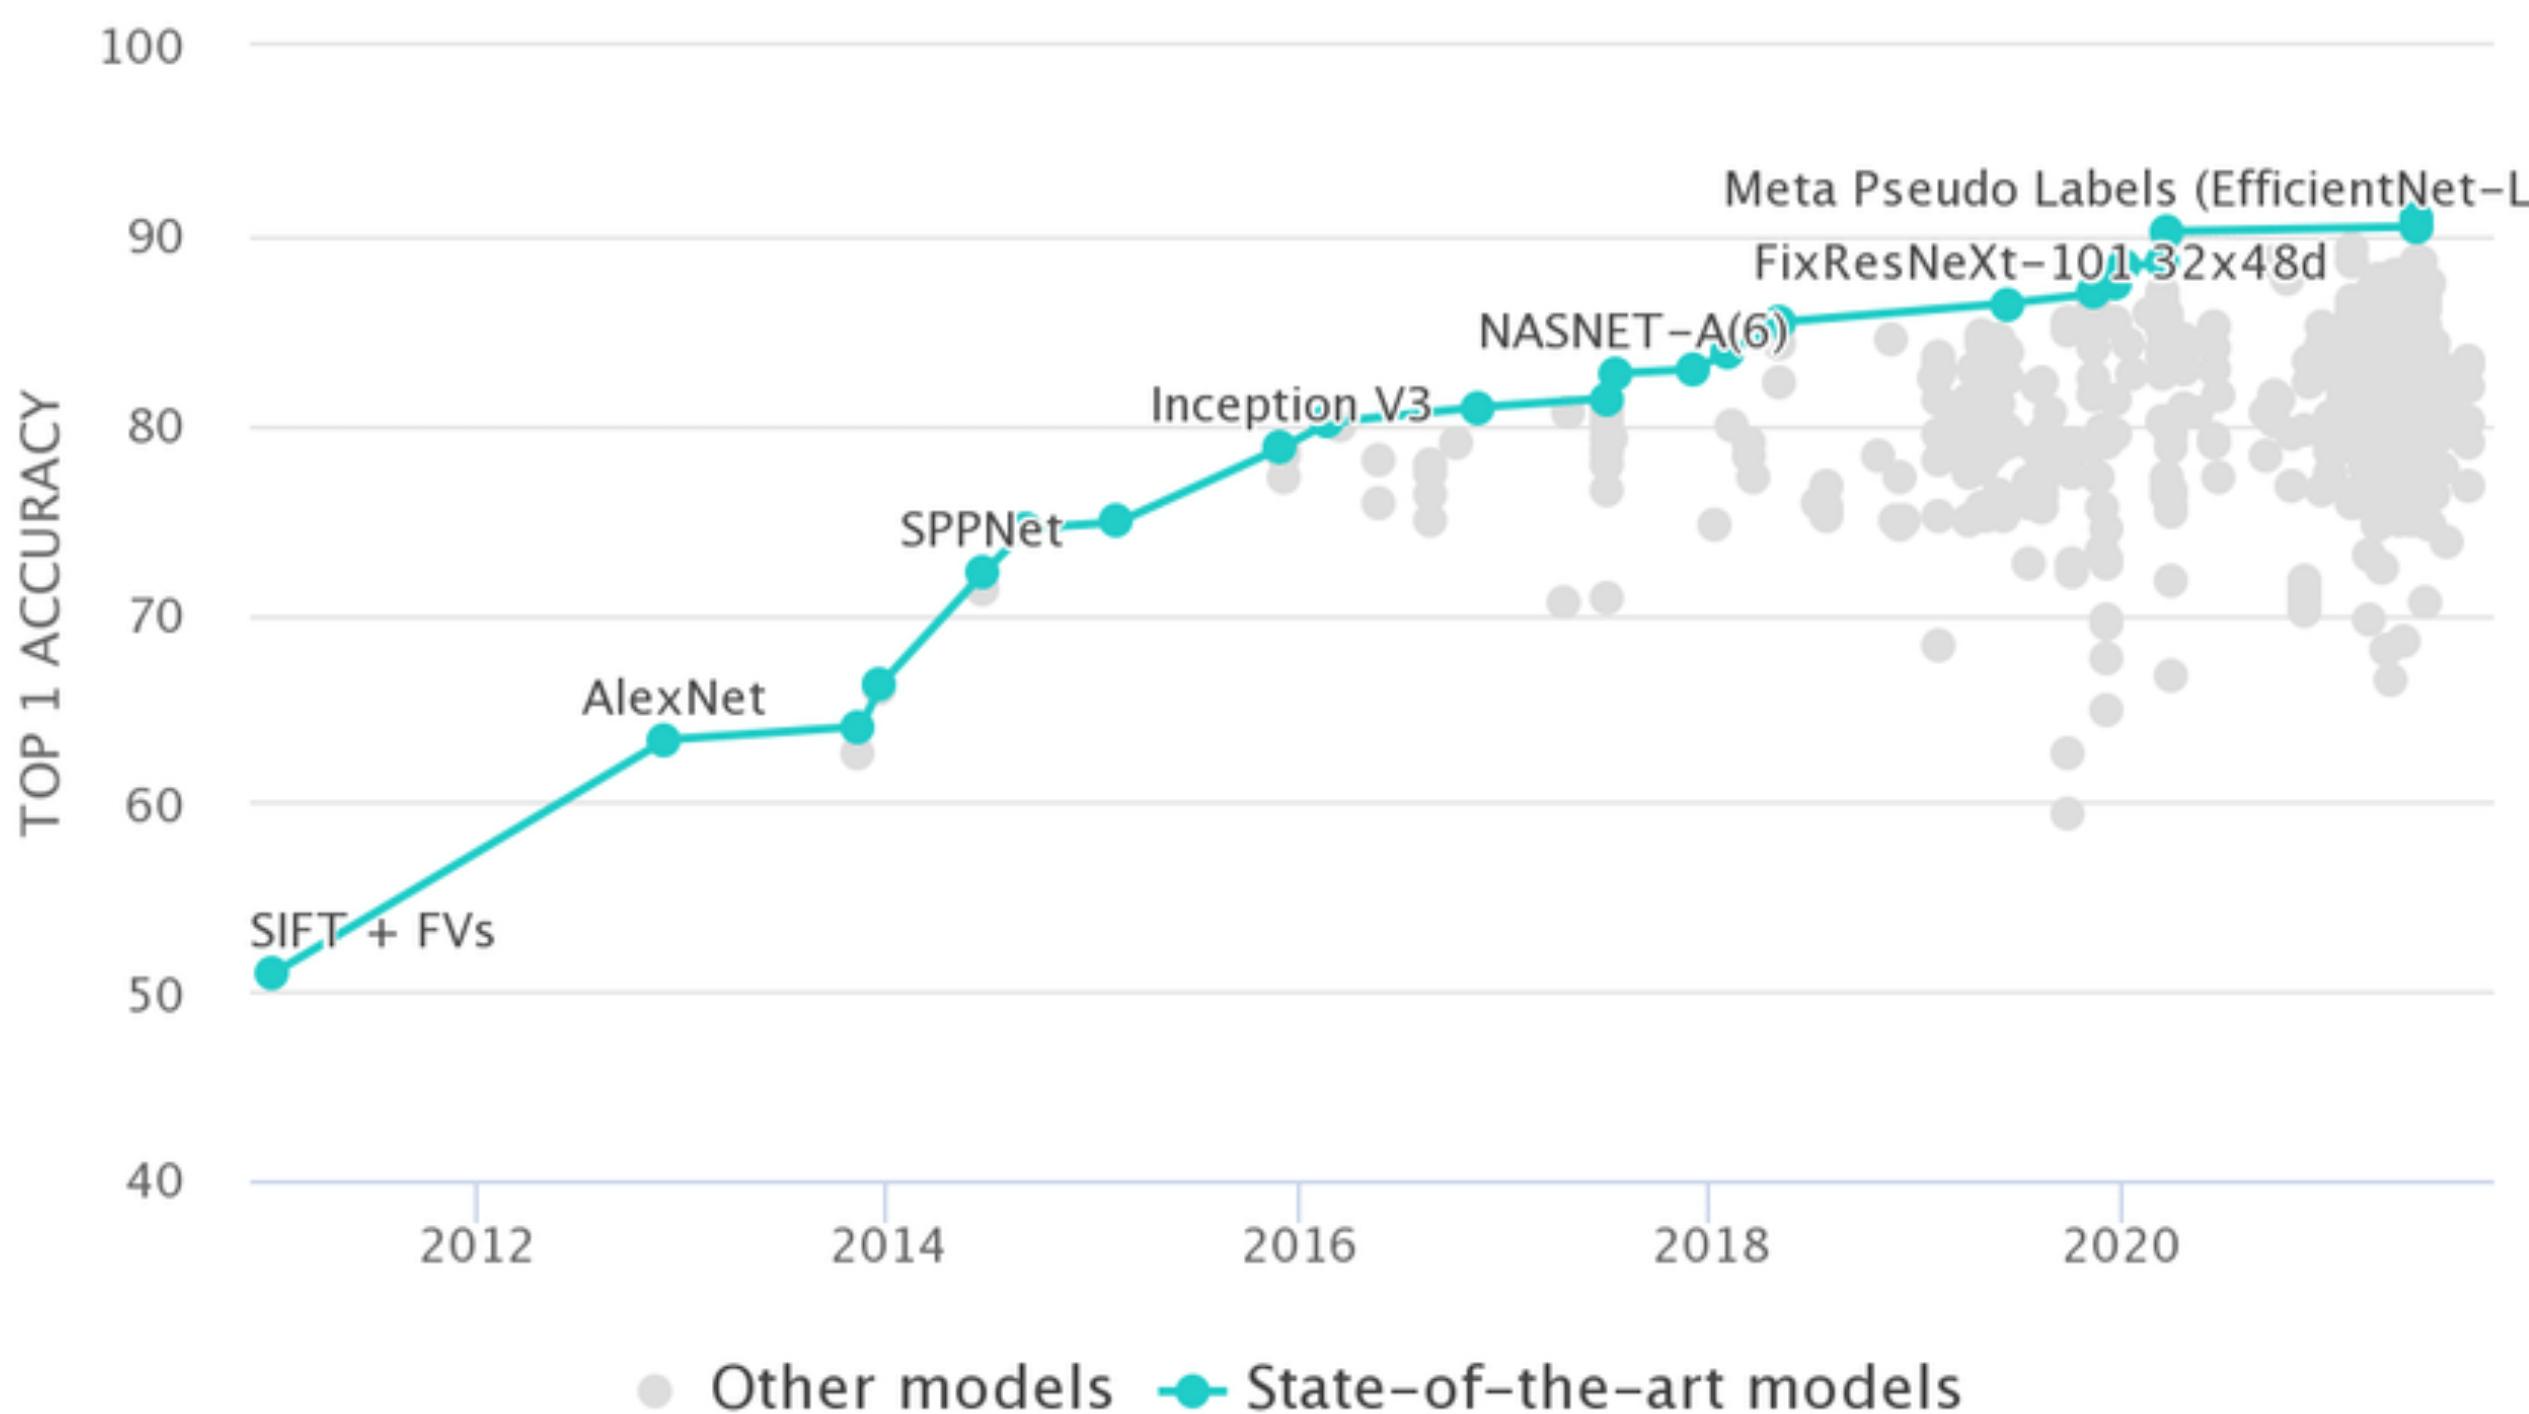
[Kaiming He, et al. “Deep residual learning for image recognition”, CVPR 2016]



# Summary

- ▶ Inspired by the visual system of vertebrates.
  - ▶ Hierarchical organization, increasing receptive fields, increasingly complex features.
- ▶ CNNs are feed-forward networks.
- ▶ Weight sharing reduces the # of parameters.
- ▶ Suitable for data with local structure and invariance properties.
- ▶ Deeper networks are hard to train:
  - ▶ Vanishing gradients problem.
  - ▶ Residual layers help.
- ▶ State-of-the-art in image classification.

Source: <https://paperswithcode.com/sota/image-classification-on-imagenet/>



# Today

---

- Convolutional Neural Networks
- Background
- CNN Architecture
- Going Deeper with CNNs
- Residual Networks (ResNets) & Extensions

## Questions?