

# Deep Learning: Recurrent Neural Networks

**Robert Legenstein**

Institute of Machine Learning and Neural Computation  
[robert.legenstein@tugraz.at](mailto:robert.legenstein@tugraz.at)

Deep Learning VO - WS 25/26  
Lecture 8

# Today

---


- ❑ Recurrent Neural Networks
  - ❑ Simple RNNs
  - ❑ Backpropagation Through Time
  - ❑ Architectural Variants
  - ❑ Gated RNNs (LSTM & GRU)

# Recurrent Neural Networks (RNNs)

**Typical input:** Sequence of values (vectors)  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(\tau)}$

**Typical output:** Sequence of values (vectors)  $\mathbf{o}^{(1)}, \dots, \mathbf{o}^{(\tau)}$

Network graph includes recurrent connections.



At each discrete time step  $t$ ,

- input  $\mathbf{x}^{(t)}$  is presented,
- hidden state  $\mathbf{h}^{(t)}$  is updated,
- output  $\mathbf{o}^{(t)}$  is computed,
- loss  $L^{(t)}$  is evaluated, and
- gradient is backpropagated through time.

**(Simple) RNNs were introduced in the 1980s.**

**Advantages:**


- Can process inputs of variable length.
- Implements parameter sharing over time.
- Can therefore generalize over time (compare to CNNs).

# Simple RNNs: Computation

**Typical input:** Sequence of values (vectors)  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(\tau)}$

**Typical output:** Sequence of values (vectors)  $\mathbf{o}^{(1)}, \dots, \mathbf{o}^{(\tau)}$

Network graph includes recurrent connections.



At each discrete time step  $t$ ,

- input  $\mathbf{x}^{(t)}$  is presented,
- hidden state  $\mathbf{h}^{(t)}$  is updated,
- output  $\mathbf{o}^{(t)}$  is computed,
- loss  $L^{(t)}$  is evaluated, and
- gradient is backpropagated through time.

$$\mathbf{a}^{(t)} = \mathbf{b} + W\mathbf{h}^{(t-1)} + U\mathbf{x}^{(t)}$$

$$\mathbf{h}^{(t)} = \tanh(\mathbf{a}^{(t)})$$

$$\mathbf{o}^{(t)} = \mathbf{c} + V\mathbf{h}^{(t)}$$

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{o}^{(t)})$$


# Unfolding in Time

View the RNN as a dynamical system driven by external input:


$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)}; \theta)$$

We can unfold the equation in time:

$$\mathbf{h}^{(3)} = f(\mathbf{h}^{(2)}, \mathbf{x}^{(3)}; \theta) = f(f(\mathbf{h}^{(1)}, \mathbf{x}^{(2)}; \theta), \mathbf{x}^{(3)}; \theta) = f(f(f(\mathbf{h}^{(0)}, \mathbf{x}^{(1)}; \theta), \mathbf{x}^{(2)}; \theta), \mathbf{x}^{(3)}; \theta)$$




Unfolding the computational graph:



Network typically learns to use  $\mathbf{h}^{(t)}$  as a lossy summary of the task-relevant aspects of the past sequence up to  $t$ .

# RNN Equations and Training



$$\mathbf{a}^{(t)} = \mathbf{b} + W\mathbf{h}^{(t-1)} + U\mathbf{x}^{(t)}$$

$$\mathbf{h}^{(t)} = \tanh(\mathbf{a}^{(t)})$$

$$\mathbf{o}^{(t)} = \mathbf{c} + V\mathbf{h}^{(t)}$$

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{o}^{(t)})$$

Training is via **backpropagation through time**: Compute the gradients in the unfolded computational graph.


# Today

---

- Recurrent Neural Networks
  - Simple RNNs
- Backpropagation Through Time
- Architectural Variants
- Gated RNNs (LSTM & GRU)

# Backpropagation Through Time

An example: a single neuron with a recurrent connection to itself.




$$h^0 = 0$$

$$h^t = \sigma \left( \underbrace{ux^t + wh^{t-1}}_{a^t} \right)$$

$$a^t = ux^t + wh^{t-1}$$

Unfolded Network:




# Unfolded Computational Graph

$$L = \sum_{t=1}^T L^t \quad L^t = \frac{1}{2} (h^t - h^{*t})^2$$

$$h^t = \sigma(a^t)$$

$$a^t = ux^t + wh^{t-1}$$




# Backpropagation Through Time

$$L = \sum_{t=1}^T L^t \quad L^t = \frac{1}{2} (h^t - h^{*t})^2$$

$$h^t = \sigma(a^t)$$

$$a^t = ux^t + wh^{t-1}$$




- We need  $\frac{\partial L}{\partial w}$  and  $\frac{\partial L}{\partial u}$  for parameter updates.
- To sketch the idea, we will show for  $\frac{\partial L}{\partial w}$ .

# Backpropagation Through Time

$$L = \sum_{t=1}^T L^t \quad L^t = \frac{1}{2} (h^t - h^{*t})^2$$

$$h^t = \sigma(a^t)$$

$$a^t = ux^t + wh^{t-1}$$




$$\frac{\partial L}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} \frac{\partial a^t}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} h^{t-1} = \sum_{t=1}^T \delta^t h^{t-1}$$

$\delta^t \stackrel{\text{def}}{=} \frac{\partial L}{\partial a^t}$

# Backpropagation Through Time

$$L = \sum_{t=1}^T L^t \quad L^t = \frac{1}{2} (h^t - h^{*t})^2$$




$$\frac{\partial L}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} \frac{\partial a^t}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} h^{t-1} = \sum_{t=1}^T \delta^t h^{t-1}$$

$\delta^t \stackrel{\text{def}}{=} \frac{\partial L}{\partial a^t} = \frac{\partial L}{\partial h^t} \frac{\partial h^t}{\partial a^t} = \frac{\partial L}{\partial h^t} \dot{\sigma}(a^t)$

$\frac{\partial L}{\partial h^t} = \underbrace{\frac{\partial L}{\partial L^t} \frac{\partial L^t}{\partial h^t}}_1 + \underbrace{\frac{\partial L}{\partial a^{t+1}} \frac{\partial a^{t+1}}{\partial h^t}}_{\delta^{t+1}} = \frac{\partial L^t}{\partial h^t} + \delta^{t+1} \cdot w$

# Backpropagation Through Time

$$L = \sum_{t=1}^T L^t \quad L^t = \frac{1}{2} (h^t - h^{*t})^2$$



$$\frac{\partial L}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} \frac{\partial a^t}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} h^{t-1} = \sum_{t=1}^T \delta^t h^{t-1}$$

↓

$$\delta^t \stackrel{\text{def}}{=} \frac{\partial L}{\partial a^t} = \frac{\partial L}{\partial h^t} \frac{\partial h^t}{\partial a^t} = \frac{\partial L}{\partial h^t} \dot{\sigma}(a^t)$$

↓

$$\frac{\partial L}{\partial h^t} = \frac{\partial L}{\partial L^t} \frac{\partial L^t}{\partial h^t} + \frac{\partial L}{\partial a^{t+1}} \frac{\partial a^{t+1}}{\partial h^t} = \frac{\partial L^t}{\partial h^t} + \delta^{t+1} \cdot w$$


$$\delta^t = \left( \frac{\partial L^t}{\partial h^t} + \delta^{t+1} \cdot w \right) \cdot \dot{\sigma}(a^t)$$

↑

Backpropagating error through time

# Backpropagation Through Time

$$L = \sum_{t=1}^T L^t \quad L^t = \frac{1}{2} (h^t - h^{*t})^2$$



**Notes on backpropagation through time:**

- ▶ Depends on the number of time steps.
- ▶ Before each parameter update, we need to wait and store everything until all computations are done.
- ▶ Memory demanding for longer sequences.
- ▶ Alternative: "truncated backprop through time".

$$\frac{\partial L}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} \frac{\partial a^t}{\partial w} = \sum_{t=1}^T \frac{\partial L}{\partial a^t} h^{t-1} = \sum_{t=1}^T \delta^t h^{t-1}$$

$\delta^t \stackrel{\text{def}}{=} \frac{\partial L}{\partial a^t} = \frac{\partial L}{\partial h^t} \frac{\partial h^t}{\partial a^t} = \frac{\partial L}{\partial h^t} \dot{\sigma}(a^t)$

$$\frac{\partial L}{\partial h^t} = \frac{\partial L}{\partial L^t} \frac{\partial L^t}{\partial h^t} + \frac{\partial L}{\partial a^{t+1}} \frac{\partial a^{t+1}}{\partial h^t} = \frac{\partial L^t}{\partial h^t} + \delta^{t+1} \cdot w$$

$$\delta^t = \left( \frac{\partial L^t}{\partial h^t} + \delta^{t+1} \cdot w \right) \cdot \dot{\sigma}(a^t)$$

**Backpropagating error through time**

# Today

---

## Recurrent Neural Networks


Simple RNNs

Backpropagation Through Time


Architectural Variants

Gated RNNs (LSTM & GRU)

## (1) Deep Recurrent Network




## (2) Single output at the end of the sequence



Some potential use cases:

- Text classification.
- Detecting anomalies from a given measurement sequence.
- Summarize a sequence and produce a fixed-size representation, then use this as input for further processing.

# (3) Encoder-Decoder (Sequence-to-Sequence) Architecture



- If the output sequence does not have the same length as input sequence, e.g. in language translation.

**Input:** Sequence  $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n_x)}$

**Output:** Sequence  $\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(n_y)}$

**Encoder:** Processes input and emits the context  $C$ , typically a simple function of its final hidden state.

**Decoder:** Generates the output sequence conditioned on this context  $C$ .


# Today

---

- Recurrent Neural Networks
  - Simple RNNs
  - Backpropagation Through Time
  - Architectural Variants
- Gated RNNs (LSTM & GRU)

# Problem: Challenging long-term dependencies

---



- *Vanishing/exploding gradients* problem:

**Gradients propagated over many stages tend to either vanish or explode.**

We can mitigate *exploding gradients* by gradient clipping.

What about *vanishing gradients*?

- In order to store memories in a way that is robust to small perturbations, the RNN must enter a region of parameter space where gradients vanish.
- The gradient of a long term interaction has exponentially smaller magnitude than that of a short term interaction.

# Long Short-Term Memory (LSTM)

Idea: Create paths through time that have derivatives that neither vanish nor explode.

## LSTM basic unit: A memory cell

- A linear neuron with a unit-weight self loop, where:
  - an **input gate** controls whether to load something in,
  - an **output gate** controls whether to make the content available to others,
  - a **forget gate** controls whether to forget the content.




[S. Hochreiter and J. Schmidhuber. "Long short-term memory." *Neural Computation*, 1997.]

**Abstract**  
Learning to store information over extended time intervals via recurrent backpropagation takes a very long time, mostly due to insufficient, decaying error back flow. We briefly review Hochreiter's 1991 analysis of this problem, then address it by introducing a novel, efficient, gradient-based method called "Long Short-Term Memory" (LSTM). Truncating the gradient

# LSTM - Memory Cell

---




$$s^t = s^{t-1} + x^t$$

[S. Hochreiter and J. Schmidhuber. "Long short-term memory." *Neural Computation*, 1997.]

# LSTM - Memory Cell


**Input Gate:**



[S. Hochreiter and J. Schmidhuber. "Long short-term memory." *Neural Computation*, 1997.]

# LSTM - Memory Cell


**Forget Gate:**



[S. Hochreiter and J. Schmidhuber. "Long short-term memory." *Neural Computation*, 1997.]


# LSTM - Memory Cell

## Output Gate:



[S. Hochreiter and J. Schmidhuber. "Long short-term memory." *Neural Computation*, 1997.]


# LSTM - Memory Cell



$$s^t = f^t s^{t-1} + g^t x^t$$
$$h^t = q^t \cdot \tanh(s^t)$$

[S. Hochreiter and J. Schmidhuber. "Long short-term memory." *Neural Computation*, 1997.]

# LSTM - Memory Cell - Details



$$h_i^{(t)} = \tanh(s_i^{(t)}) q_i^{(t)}$$

$$s_i^{(t)} = f_i^{(t)} s_i^{(t-1)} + g_i^{(t)} \sigma \left( b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right)$$

$$g_i^{(t)} = \sigma \left( b_i^g + \sum_j U_{i,j}^g x_j^{(t)} + \sum_j W_{i,j}^g h_j^{(t-1)} \right)$$


$$f_i^{(t)} = \sigma \left( b_i^f + \sum_j U_{i,j}^f x_j^{(t)} + \sum_j W_{i,j}^f h_j^{(t-1)} \right)$$

$$q_i^{(t)} = \sigma \left( b_i^o + \sum_j U_{i,j}^o x_j^{(t)} + \sum_j W_{i,j}^o h_j^{(t-1)} \right)$$

$\mathbf{x}^{(t)}$  is the current input vector.

$\mathbf{h}^{(t)}$  is the current hidden layer vector.

# LSTM - Memory Cell - Details



$$h_i^{(t)} = \tanh(s_i^{(t)}) q_i^{(t)}$$

$$s_i^{(t)} = f_i^{(t)} s_i^{(t-1)} + g_i^{(t)} \sigma \left( b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right)$$

$$g_i^{(t)} = \sigma \left( b_i^g + \sum_j U_{i,j}^g x_j^{(t)} + \sum_j W_{i,j}^g h_j^{(t-1)} \right)$$

$$f_i^{(t)} = \sigma \left( b_i^f + \sum_j U_{i,j}^f x_j^{(t)} + \sum_j W_{i,j}^f h_j^{(t-1)} \right)$$

$$q_i^{(t)} = \sigma \left( b_i^o + \sum_j U_{i,j}^o x_j^{(t)} + \sum_j W_{i,j}^o h_j^{(t-1)} \right)$$

$\mathbf{x}^{(t)}$  is the current input vector.

$\mathbf{h}^{(t)}$  is the current hidden layer vector.

# Gated Recurrent Unit (GRU)

A simplified version of an LSTM cell.

- an **update gate** :  $u$  simultaneously controls the forgetting factor and the input gating.
- a **reset gate** :  $r$  controls which parts of the hidden state are used for the update.

$$h_i^{(t)} = u_i^{(t-1)} h_i^{(t-1)} + (1 - u_i^{(t-1)}) \sigma \left( b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} r_j^{(t-1)} h_j^{(t-1)} \right)$$

- The output values of  $u_i^t$  and  $r_j^t$  are computed in the usual way.

**Recall LSTMs:**  $h_i^{(t)} = \tanh(s_i^{(t)}) q_i^{(t)}$

$$s_i^{(t)} = f_i^{(t)} s_i^{(t-1)} + g_i^{(t)} \sigma \left( b_i + \sum_j U_{i,j} x_j^{(t)} + \sum_j W_{i,j} h_j^{(t-1)} \right)$$

- ▶ GRUs have only two gates (instead of three with LSTMs).
- ▶ A single "update gate". (LSTMs: input & forget gates).
- ▶ Therefore usually less parameters than LSTMs.

# Today

---

- Recurrent Neural Networks
- Simple RNNs
- Backpropagation Through Time
- Architectural Variants
- Gated RNNs (LSTM & GRU)

## Questions?