

MARKOV CHAIN MONTE CARLO II

PROBABILISTIC DECISION MAKING VU

(REINFORCEMENT LEARNING VO)

Thomas Wedenig

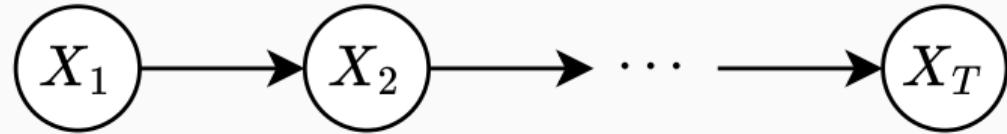
Dec 17, 2025

Institute of Machine Learning and Neural Computation
Graz University of Technology, Austria

1. MCMC - Recap
2. Estimating Expectations with MCMC Samples
3. Gibbs Sampling
4. Langevin MCMC
5. Hamiltonian Monte Carlo
6. Diagnosing MCMC

MCMC - RECAP

- Goal: Draw **samples** from **distribution** $\pi(x)$
- Typically, only $\tilde{\pi}(x) \propto \pi(x)$ is known
 - e.g., Bayesian inference: $\tilde{\pi}_y(x) := p(x, y) \propto p(x | y)$
 - e.g., Boltzmann distributions: $\tilde{\pi}(x) := \exp(-E(x)), E : \mathcal{X} \rightarrow \mathbb{R}$
- We care about **samples** because...
 - they might be **the object of interest** (e.g., sample image from generative model), or
 - we may need samples to **approximate expectations** (Monte Carlo)



- Idea: Construct Markov Chain

$$p(x_1, x_2, \dots, x_T) = p(x_1) p(x_2 | x_1) p(x_3 | x_2) \dots p(x_T | x_{T-1})$$

- with unique stationary distribution π ,
- its marginal distribution converges to π when $T \rightarrow \infty$,
- $p(x_1)$ and $p(x_t | x_{t-1})$ are easy to sample from

Sampling from the chain for “long enough” yields samples with marginal distribution close to π !



Initialize x

repeat:

$x \leftarrow$ some (random) update to x



Running a
Markov Chain Monte Carlo
Algorithm

Metropolis-Hastings

Given: Target $\tilde{\pi}$ and proposal distribution $q(y | x)$.

Initialize $X_1 \sim p(X_1)$ and repeat for $t = 1, \dots, T - 1$:

- 1 Given current state $X_t = x$, propose:

$$y \sim q(y | x)$$

- 2 Compute acceptance probability:

$$\alpha(x \rightarrow y) = \min \left(1, \frac{\tilde{\pi}(y) q(x | y)}{\tilde{\pi}(x) q(y | x)} \right)$$

- 3 Accept or reject:

Sample $U \sim \text{Uniform}(0, 1)$,

$$X_{t+1} = \begin{cases} y & \text{if } U \leq \alpha(x \rightarrow y) \\ x & \text{else} \end{cases}$$

Metropolis-Hastings Update

- 1 Propose $y \sim q(y | x)$
- 2 Compute $\alpha = \min\left(1, \frac{\tilde{\pi}(y) q(x|y)}{\tilde{\pi}(x) q(y|x)}\right)$
- 3 Accept y with probability α

- For continuous state spaces, we could choose a **Gaussian proposal**

$$q(y | x) := \mathcal{N}(y; x, \sigma^2 I)$$

- Demo last week: Discrete state space with **proposal**

$$q(y | x) = \text{Uniform}(\{\text{bitstrings that differ from } x \text{ in exactly one bit}\})$$

ANALOGY: OPTIMIZATION

- Assume we want to find the **maximum** of $\tilde{\pi}$
- Idea: With current state x , repeat
 - 1 Propose $y \sim q(y | x)$
 - 2 Set $\alpha = \tilde{\pi}(y)/\tilde{\pi}(x)$
 - 3 If $\alpha \geq 1$, accept y , else stay at x
- “Greedy random search”

RANDOM-WALK METROPOLIS-HASTINGS DEMO

<https://chi-feng.github.io/mcmc-demo/app.html>

ESTIMATING EXPECTATIONS WITH MCMC SAMPLES

Recall Goal: Draw i.i.d. samples from π .

Are samples from our Markov Chain **i.i.d.** samples from π ? 🤔

No !

- If samples were i.i.d., X_t and X_{t+1} would be **independent**
- In the beginning, X_t is not even **marginally** distributed like π (not converged yet)
- Even when the chain has perfectly converged to π , and thus $X_t \sim \pi$, $X_{t+1} \sim \pi$ **marginally**, we still have **possibly large correlation** between X_t and X_{t+1}
- When the finite-state chain has converged, we have

$$p(X_t = j, X_{t+1} = i) = p(X_t = j)p(X_{t+1} = i | X_t = j) = \pi_j P_{ij} \neq \pi_j \pi_i = p(X_t = j)p(X_{t+1} = i)$$

Are we doomed? Monte Carlo estimators need i.i.d. samples 😢

Thankfully, no! 🎉

- There is an **MCMC counterpart** to classic i.i.d. Monte Carlo
- It assumes our Markov Chain has π as its **unique stationary distribution**
- And it assumes **certain properties**¹

¹Irreducible (when state space finite), irreducible and positive recurrent (state space countable), π -irreducible and Harris recurrent (state space uncountable.)

Ergodic Theorem

Assume X_1, X_2, \dots is a Markov Chain with above properties and stationary distribution π . Then,

$$\hat{\mathbb{E}}_T := \frac{1}{T} \sum_{t=1}^T f(X_t) \xrightarrow{\text{a.s.}} \mathbb{E}_{X \sim \pi}[f(X)]$$

as $T \rightarrow \infty$ for every initial state X_1 .

- Almost sure (a.s.) convergence means

$$\mathbb{P}\left(\lim_{T \rightarrow \infty} \hat{\mathbb{E}}_T = \mathbb{E}_{X \sim \pi}[f(X)]\right) = 1$$

Recap: Properties of MC Estimate

Given i.i.d. samples $Y_1, \dots, Y_T \sim \pi$, the Monte Carlo estimate of $\mathbb{E}_{X \sim \pi} [f(X)]$ is

$$\hat{\mathbb{E}}_T^{(\text{MC})} := \frac{1}{T} \sum_{t=1}^T f(Y_t).$$

- Unbiased for every T : $\mathbb{E} [\hat{\mathbb{E}}_T^{(\text{MC})}] = \mathbb{E}_{X \sim \pi} [f(X)]$

Properties of MCMC Estimate

- MCMC estimate is generally² **biased** for any finite T
 - We can reduce bias by **discarding the first B samples** (burn-in phase)
 - Bias goes to 0 as $T \rightarrow \infty$, i.e., **asymptotically unbiased**

²It would be unbiased if the initial distribution was already π , but we can't do that in practice.

Recap: Properties of MC Estimate

- Variance decreases linearly with T : $\text{var} \left[\hat{\mathbb{E}}_T^{(\text{MC})} \right] = \frac{\text{var}[f(X)]}{T}$

Properties of MCMC Estimate

- As $T \rightarrow \infty$, the **variance** of the MCMC estimator is approximately

$$\text{var} \left[\hat{\mathbb{E}}_T \right] \approx \frac{\sigma_{\text{as}}^2}{T}$$

with

$$\sigma_{\text{as}}^2 = \text{var}_\pi[f(X)] + 2 \sum_{n=1}^{\infty} \text{cov}_\pi(f(X_0), f(X_n)) .$$

where $X_0 \sim \pi, X_1 \sim \pi, \dots$ (i.e., chain in **stationary** regime)

- **Note:** σ_{as}^2 typically **larger** than $\text{var}_\pi[f(X)]$, but still decreases **linearly** with T

GIBBS SAMPLING

- Target: $\pi(\mathbf{x})$ with $\mathbf{x} = (x_1, \dots, x_d)^\top \in \mathcal{X}$
 - e.g., $\mathcal{X} = \mathbb{R}^d$ or $\mathcal{X} = \{-1, +1\}^d$
- Let $\mathbf{x}_{-i} := (x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_d)$
- Assume conditionals $\pi(x_i | \mathbf{x}_{-i})$ can be sampled easily for every i

Gibbs Sampling

Initialize $\mathbf{x}^{(1)} \sim p(\mathbf{x}^{(1)})$ and repeat for $t = 1, \dots, T - 1$:

- Given current state $\mathbf{x}^{(t)} := (x_1^{(t)}, \dots, x_d^{(t)})$, sample

$$x_1^{(t+1)} \sim \pi(x_1 | x_2^{(t)}, x_3^{(t)}, \dots, x_d^{(t)})$$

$$x_2^{(t+1)} \sim \pi(x_2 | x_1^{(t+1)}, x_3^{(t)}, \dots, x_d^{(t)})$$

$$x_3^{(t+1)} \sim \pi(x_3 | x_1^{(t+1)}, x_2^{(t+1)}, x_4^{(t)}, \dots, x_d^{(t)})$$

$$\vdots$$

$$x_d^{(t+1)} \sim \pi(x_d | x_1^{(t+1)}, x_2^{(t+1)}, \dots, x_{d-1}^{(t+1)})$$

GIBBS SAMPLING DEMO

<https://chi-feng.github.io/mcmc-demo/app.html>

Metropolis-Hastings Update

- 1 Propose $y \sim q(y | x)$
- 2 Compute $\alpha = \min \left(1, \frac{\tilde{\pi}(y) q(x|y)}{\tilde{\pi}(x) q(y|x)} \right)$
- 3 Accept y with probability α

- Gibbs Sampling can be viewed as an instance of Metropolis-Hastings (MH):

$$y \sim q_i(y | x) := \pi(y_i | x_{-i}) \mathbf{1}\{y_{-i} = x_{-i}\}.$$

- Hence, $y_{-i} = x_{-i}$, and thus

$$\pi(x) = \pi(x_{-i}) \pi(x_i | x_{-i}), \quad \pi(y) = \pi(x_{-i}) \pi(y_i | x_{-i}),$$

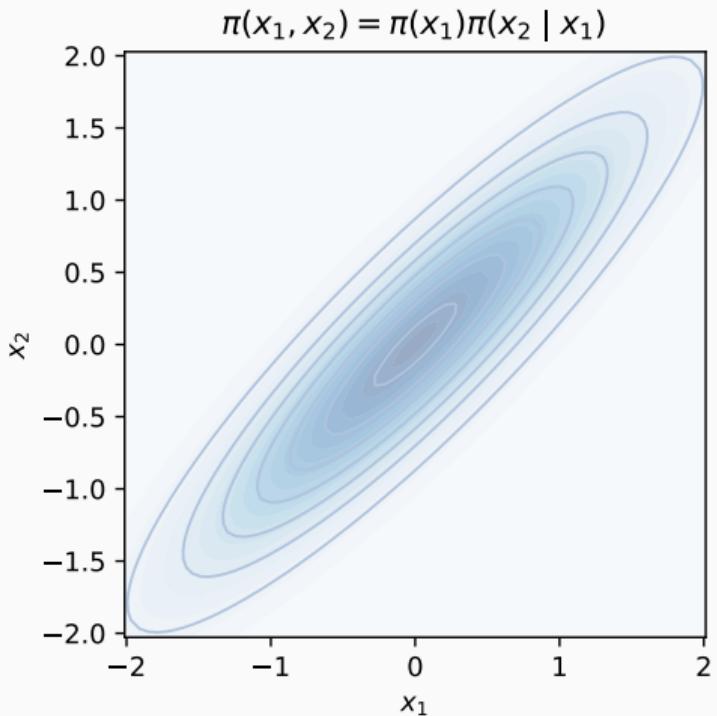
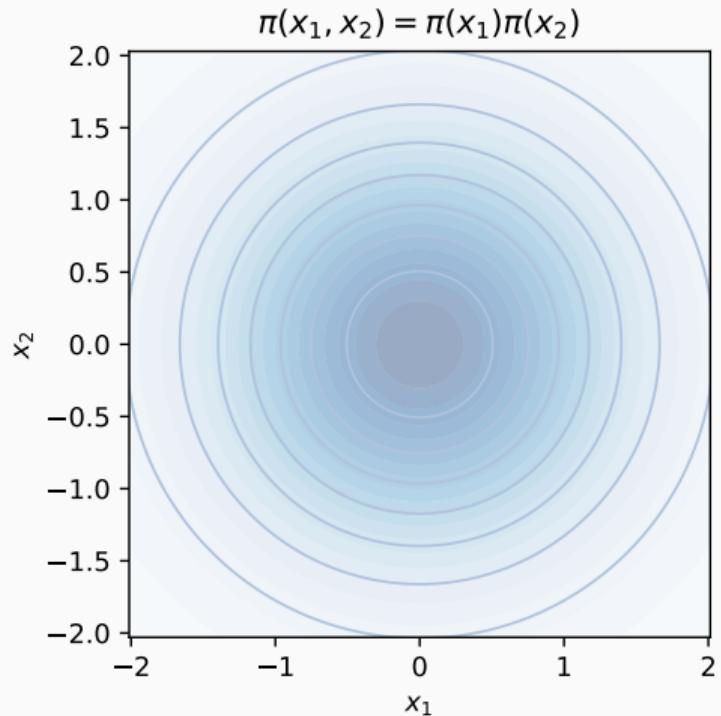
- The MH ratio becomes

$$\frac{\pi(y) q_i(x | y)}{\pi(x) q_i(y | x)} = \frac{[\pi(x_{-i}) \pi(y_i | x_{-i})] \pi(x_i | x_{-i})}{[\pi(x_{-i}) \pi(x_i | x_{-i})] \pi(y_i | x_{-i})} = 1$$

and we accept with probability $\alpha = 1$.

- Every coordinate update can be seen as a MH step that always accepts 😍
- Gibbs Sampling typically only stores **every d^{th} MH sample**
 - Technically, this is not necessary if we e.g. apply the **ergodic theorem**
 - But it **reduces autocorrelation**
 - ...at the expense of “throwing away” samples
- Theoretical caveat: In this view, q_i is **time-dependent**
 - Issue: the ergodic theorem assumes a **time-homogeneous** chain
 - Luckily, we can make it time-homogeneous by augmenting the state space $\mathbf{z} := (\mathbf{x}, i)$
 - We then consider $q(\mathbf{y} \mid \mathbf{z})$ in this augmented space, which also updates i accordingly

GIBBS SAMPLING FROM A GAUSSIAN



- Per-coordinate updates may yield **large autocorrelation**
 - Especially for **highly correlated variables**
 - We **cannot “move diagonally” in the state space**
- In the above example, we of course could sample both coordinates at once !
 - In general, we can't sample **all coordinates at once**
 - ...otherwise we would not need MCMC 😂
- Is there something in the middle? 🤔
- **Blocked Gibbs Sampling:** Sample multiple coordinates at once

Blocked Gibbs Sampling

Let

$$\mathbf{x} = (\mathbf{x}_{B_1}, \mathbf{x}_{B_2}, \dots, \mathbf{x}_{B_K}), \quad B_1 \cup \dots \cup B_K = \{1, \dots, d\}.$$

Initialize $\mathbf{x}^{(1)} \sim p(\mathbf{x}^{(1)})$ and repeat for $t = 1, \dots, T - 1$:

- Given current state $\mathbf{x}^{(t)} := (\mathbf{x}_{B_1}^{(t)}, \dots, \mathbf{x}_{B_K}^{(t)})$, sample

$$\mathbf{x}_{B_1}^{(t+1)} \sim \pi\left(\mathbf{x}_{B_1} \mid \mathbf{x}_{B_2}^{(t)}, \mathbf{x}_{B_3}^{(t)}, \dots, \mathbf{x}_{B_K}^{(t)}\right)$$

$$\mathbf{x}_{B_2}^{(t+1)} \sim \pi\left(\mathbf{x}_{B_2} \mid \mathbf{x}_{B_1}^{(t+1)}, \mathbf{x}_{B_3}^{(t)}, \dots, \mathbf{x}_{B_K}^{(t)}\right)$$

$$\mathbf{x}_{B_3}^{(t+1)} \sim \pi\left(\mathbf{x}_{B_3} \mid \mathbf{x}_{B_1}^{(t+1)}, \mathbf{x}_{B_2}^{(t+1)}, \mathbf{x}_{B_4}^{(t)}, \dots, \mathbf{x}_{B_K}^{(t)}\right)$$

$$\vdots$$

$$\mathbf{x}_{B_K}^{(t+1)} \sim \pi\left(\mathbf{x}_{B_K} \mid \mathbf{x}_{B_1}^{(t+1)}, \mathbf{x}_{B_2}^{(t+1)}, \dots, \mathbf{x}_{B_{K-1}}^{(t+1)}\right)$$

What if we cannot directly sample from some **conditional distribution** ?

$$\mathbf{x}_{B_i}^{(t+1)} \sim \pi\left(\mathbf{x}_{B_i} \mid \mathbf{x}_{B_1}^{(t+1)}, \dots, \mathbf{x}_{B_{i-1}}^{(t+1)}, \mathbf{x}_{B_{i+1}}^{(t)}, \dots, \mathbf{x}_{B_K}^{(t)}\right)$$

- We could use MCMC again **just for this block !**
 - e.g., random walk MH (or something smarter)
- This is called **Metropolis-within-Gibbs**

LANGEVIN MCMC

Maximizing $\tilde{\pi}$ via greedy random search

With current state x , repeat:

- 1 Propose $y \sim q(y | x)$
- 2 Set $\alpha = \tilde{\pi}(y)/\tilde{\pi}(x)$
- 3 If $\alpha \geq 1$, accept y , else stay at x

This is clearly **ridiculous** in most cases 😂

- If $\tilde{\pi}$ is a **differentiable** PDF, we can use **gradient information** to find a maximum
 - e.g., **gradient ascent**
- In **MCMC sampling**, we want to **explore regions of high probability**
 - Gradients can guide us towards regions of **high density**

- Metropolis-Adjusted Langevin Algorithm (MALA)

- MALA is an instance of Metropolis-Hastings with

$$q(y | x) = \mathcal{N} \left(y; x + \frac{\eta}{2} \nabla_x \log \pi(x), \eta I \right)$$

with step size $\eta > 0$.

- As usual, we accept y with probability

$$\alpha(x \rightarrow y) = \min \left(1, \frac{\tilde{\pi}(y) q(x | y)}{\tilde{\pi}(x) q(y | x)} \right)$$

- Note: We have to be able to compute $\nabla_x \log \tilde{\pi}(x) = \nabla_x \log \pi(x)$
 - And evaluate $\tilde{\pi}$ in the acceptance step

MALA DEMO

<https://chi-feng.github.io/mcmc-demo/app.html>

- This choice of q is motivated by the Langevin Stochastic Differential Equation (SDE):

$$dX_t = \frac{1}{2} \nabla \log \pi(X_t) dt + dW_t$$

where W_t is a Wiener process (Brownian motion)

- Discretizing this SDE using Euler-Maruyama gives:

$$X_{t+1} = X_t + \frac{\eta}{2} \nabla \log \pi(X_t) + \sqrt{\eta} z, \quad z \sim \mathcal{N}(0, I)$$

which is exactly a sample from this $q(y | x)$

- For MALA, we need to know $\nabla_x \log \tilde{\pi}(x)$ and $\tilde{\pi}(x)$
- Sometimes, we **do not have access to** $\tilde{\pi}$, but **only to** $\nabla_x \log \tilde{\pi}(x)$
 - e.g., Score-based Generative Models (Diffusion Models)
- Idea: Let's just **always accept** 😂
- This is called the **Unadjusted Langevin Algorithm (ULA)**
- Not a MH algorithm anymore, and the stationary distribution is $\pi_\eta \neq \pi$!
 - It has stationary bias $\mathbb{E}_{\pi_\eta} [f] - \mathbb{E}_\pi [f]$ of order $O(\eta)$ in general (even as $T \rightarrow \infty$)

HAMILTONIAN MONTE CARLO

- To explore π , we want to make “**large steps**”
 - while still staying in **high probability regions**
- Also, we wish to incorporate **momentum**
 - Don’t forget direction immediately
 - Physical analog: ball with **mass**

- If $x \in \mathbb{R}^d$, augment the state space by a new **momentum variable** $p \in \mathbb{R}^d$
- Define the **Hamiltonian**

$$H(x, p) = \underbrace{-\log \tilde{\pi}(x)}_{E(x)} + \underbrace{\frac{1}{2} p^\top M^{-1} p}_{K(p)}.$$

where M is the **mass matrix**.

- This defines a **new joint distribution**

$$\pi_{\text{aug}}(x, p) = Z_H^{-1} \exp(-H(x, p)) = Z_H^{-1} \tilde{\pi}(x) \exp\left(-\frac{1}{2} p^\top M^{-1} p\right).$$

with **marginal** $\pi_{\text{aug}}(x) = \int \pi_{\text{aug}}(x, p) dp = \pi(x)$

- HMC is again an instance of **Metropolis-Hastings**, but **sampling from** $\pi_{\text{aug}}(x, p)$
 - Since we want to **sample** from π (the marginal), we **throw away** the sampled p

- In a single MH proposal step, we will consider \mathbf{x}, \mathbf{p} to change over “time” t
 - This is **not** time t in the Markov chain
- We want that the **Hamiltonian is conserved over time**:

$$\frac{dH(\mathbf{x}, \mathbf{p})}{dt} = \frac{\partial H}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial t} + \frac{\partial H}{\partial \mathbf{p}} \frac{\partial \mathbf{p}}{\partial t} \stackrel{!}{=} 0$$

- Which we enforce by letting \mathbf{x}, \mathbf{p} evolve via **Hamiltonian dynamics**:

$$\frac{\partial \mathbf{x}}{\partial t} := \frac{\partial H}{\partial \mathbf{p}} \quad \frac{\partial \mathbf{p}}{\partial t} := -\frac{\partial H}{\partial \mathbf{x}}$$

- Let $(\mathbf{y}, \mathbf{p}') = \Phi_\tau(\mathbf{x}, \mathbf{p})$ be a **ODE solver** that **integrates this ODE** over some time τ
 - If we made no numerical error, then $H(\mathbf{x}, \mathbf{p}) = H(\Phi_\tau(\mathbf{x}, \mathbf{p}))$ for any τ and (\mathbf{x}, \mathbf{p})
 - In practice, we set step size ε and number of steps L , where then $\tau = L\varepsilon$

Hamiltonian Monte Carlo

Initialize $X_1 \sim p(X_1)$ and repeat for $t = 1, \dots, T - 1$:

- 1 Given current state $X_t = \mathbf{x}$, sample $\mathbf{p} \sim \mathcal{N}(\mathbf{p}; 0, M)$ and **propose** (deterministically)

$$(\mathbf{y}, \mathbf{p}') = \Phi_\tau(\mathbf{x}, \mathbf{p}) \quad \text{i.e., } q((\mathbf{y}, \mathbf{p}') \mid (\mathbf{x}, \mathbf{p})) = \delta_{\Phi_\tau(\mathbf{x}, \mathbf{p})}(\mathbf{y}, \mathbf{p}')$$

- 2 Compute acceptance probability:

$$\begin{aligned} \alpha((\mathbf{x}, \mathbf{p}) \rightarrow (\mathbf{y}, \mathbf{p}')) &= \min(1, \exp(-H(\mathbf{y}, \mathbf{p}') + H(\mathbf{x}, \mathbf{p}))) \\ &= 1 \quad \text{if ODE solver has no numerical error} \end{aligned}$$

- 3 Accept or reject:

$$\text{Sample } U \sim \text{Uniform}(0, 1), \quad X_{t+1} = \begin{cases} \mathbf{y} & \text{if } U \leq \alpha \\ \mathbf{x} & \text{else} \end{cases}$$

- Issue: L and ε are sometimes **hard to tune**
 - Problem specific, brittle hyperparameters
- NUTS (**No U-Turn Sampler**): Variant of HMC which
 - Auto-tunes ε in a warm-up phase
 - Picks L adaptively per MH step (integration stops when trajectory want to “turn around”)

HMC DEMO

<https://chi-feng.github.io/mcmc-demo/app.html>

DIAGNOSING MCMC

- Recall that for a **stationary** Markov Chain, the MCMC estimator

$$\hat{\mathbb{E}}_T = \frac{1}{T} \sum_{t=1}^T f(X_t), \quad X_t \sim \pi$$

has variance

$$\text{var}[\hat{\mathbb{E}}_T] \approx \frac{\sigma_{\text{as}}^2}{T}, \quad \sigma_{\text{as}}^2 = \text{var}_{\pi}[f(X)] + 2 \sum_{k=1}^{\infty} \text{cov}_{\pi}(f(X_0), f(X_k)).$$

- Note:** This formula is only valid **after burn-in**

Diagnostics

When running MCMC at time t , we want to answer **two questions**:

- Are we **in stationarity**, i.e., is $p_t \approx \pi$?
- How **large** is σ_{as}^2 , i.e., how **inefficient** is the chain ?

- We suffer from **initialization bias** since $p_1 \neq \pi$
- Thus, discard the first B samples (**burn-in phase**)
- Unfortunately, there is no good data-driven way to pick B 😞
- Note: B only affects **bias** and is **not related** to σ_{as} !

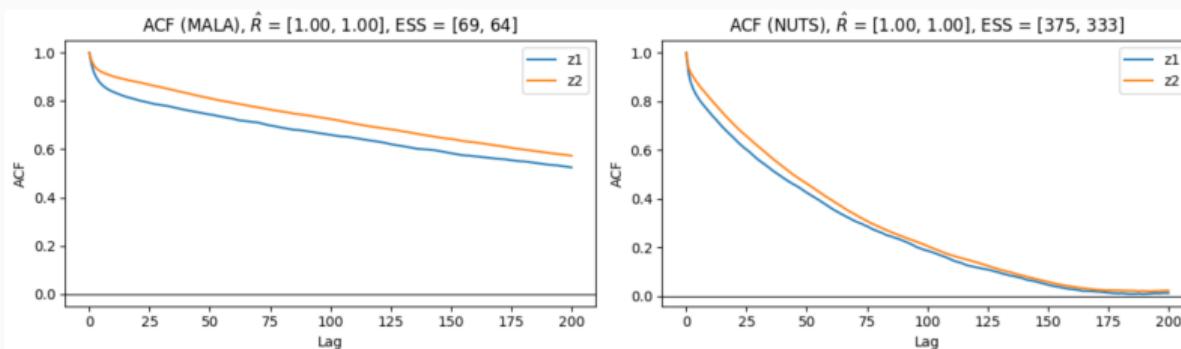
AUTOCORRELATION FUNCTION (ACF)

- For a **stationary** chain, we can write

$$\rho_k = \frac{\text{cov}_\pi(f(X_0), f(X_k))}{\text{var}_\pi(f(X))}, \quad \sigma_{\text{as}}^2 = \text{var}_\pi[f(X)] \left(1 + 2 \sum_{k=1}^{\infty} \rho_k\right)$$

where ρ_k is the **autocorrelation** at lag k .

- We can plot an estimate $\hat{\rho}_k$ over k after burn-in:



- ACF diagnoses **within-chain dependence**, not convergence to π !

EFFECTIVE SAMPLE SIZE (ESS)

- We would like to define **Effective Sample Size (ESS)** such that

$$\text{var} \left[\hat{\mathbb{E}}_T \right] = \frac{\text{var}_{\pi}[f(X)]}{\text{ESS}}$$

- ESS tells us **how many i.i.d. samples would give the same variance !**
- For a **stationary** chain, we thus have

$$\text{ESS} = \frac{\text{var}_{\pi}[f(X)]}{\sigma_{\text{as}}^2/T} = \frac{T}{1 + 2 \sum_{k=1}^{\infty} \rho_k}$$

which we can estimate from **samples post burn-in**

- We want to check if MCMC has **converged** to “stationary” distribution
 - By construction, we will **converge** to the **unique stationary distribution** π as $T \rightarrow \infty$
 - However, for finite T , we can get **stuck** in quasi-stationary distributions $\neq \pi$!
- Run $M \geq 2$ **independent chains** and discard **burn-in**
- Idea: Compare **within-chain variance** W with **between-chain variance** B :

$$\hat{R} := \sqrt{\frac{(1 - \frac{1}{T})W + \frac{1}{T}B}{W}}$$

- If all chains converged to the same distribution, we should have $\hat{R} \approx 1$
 - Often, one checks if $\hat{R} \leq 1.01$

- Goal: Sample from **complicated, high-dimensional distribution π**
 - Where typically only $\tilde{\pi}$ is known
- Almost all MCMC algorithms presented are **instances of Metropolis-Hastings**
 - Random-walk MH, Gibbs Sampling, MALA, HMC, NUTS
- In **high dimensions**, HMC/NUTS is the **go-to choice**
 - If you have **tractable conditionals**, Gibbs sampling may be a **good option**
- However, even well-tuned MCMC methods **can take a long time to converge**
 - This is called **mixing time**
- Diagnosis of MCMC is **hard in general**
 - Common tools: ACF, ESS, \hat{R}